

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

Tuowen Zhao
School of Computing
University of Utah
ztuowen@cs.utah.edu

Protonu Basu
Facebook
protonu@fb.com

Samuel Williams
Computational Research Division
Lawrence Berkeley National Lab
swwilliams@lbl.gov

Mary Hall
School of Computing
University of Utah
mhall@cs.utah.edu

Hans Johansen
Computational Research Division
Lawrence Berkeley National Lab
hjohansen@lbl.gov

ABSTRACT

Stencil computations in real-world scientific applications may contain multiple interrelated stencils, have multiple input grids, and use higher order discretizations with high arithmetic intensity and complex expression structures. In combination, these properties place immense demands on the memory hierarchy that limit performance. Blocking techniques like tiling are used to exploit reuse in caches. Additional fine-grain data blocking can also reduce TLB, hardware prefetch, and cache pressure.

In this paper, we present a code generation approach designed to further improve tiled stencil performance by exploiting reuse within the block, increasing instruction-level parallelism, and exposing opportunities for the backend compiler to eliminate redundant computation. It also enables efficient vector code generation for CPUs and GPUs. For a wide range of complex stencil computations, we are able to achieve substantial speedups over tiled baselines for the Intel KNL, Intel Skylake-X, and NVIDIA P100 architectures.

CCS CONCEPTS

• **Software and its engineering** → **Source code generation**; • **Computing methodologies** → *Parallel programming languages*.

KEYWORDS

Compiler Optimization, Stencil, vectorization

ACM Reference Format:

Tuowen Zhao, Protonu Basu, Samuel Williams, Mary Hall, and Hans Johansen. 2019. Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs. In *The International Conference for High Performance Computing, Networking, Storage, and Analysis (SC '19)*, November 17–22, 2019, Denver, CO, USA. ACM, New York, NY, USA, 13 pages. <https://doi.org/10.1145/3295500.3356210>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SC '19, November 17–22, 2019, Denver, CO, USA

© 2019 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-6229-0/19/11...\$15.00

<https://doi.org/10.1145/3295500.3356210>

1 INTRODUCTION

Stencil computations are ubiquitous in scientific applications that solve partial differential equations using the finite difference or finite volume methods, where the derivative at each point in space is calculated as a weighted sum of neighboring point values (a “stencil”). A stencil’s *order of accuracy* is the exponent on the relationship between grid spacing (array size) and error — both small grid spacings (large arrays) and high order can result in low error. A stencil’s order greatly impacts the optimizations needed to achieve high performance. Low-order discretizations result in smaller stencils that have limited data reuse, are typically bound by memory bandwidth, and thus underutilize the compute capability afforded by manycore, wide vector, and GPU architectures. Much of the prior work in this field has been based on lower order stencils and has thus focused on techniques to reduce main memory data movement [6, 13, 14, 20, 21, 23, 28, 33, 36, 38, 41, 45].

As processor architectures become more compute-intensive [37], computational scientists are increasingly turning to high-order schemes that perform more computation per point (more compute-intensive) but can attain equal error with larger grid spacings (smaller arrays). Although higher-order stencils inherently result in higher arithmetic intensity, they also place immense pressure on register file, cache, TLB, and hardware prefetchers. Worse still, to further utilize available compute capability, stencil computations are often a composition of multiple high-order stencils, such as the 8th-order hypterm kernel, described in [12], and depicted in Figure 1. It computes five stencils that operate on eight input fields.

Prior work on optimizing high-order stencils leverages the associativity of the weighted sums in a stencil computation; such operations can therefore be safely reordered to achieve the same result within round-off tolerances. Consequently, execution order can be optimized to exploit data reuse, and thus reduce memory load/store operations and reduce register pressure [3, 10, 24, 25, 29]. Prior associative reordering methods for stencils are limited in several ways. Most focus on reuse of individual data elements, with an eye towards optimizing scalar registers [24, 25]. Where reuse of vectors is considered to support vector code generation, it is limited to isotropic, constant-coefficient stencils [3], or arises from a post-pass vectorization, preceded by DLT (data-layout transformation) optimization [29]. In some cases, cross-iteration reuse is identified as a byproduct of loop unrolling [10, 25]. Only one of

these approaches targets GPUs, and it exploits reuse just within an expression [24].

This paper addresses these limitations, describing a vector code generator for general stencil computations targeting both CPUs and GPUs. It identifies data reuse *without unrolling* within a fine-grain block of a stencil computation. For further optimization gains, this approach to reuse analysis and vectorization can also work in tandem with a fine-grained blocked data layout that decomposes the original grid domain into small, fixed-size multi-dimensional subdomains [2, 17, 40], such as *bricks* [44], which have been shown to achieve performance portability across CPU and GPU. Bricks are stored contiguously in memory to enable a number of optimizations. First, accesses within a brick are part of a single address stream, mitigating the negative impact of blocking on hardware prefetchers and TLB. Second, when combined with *vector folding* [39], an individual dimension can be smaller than the vector width; this flexibility can reduce cache and register pressure for complex stencils like *hyperterm*. Other stencil optimizations such as temporal blocking and wavefront parallelism are beyond the scope of this paper, but are complementary and can be combined with our method.

This paper makes the following contributions: (1) it presents a vector code generation algorithm for general stencil computations that exploits data reuse within a block without unrolling, and targets both CPUs and GPUs; (2) it compares the effectiveness of the code generation approach for iteration space tiling vs. bricks on CPUs, isolating the benefits of each; (3) it offers the first description of node-level vector code generation for bricks; (4) it presents performance results on 24 stencils, including real-world proxy stencils such as *hyperterm*, demonstrating performance gains on Intel Knights Landing (Xeon Phi) processors (up to 3.4 \times), Intel Xeon Skylake-X (1.3 \times), and NVIDIA P100 (1.6 \times).

2 BACKGROUND AND MOTIVATION

In this section, we motivate our approach, using the *hyperterm* kernel, with code and the compiler’s expression tree shown in Figure 1. Stencils like *hyperterm* exhibit high temporal reuse across stencil iterations, e.g., `cons[imx][k][j][i+1]` and `cons[imx][k][j][i-1]` two iterations later. It is common to use tiling to exploit this reuse in caches or unrolling/unroll-and-jam to enable optimizations for reuse in registers. Additional array common subexpressions within and across expressions, such as the results of the shaded operators at the bottom of Figure 1, can also be reused in registers. Due to the complexity of *hyperterm*, exploiting such register reuse can lead to severe register pressure; exposing cross-iteration reuse using unrolling may increase register pressure, and even cause instruction cache misses. In addition, *hyperterm* has high arithmetic intensity, with 358 floating-point operations per iteration. Achieving high performance also demands efficient use of wide SIMD units in CPUs and SIMT threads in GPUs.

Another consideration is that *hyperterm* places immense pressure on the TLB and hardware prefetcher due to the number of independent data streams. One k - j plane of *hyperterm* requires 133 simultaneously active read or write data streams (corresponding to different registers, cache lines and potentially, TLB entries). Tiling will exacerbate this problem. To reduce the number of data streams for such stencils, prior work has developed variations of *blocked*

data layouts, where the original grid domain is decomposed into small, fixed-sized multi-dimensional subdomains [2, 17, 40]. In this paper, we expand on the concept of *bricks*, where these subdomains are stored contiguously in memory [44]. Using an $8 \times 8 \times 8$ brick size and stencil radius ≤ 8 , we access the elements within a brick using a single stream as opposed to 64 streams for a tiled code. The computation inside one brick would be similar to an $8 \times 8 \times 8$ tiled stencil.

Taken together, this paper describes a vector code generator that can balance the aforementioned optimization requirements of high-order stencils such as *hyperterm*. Our approach exploits reuse within a multi-dimensional data block, arising from either tiling, which reorders the computation, or bricks, which also reorganizes the data layout. As stencils are known to pose challenges to vectorization due to issues of alignment [15], the approach must expose aligned vector operations. Additionally, our approach further reduces arithmetic intensity by exposing opportunities for array common subexpression elimination [10]. The remainder of this section provides the foundation for the code generation approach.

2.1 Stencils as Gather or Scatter Operations

The kernel of a stencil computation typically contains a weighted sum of neighboring points. Such sums are most commonly expressed as *gather* operations, as in the 5-point 2D stencil code of Figure 2(a), where the value of this sum is calculated for each iteration of a loop nest by gathering its neighboring inputs (some of which are widely spaced in memory), individually weighting them, and summing them. Figure 3(a) visualizes this gather computational pattern for the 5-point stencil code.

However, one can observe that these weighted sums are associative and can be reordered without changing the meaning of the computation. This concept is *associative reordering*. Therefore, an alternative implementation of the 5-point stencil is a *scatter* operation, where one input is weighted and scattered to all the neighboring points that use it as a term in the sum. Figure 3(b) shows the resultant scatter pattern for the 5-point stencil. Scatters have several advantages including minimizing the number of loads, and improving instruction level parallelism, but may increase the number of stores. For high-order stencils that access a large number of inputs to compute each output point, an approach that favors reducing loads is preferable to one that reduces stores. We also find that the output data often resides in registers, particularly on GPUs, or in L1 cache, so store cost is typically low. Scatter also matches the strengths and weaknesses of bricks. Loads are more costly, because accesses that cross brick boundaries introduce indexing overhead due to the adjacency list, and unaligned loads are not applicable.

In the following, we will select some portions of the computation with high reuse to be computed using scatter.

2.2 Overview of Approach

The current code generator uses a domain-specific frontend, implemented in Python, that accepts stencil descriptions as input, shown in Figure 2(c). The output of the code generator is integrated into C code as in Figure 2(d). From this specification, we can generate either tiled or brick code that incorporates scatter, as used in the experiments of Section 5. The data layout for the tiled code is a

```

flux[irho][k][j][i] =
-dxinvo0 * (
    ALP * cons[imx][k][j][i + 1] + -ALP * cons[imx][k][j][i - 1] +
    BET * cons[imx][k][j][i + 2] + -BET * cons[imx][k][j][i - 2] +
    GAM * cons[imx][k][j][i + 3] + -GAM * cons[imx][k][j][i - 3] +
    DEL * cons[imx][k][j][i + 4] + -DEL * cons[imx][k][j][i - 4]
);
flux[irho][k][j][i] +=
-dxiniv1 * (
    ALP * cons[imy][k][j + 1][i] + -ALP * cons[imy][k][j - 1][i] +
    BET * cons[imy][k][j + 3][i] + -BET * cons[imy][k][j - 2][i] +
    GAM * cons[imy][k][j + 3][i] + -GAM * cons[imy][k][j - 3][i] +
    DEL * cons[imy][k][j + 4][i] + -DEL * cons[imy][k][j - 4][i]
);
flux[irho][k][j][i] += -dxiniv2 * ...

flux[imy][k][j][i] =
-dxinvo0 * (
    ALP * cons[imy][k][j][i + 1] * qup1 + -ALP * cons[imy][k][j][i - 1] * qum1 +
    BET * cons[imy][k][j][i + 2] * qup2 + -BET * cons[imy][k][j][i - 2] * qum2 +
    GAM * cons[imy][k][j][i + 3] * qup3 + -GAM * cons[imy][k][j][i - 3] * qum3 +
    DEL * cons[imy][k][j][i + 4] * qup4 + -DEL * cons[imy][k][j][i - 4] * qum4
);
flux[imy][k][j][i] +=
-dxiniv1 * (
    ALP * cons[imy][k][j + 1][i] * qvp1 + -ALP * cons[imy][k][j - 1][i] * qvm1 +
    ALP * q[qpres][k][j + 1][i] + -ALP * q[qpres][k][j - 1][i] +
    BET * cons[imy][k][j + 2][i] * qvp2 + -BET * cons[imy][k][j - 2][i] * qvm2 +
    BET * q[qpres][k][j + 2][i] + -BET * q[qpres][k][j - 2][i] +
    ... );

```

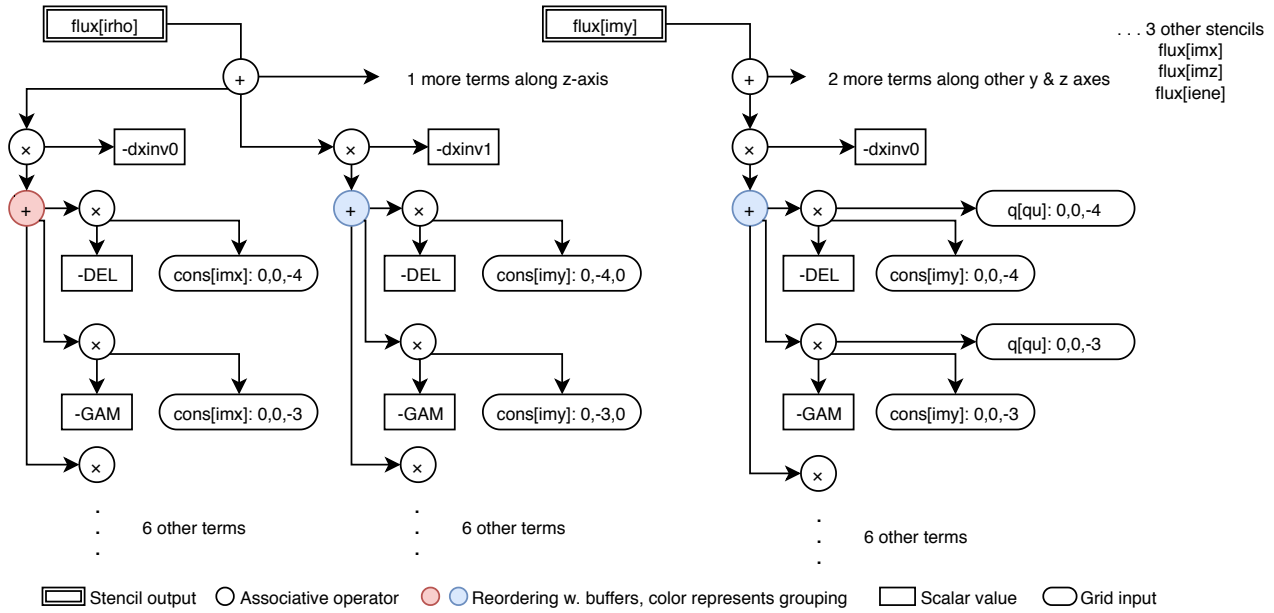


Figure 1: CNS’s hypterm stencil excerpt (top) and derived expression trees (bottom).

3D array. Individual bricks are stored in contiguous memory, and a collection of bricks that represents a domain are organized as a graph using an adjacency list. To compute a stencil, one iterates over all indices of bricks and, for each brick, computes the stencil at all (k, j, i) in the dimensions of a brick.

To detect reuse within and across stencils, we formulate the problem on an expression directed acyclic graph (DAG) of the stencil kernel as in Figure 1. We identify the operands in the DAG that are reused within or across iterations of the same expression. A profitability analysis determines whether a scatter should be used to optimize the redundant loads, and derives an iteration schedule. Operators containing operands for which a scatter is profitable are marked for reordering. The unmarked portions of the computation will use gather operations. We then group the marked subexpressions into stages to be computed together to capitalize on reuse across subexpression DAGs. Common subexpressions, which by definition use the same input, are likely to be grouped

together. Indirectly, this may result in common subexpression elimination (ASE) [3, 10], as the backend compiler can more easily detect such common subexpressions if they are adjacent instructions.

With the dimensions of the block, we use the stages and scatter schedule to produce vectorized code. Given the results of analysis, the code generator derives new loop bounds for the resulting tile, constrained by the boundaries of the buffers, with loop peeling as needed to implement the full block. Thus, instead of unrolling first and identifying reuse patterns in the unrolled code, we identify reuse based on indexing expressions of operands, and create the loops indicated by the profitability analysis. The code generator performs a few additional optimizations during vectorization. Our approach, with vectorization, is portable across both CPUs and GPUs. The code generation technique is detailed in the next two sections. Section 3 discusses the analysis that identifies reuse and decides how to split the computation into stages. Section 4 describes the actual vector code generation.

```

1  for (j = tj; j < tj + 4; ++j)
2  for (i = ti; i < ti + 4; ++i) {
3      c = In[j][i] * coeff[0]
4        + In[j][i+1] * coeff[1]
5        + In[j][i-1] * coeff[2]
6        + In[j+1][i] * coeff[3]
7        + In[j-1][i] * coeff[4];
8      Out[j][i] = c * vel[j][i];
9  }

```

(a) Stencil in a tiled region.

```

1  # Declarations
2  i = Index(0) ...
3  In = Grid("In", 2) ...
4  coeff = [ConstRef('coeff[0]'), ...]
5
6  c = In(i, j) * coeff[0] + In(i+1, j) * coeff
7    [1] + In(i-1, j) * coeff[2] + In(i, j+1) *
8    coeff[3] + In(i, j-1) * coeff[4]
9
10 Out(i, j).assign(c)

```

(c) Input to the code generator: kernel.py

```

1  float buf[4][4];
2  /*
3   * buf computed using vector scatter
4   * as in Figure 5.
5   */
6  // Vectorization directive
7  for (j = tj; j < tj + 4; ++j)
8  for (i = ti; i < ti + 4; ++i) {
9      Out[j][i] = buf[j-tj][i-ti] * vel[j][i];
10 }

```

(b) Split stencil with buffer to expose reuse.


```

1  // tile control loops
2  for (...) // Tile starting at tj, ti
3      tile("kernel.py", "FLEX", (4,4), // Tile
4          dimension
5            ("tj", "ti"));
6
7  // iterating over all brick
8  for (...) // brick index b
9      brick("kernel.py", "AVX512", (4,4),
10          (2,4) /* Folding */ , b);

```

(d) Adding kernel to C code.

Figure 2: A 5-point stencil example.



(a) Gather

(b) Scatter

Figure 3: Gather vs. scatter operations for 5 points in 2D.

3 REUSE-BASED EXPRESSION SPLITTING

This section describes how to split a stencil kernel into compute stages based on its reuse pattern. We first build an expression directed acyclic graph (DAG) from the code. For simplicity, we illustrate the algorithm using the running example of Figure 2 whose expression DAG is shown in Figure 4, but also refer to the DAG for hyterm in Figure 1 when discussing operator grouping. The code generation framework obtains this graph by identifying the assignments to grids in the original abstract syntax tree (AST), and the operators, constants and grid references on the right hand side. The output grid is the target of the DAG, and the operand grids are annotated with their name and offset from the iterator. In Figure 2, with iterator $[j, i]$, reference $\text{In}[j][i-1]$ is represented by node $\text{In}: 0, -1$. Neighboring associative operators of the same type are combined to a single operator with three or more terms, resulting in the 5-way addition in Figure 4. Observe that such DAGs can represent a broad variety of stencil codes: variable coefficients, multiple stencils, and several inputs multiplied by the same coefficient.

In this phase we start from the expression DAG using two major steps (1) identify reuse profitability and select associative operators to be reordered by postorder traversal of the expression DAG; (2) identify opportunities to group operators across expressions into stages to further improve data reuse.




Figure 4: Expression DAG for the 5-point stencil of Figure 2.

3.1 Reuse Profitability in Operators

The first step of the code generation algorithm is to mark associative operators in the expression DAG for reordering based on a profitability analysis. We define the profitability of scatter using the reduction of the number of inputs that are simultaneously live for each iteration step. We calculate profitability using a postorder traversal of the expression DAG, and we then mark the operators to be reordered that exceed a profitability threshold.

The postorder traversal collects $R(E)$, grids and offsets from the DAG for expression E , as in Equation 1. The number of elements in this set is then the number of unique reads when calculating this subexpression as how they appear in the expression DAG.

$$R(E) = \{ \langle g, \vec{o} \rangle | \text{grid } g \text{ appears in } E \text{ with offset } \vec{o} \} \quad (1)$$

When E is an associative operator, it consists of multiple terms (subexpressions), E_i . When we shift the terms of the associative operator between iterations we are effectively adding a per-term constant $\vec{\delta}_i$ to the offset of the corresponding term. This shift, $\vec{\delta}_i$, is sometimes referred to as the retiming vector in loop shifting [29]. If we add the shifts to all terms of the operator then we can collect the set of grids and offsets with Equation 2. The number of elements in this set is then the number of unique reads when calculating the associative operator with each term shifted by $\Delta = \{ \vec{\delta}_i \}$.




Figure 5: Illustration of profitability analysis and operator marking for the 5-point stencil in Figure 2. During DAG traversal (dashed arrow), Algorithm 1 is used to try to minimize distinct references by adding shifts to each term (numbers on edge). If the operator should be reordered (per Equation 3), the DAG is marked with shift δ_i .

$$R(E^\Delta) = \bigcup_i \{ \langle g, \vec{o} + \vec{\delta}_i \rangle \mid \langle g, \vec{o} \rangle \in R(E_i) \} \quad (2)$$

Using the cardinality of the two sets from Equation 1 and Equation 2, we can evaluate the profitability P of a reordered expression E^Δ as in Equation 3. Reordering is marked as profitable if reads are reduced by a large fraction; that is, when P exceeds a global threshold $t_1 \geq 1$.

$$P(E, E^\Delta) = \frac{|R(E)|}{|R(E^\Delta)|} \geq t_1 \quad (3)$$

Note that $|R(E)|$ is a property of the stencil computation. In order to maximize P , we only need to find a set of shifts, Δ , that minimize $|R(E^\Delta)|$.

This process is illustrated in Figure 5 for the 5-point stencil. During post-order traversal, when an associative operator is encountered, Algorithm 1 is used to inspect its terms; this is denoted in the figure by a dashed arrow. We note that in the original computation there are five distinct references, thus $|R(E)| = 5$. The algorithm assigns each term with one shift value that is marked on the edges; with shift this produces the same reference: In: 0, -1. Thus, from Equation 2, the number of distinct references after the shift is 1. This gives us a reuse profitability of 5 from Equation 3. Assuming this profitability is above the predefined threshold t_1 , this addition is marked in the original graph with the shift values produced from Algorithm 1, δ , annotated to edges leading to each term.

The number of possible shift amounts for one term is related to the radius of the stencil; the total search space is exponential in the number of terms. This search space is potentially too large to search exhaustively. However, the set of offsets that minimizes $P(E^\Delta)$ has the properties of Equation 4: the minimum only arises when for each term, it either has no common grid input with any other terms, or it has at least one read that can be reused after shifting.

$$\forall E_i, \text{ either} \quad (4)$$

$$\begin{cases} \forall \langle g, \vec{o} \rangle \in R(E_i), g \notin R(E - \{E_i\}), \text{ or} \\ \exists \langle g, \vec{o} \rangle \in R(E_i), \langle g, \vec{o} + \vec{\delta}_i \rangle \in R((E - \{E_i\})^{\Delta - \{\delta_i\}}) \end{cases}$$

PROOF. Equation 4 can be proven using contradiction by assuming the negation: $|R(E^\Delta)|$ is the minimum but has an E_i that shares a common grid input and does not have reuse with shift δ_i . In this case, we can change δ_i so that at least one read is reused with some other term. We have the new $|R(E^\Delta')|$, which will be smaller by at

least one. It contradicts the assumption that Δ^* gives the minimum. Thus, Equation 4 must be true. \square

Based on (4), we developed a fast greedy algorithm in Algorithm 1. The complexity of this algorithm is $O(n^4 \log n)$ where $n = \sum_i |R(E_i)|$. This n is bounded by the number of reads in the original stencil code, N . In fact, this is only a loose upper bound for the complexity, and for associative operators that only have one offset for all reads in one term, such as the hypertext stencil in Figure 1 and many of other stencils, its complexity is only $O(n)$. After the scatter is decided, we can then compute the profitability and mark associative operators for grouping.

Algorithm 1 Greedy algorithm for deciding shift amounts for child subexpressions of E

```

Initialize  $\Delta = \{\vec{\delta}_i\} = \emptyset$ 
repeat
  for all child subexpressions  $E_i$  do
    for all  $\langle g, \vec{o} \rangle \in R(E_i)$  do
      for all  $\langle g', \vec{o}' \rangle \in R(E_{j \neq i}), g = g'$  do
         $\vec{\delta}'_i = \vec{o}' + \vec{\delta}_j - \vec{o}$   $\triangleright$  Compute new shift
        If obtains a lower  $|R(E^\Delta)|$ ,  $\vec{\delta}_i = \vec{\delta}'_i$   $\triangleright$  Update  $\Delta$ 
      end for
    end for
  end for
until  $\Delta$  isn't updated
return  $\Delta$ 

```

3.2 Cross-operator Reuse

Step 2 in our framework will enable optimization of loads across parts of the expression DAG by attempting to group reordered operators utilizing the same data as input to be computed together. Each marked computation can be computed once all the values from its subexpression are available, which includes all subexpressions that are marked. This creates a dependence graph that includes all marked subexpressions where edges are contracted from the original expression DAG.

Beyond optimizing for reuse across subexpressions, operator grouping also impacts the number of buffers that are simultaneously

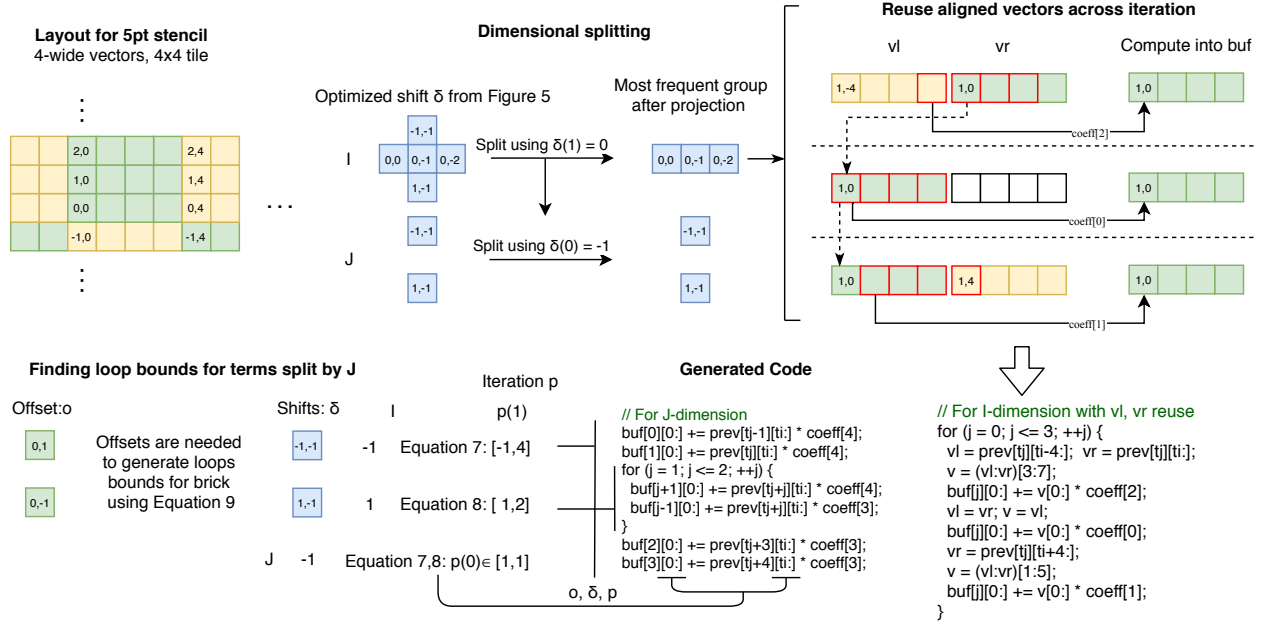


Figure 6: Vector code generation for 5-point stencil example in Figure 2,5. The generated code employs dimensional splitting and reuses aligned vector load.

live. We use a modified topological sort based on a priority tuple measure similar to the one used in [25] to break ties during the sort when multiple alternatives are present. Applying this process results in a linearized sequence that tries to balance buffer usage and the number of operators that can be computed concurrently. This step has a complexity of at most $O(m^2)$ where m is the number of marked subexpressions.

We can then scan the sequence and merge nearby operators into one stage based on a measure of hardware pressure, which is derived from (1) the number of distinct inputs; and, (2) the number of buffers accessed. Both of them can represent the grouped operators' pressure on the registers. The first value can be computed using Algorithm 1. To determine viability, we compute a weighted sum of the two measures after the grouping and compare against a global threshold t_2 representing the architectural constraint, as in Equation 5, where \mathbb{E} represents the group of operators' expressions and $|B|$ is the number of buffers in the current group. Merging can be attempted at most twice for each buffer. This results in a similar complexity of Algorithm 1 where $n = N$, the number of reads in the original stencil code. Since $N > 1$, $m < N$, we have a total complexity of $O(N^4 \log N)$. As noted previously, for stencils whose associative operator has only one offset for each term, this complexity reduces to $O(m^2 + N)$.

$$|R(\mathbb{E}^\Delta)| + k|B| \leq t_2 \quad (5)$$

Consider the example in Figure 1. Our approach recognizes that there is potential for reuse within the group of associative operators colored by blue. Note that adding the blue operator on the left to a group consisting of the blue operator on the right will not affect the first term of Equation 5. The equation only approaches the threshold

by the extra buffer. However, when adding the associative operator colored red to this group, both terms increase and the threshold is reached faster, preventing aggressive grouping.

4 VECTOR CODE GENERATION

Once all compute stages are created, we obtain a sequence of groups of operators in the expression DAG. For any groups not identified by the algorithm as profitable for reordering, a gather, using the original computation, will be generated. Other stages contain associative operators that are profitable for reordering; these operators will have a shift, $\tilde{\delta}_i$, associated with each term, $\{E_i\}$, of the operator. We can then generate code for each stage in the sequence using either gather or scatter. In this section, we describe how vector code is generated from $\{E_i\}$ and Δ with loops. We also present additional optimizations that we used. A walk through of the code generation is shown in Figure 6 for the 5-point stencil.

The techniques in this section are based on scanning the iteration space and the expression DAG side-by-side, which results in a complexity of $O(|S||V|)$, where $|S|$ is the size of the tile and $|V|$ is the number of nodes in the expression DAG.

4.1 Vector Scatter

While we focus on scattered computation with vectors, this is derived from reordering using scalar values. We first use the shifts computed for each term. For iteration \tilde{p} within the tile, we can determine the shifted loop index, the source of the scatter, by subtracting $\tilde{\delta}_i$ as in Equation 6. Each of the sources will scatter the corresponding subexpression that is shifted by adding $\tilde{\delta}_i$ to all the original array indices. The scatter can then walk through all such \tilde{p}'

and compute the value for each term to the destination with $\vec{p}' + \vec{\delta}_i$.

$$\vec{p}' = \vec{p} - \vec{\delta}_i \quad (6)$$

When generating vector code, we add restrictions to the destination so that only destinations that are aligned are written. To accommodate wide vectors and small data blocking, we assume the vector is a multidimensional vector-sized rectangle, that exhibits a length $v(d)$ on each tile dimension d . Then we have an aligned destination whose offset within the tile is a multiple of the vector length on all dimensions. With this reduced set of destinations, the meaning of the calculation is preserved since every location in the block is in one of the aligned vectors.

We notice that we can generate some loops to reduce the code size, since $\vec{\delta}_i$ reflects the amount of loop shifting applied to the term. $T(d)$ represents the tile size on dimension d . We can formulate the loop bounds as in Equations 7-9. Equation 7 represents the total range of the shifted iteration space. All or part of this range can use loops while some iterations might need to be peeled from either side of this range.

$$\begin{cases} 0 - \max_i(\delta_i(d)), \\ T(d) - v(d) - \min_i(\delta_i(d)) \end{cases} \quad (7)$$

Equation 8 represents when the writes for all terms are inside the original tiled space. A loop can be generated on the tiled code for all iterations in this range.

$$\begin{cases} 0 - \min_i(\delta_i(d)), \\ T(d) - v(d) - \max_i(\delta_i(d)) \end{cases} \quad (8)$$

Equation 9 represents when all reads are inside the blocked data region. For bricks, a loop can be generated by combining Equation 8 and Equation 9.

$$\begin{cases} 0 - \min_i(o_i(d) + \delta_i(d)), \\ T(d) - v(d) - \max_i(o_i(d) + \delta_i(d)) \end{cases} \quad (9)$$

The lower left of Figure 6 shows how loop bounds are inferred for two of the terms that are split from the original stencil using Equations 7 and 8.

4.2 Optimizations

We employ two additional optimizations when generating scatter for one stage, (1) further data reuse within vector registers using aligned vector load and vector align instructions; and, (2) dimensional splitting which divides the stencil expression and calculates it along one dimension at a time [18] to further reduce peeling. These optimizations are illustrated in the upper half of Figure 6.

We noticed that more reuse can be exploited with an aligned vector load. This optimization is especially useful for bricks where a logical vector may cross the brick boundary and must be merged using vector aligns. This merging, using `alignr` intrinsics on the CPU, is happening for almost every iteration when we are traversing the contiguous direction while the two aligned vectors will be the same for several iterations. These aligned loads can then be cached and reused for consecutive iterations. Upper right of Figure 6 shows that we can create two temporary vectors `v1` and `vr` to cache the aligned vectors. With this optimization, we only load these values three times, $2 + 0 + 1$, instead of five, $2 + 1 + 2$.

We also employ dimensional splitting of the stencil to further reduce the final code size. We observe that Equation 8 is only related to the shift selected for each term, and it is often more constrained than Equation 9. If we directly generate code for the 5-point stencil example in Figure 6, the calculation for terms on the I-dimension is unnecessarily peeled because of the two terms on the J-dimension. We can then group the terms based on their values to increase the range of Equation 8. We achieve this by picking one dimension at a time and group the terms based on their shifts on the other dimensions. We then select the most frequent groups to be computed together. Larger loops can be created for the other dimensions where they have a common shift value. We repeat this process for each of the dimensions to fully split the stencils. This process is applied to the 5-point stencil in Figure 6. For stencils such as CNS, Figure 1, we can achieve a perfect split as peeling only happens for one dimension at a time. Without dimensional splitting, no loops are possible for the stencil with $8 \times 8 \times 8$ tile.

4.3 Vectorizing on GPU

As observed by prior work [46], GPUs offer the same vector merging capability as `alignr` intrinsics on the CPU using either shared memory or shuffle instructions, `__shfl_up` and `__shfl_down`. This allows us to transfer our vectorization method onto NVIDIA GPUs and use each warp as one vector that has a length of 32. Also, through use of vector folding and combinations of multiple shuffles we enable support for smaller data blocks such as $8 \times 8 \times 8$ or $4 \times 4 \times 4$.

5 EXPERIMENTAL RESULTS

This section presents performance results for the generated code for both CPUs and GPUs, applied to tiled or brick code.

5.1 Target Architectures

Intel Knights Landing. The Intel Xeon Phi 7250 Knights Landing (KNL) has 68 physical cores organized into a 2D on-chip mesh of 34 tiles each with two CPU cores¹ and a shared 1MB L2 cache. Each core has a private 32KB L1 data cache, implements 4-way multi-threading, and has two AVX-512 vector processing units (VPUs). AVX-512 instructions operate on 8 double-precision or 16 single-precision floating-point data elements in a SIMD fashion. Its theoretical peak performance is 2611.2 GFLOP/s double-precision fused multiply-and-add. Each chip has both standard DDR4 DRAM memory and high-bandwidth MCDRAM memory that we configured as a last level cache using the *quadcache* mode, which yields a peak STREAM performance of 332 GB/s.

Intel Xeon Gold (Skylake-X). The Intel Xeon Gold 6130 CPU has 16 physical cores. Each core has a private 32KB L1 data cache and 1MB L2 cache, implements 4-way multithreading, and has two AVX-512 vector processing units (VPUs). Each core has a nominal frequency of 2.1GHz. The whole CPU has a theoretical peak performance is 1075.2 GFLOP/s. Concurrently, 6 DDR4 memory controllers provide a STREAM bandwidth of 85 GB/s.

NVIDIA P100. The P100 GPU has 56 streaming multiprocessors. Each streaming multiprocessor has 64 single-precision and 32 double-precision CUDA cores and has a warp size of 32. Each

¹We use 32 tiles for a total of 64 cores in all our experiments to isolate system overhead.




Figure 7: Performance on KNL and Skylake-X. For real-world stencil kernels, smaller blocking sizes may be beneficial because more input grids can put higher pressure on the cache. For higher-order and real-world stencils, the brick approach often provides the best speedup: On KNL, *hyp term* – 3.4 \times and *f3d125pt* – 1.9 \times . On Skylake-X, *hyp term* – 1.3 \times and *f3d125pt* – 2.1 \times .

streaming multiprocessor has a dedicated texture/L1 cache. The P100 has a theoretical peak single-precision performance of 9.3 TFLOP/s, a peak double-precision performance of 4.7 TFLOP/s, and a GPU-STREAM [9] bandwidth of 586 GB/s.

5.2 Benchmarks and Proxy Codes

We use three categories of stencil kernels of varying order shown in Table 1. The first category includes smoothers from the HPGMG benchmark suite [1]. The second category includes high-order stencil computations from the Compressible Navier-Stokes computation [12]. The third category includes synthetic stencils to capture the memory-compute ratio of different kinds of stencil shapes and radii. Many of the real stencils are a composition of these kernel patterns. Stencils are named according to their class (“f” or “s”) and the number of points where each point is weighted individually. Class “s” refers to stencils for the Laplacian second derivatives, that only operate on elements along each of the axes (star-shaped) in each dimension. For Class “s”, the stencil radius is just half the order; there are no off-axis points in the stencil. Class “f” refers to stencils for the “compact” Laplacian that touch all points in a cube (dense), with Manhattan distance equal to the radius. The radius is again just half the order; in some applications these stencils can produce more accurate solutions.

The baseline code is written in C. Our code generator generates code from the stencil written as python expressions in a python script, as in Figure 2(c). This specification can be used as a standalone DSL or as an intermediate output from the parser. The thresholds described in Section 3 are exposed as parameters to the

Name	Grid	FLOPS	Name	Grid	FLOPS
Stencils from HPGMG					
chebyshev	6	39	poisson	2	21
helm-v2	7	22	helm-v4	7	115
Stencils from CNS					
hyp term	13	358	diff term	11	415
2D Star-Shaped Stencils					
s2d5pt	2	9	s2d9pt	2	17
s2d13pt	2	25	s2d17pt	2	33
s2d21pt	2	41	s2d25pt	2	49
s2d29pt	2	49	s2d33pt	2	49
2D Full Stencils					
f2d9pt	2	17	f2d25pt	2	49
f2d49pt	2	97	f2d81pt	2	161
3D Star-Shaped Stencils					
s3d7pt	2	13	s3d13pt	2	25
s3d19pt	2	37	s3d25pt	2	49
3D Full Stencils					
f3d27pt	2	53	f3d125pt	2	249

Table 1: Stencils used in experiments. Grid represents the number of grids the stencil operates on, while FLOPS represents the number of FLOPS performed per point.

code generator. We used $t_1 = 1.5$ (reuse estimate), $t_2 = 20$ (simultaneously active buffers), and $k = 2$ (weight assigned to buffers)

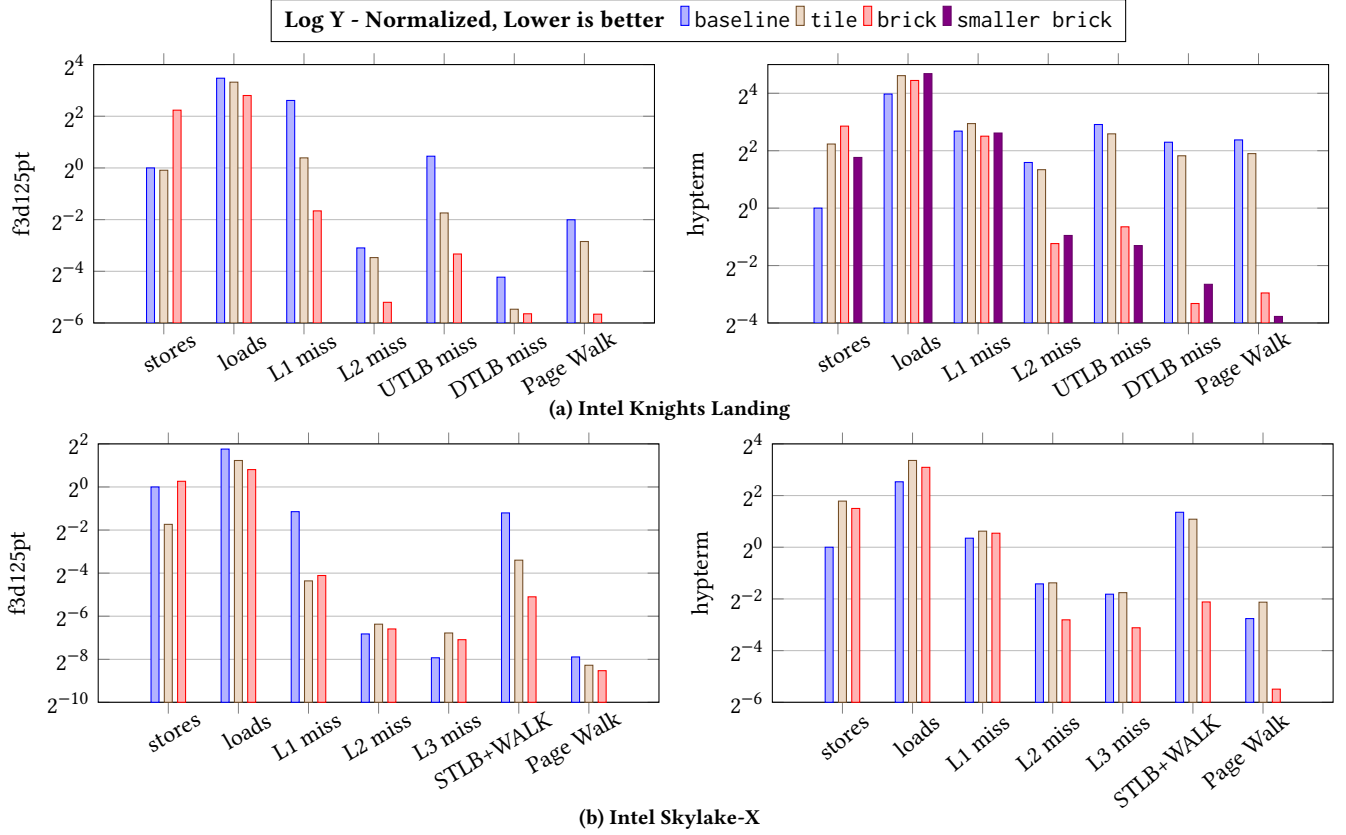


Figure 8: Profiling metrics on KNL and Skylake-X. Each bar is the raw counter value, normalized to the baseline total number of stores. The generated code provides a dramatic reduction in cache and TLB misses and page walks. On KNL, L1 L2 and TLB misses are reduced as much as 19 \times , 7 \times , and 49 \times respectively. On Skylake-X, we also reduced L1 misses by up to 8 \times and TLB-related metrics up to 14 \times . These indicates much better data locality and cache reuse.

for our experiments. For all the stencils our code generator runs in under two seconds. Further discussion on the relationship between the brick library, the code generator, and our choices for these parameters is discussed in Section 5.5.

5.3 CPU Performance: KNL and Skylake-X

Figure 7 shows performance in GFLOP/s for the 24 stencils in our experiment, running on KNL and Skylake-X. We compare the performance of three code versions:

- a baseline version where the stencils are expressed as a gather. The stencils are tiled on all dimensions and parallelized using OpenMP. 2D stencils are parallelized using threads on the outermost tile control loop (J-dimension). 3D stencils are parallelized using threads on the two outermost tile control loops (K-, J-dimensions). The innermost tiled loop (I-dimension) is vectorized using `#pragma omp simd`. For 3D stencils we use three different tile sizes: $4 \times 4 \times 4$, $4 \times 4 \times 8$, and $8 \times 8 \times 8$. For 2D stencils we use two sizes: 8×8 , and 4×8 . This version reports the best performance out of these tile sizes.

- a tile version, that uses the same thread schedule and tiling as baseline but applies our code generator. Only reusing aligned loads is not applied as it attempts to be more general by keeping the array-format references. Vectorization is done using `#pragma omp simd`.
- a brick version, where fine-grained data blocking is used [44] with our code generator. The thread schedule is kept the same, and vectorization also uses the OpenMP `simd` directive. Aligned data loads and merge are implemented using intrinsics to load vectors across brick boundaries. Reusing aligned load is applied. Here we try to separate the effect of different brick sizes by using brick to denote the performance for sizes $8 \times 8 \times 8$ and 8×8 and use smaller brick to represent other smaller sizes presents in the baseline.

The baseline and tile are compiled with either `-O2` or `-Ofast`, whichever achieves the best performance. The brick version is compiled using `-O2` only.

The results for the two CPU architectures in Figure 7 suggest several observations. The trend of performance relative to the shape and size of the stencil can be shown in the synthetic stencils. For star-shaped stencils, both 2D and 3D, the relative performance of tile to the baseline increases with larger stencil diameters due

to more temporal reuse. For lower order stencils, brick is slower than both other versions due to indexing overhead. However, since higher order stencils exhibit much higher temporal reuse, brick becomes the fastest version due to both the data blocking from bricks and operator reordering.

For the full stencils, on KNL, tile is able to obtain speedup on 3D but not on 2D stencils. This is due to 3D stencils exhibiting much higher reuse for each loaded input. Due to the alignment constraints, for each element read, f3d125pt is able to scatter to at most 25 points, while for f2d81pt it can only have up to 9 points.

Although tile is often slower than brick, it is faster for f3d125pt on KNL. Here, the generated code size becomes the bottleneck for brick because Equation 9 limits the region that can use loops. For tile, the code size for the entire kernel including outer loops is around 34KB for 4×4×8, which is when tile gives the best performance. Due to the extra indexing calculation and much wider peeled region from data indirection, the code size for bricks is 54KB. This is much higher than the L1 instruction cache size of 32KB.

We used Intel VTune Amplifier to obtain profiling results to further dissect the performance difference between different versions. We selected two of the complex stencils to show in Figure 8. When comparing brick to baseline on KNL, cache misses are reduced by up to 19×, and TLB misses are reduced by up to 49×. Similarly, on Skylake-X, brick also reduced L1 misses by up to 8× and TLB-related metrics up to 14×. In addition, even though the code generator increases the number of stores, these appear to be serviced from L1 due to the decrease in cache misses as compared to baseline. The effect of better locality is more pronounced on KNL with more threads and much smaller cache per thread. For f3d125pt, our code generation can also reduce the total number of loads.

For the real-world stencils we noticed that tile offers similar or worse performance compared to baseline. This is due to the inherently higher L1 cache pressure for these stencils. Table 1 shows the number of grids referenced for each of these stencils. While we prefer each of the grids to be located in the fastest cache available for reuse, the buffers from vector scatter increase L1 pressure and detract from the better locality it provides. As seen in Figure 8, the L1 miss count is even across all versions of the code for hyp term. As noted previously, KNL threads have less cache capacity. baseline and tile tend to achieve the best performance on 4×4×8 tiles or sometimes 4×4×4 tiles. The smaller brick version, which relies on vector folding, improves KNL performance due to reduced stores and improved TLB behavior.

5.4 GPU Performance: NVIDIA P100

Figure 9 presents NVIDIA P100 performance. We compare the performance of three code variants:

- baseline version (a gather) that is tiled using the thread-block decomposition of CUDA, where each CUDA block is comprised of 4×4×32 (K, J, I) threads for 3D or 8×32 for 2D. Each thread computes one stencil output. We also could spawn 32 threads for each block and iterate on the (K, J) dimension to further imitate the number of threads used for the brick version, but it is always slower.

- tile version that applies our code generation on a tiled thread-block, where it compute (K, J, I) elements using I threads. The tile size is the same as the baseline. Only aligned loads and reuse of those are not applied.
- brick version where fine-grained data blocking is used [44] with the code generator. The subdomain that each CUDA block will compute is the same as baseline, but only 32 threads (one warp) are in each block to compute all stencils in the subdomain, like a vector with width of 32. We also tried smaller sizes such as 4×4×8 and 4×4×4 for real-world stencils; these are reported as smaller brick. Note that the I-dimension of these sizes are smaller than the vector width on the GPU which is the warpsize, 32.

In Figure 9, the baseline shows very little performance variation. This is because FLOP/s are limited by data dependences. Our approach reduces this bottleneck by exploiting input reuse. The NVCC compiler can generate code where the buffers introduced by vector scatter reside in registers. With sufficient registers, one element is read and scattered to multiple destinations in registers so that the write cost is low compared to read. The effect of holding buffers in registers is reflected in the drop of performance between (s2d21pt, s2d25pt) and (s3d13pt, s3d19pt). Also note that due to alignment, fewer registers are required than for gather. For example, for f3d125pt only 25 destination register are used when scattering one input, whereas in gather 125 input are used.

We used NVProf to obtain several metrics related to memory performance for the f3d125pt and hyp term stencils in Figure 10. Our code generator improves register and cache reuse significantly; the number of global loads are reduced by as much as 12×. Higher cache reuse results in lower L1 pressure that contributes to reduction in the volume of global load traffic seen at L2. Some of the generated code temporarily stages values in local memory, which can be identified by the local load and store metric (L-Load/Store). The tile code also shows higher L2 traffic. The performance impact of these increases is significantly outweighed by the higher global and L2 loads of the baseline code. hyp term, one of the more complex stencils, exhibits increased register and cache pressure, resulting in extra local load and stores and HBM traffic even when using bricks. However, with a reduced brick size and vector folding, this effect is completely eliminated. It is possible that TLB behavior is improved on the GPU as it was on the CPU, but such metrics are unavailable in NVProf.

5.5 Discussion

In this subsection, we consider the performance impact of varying the configuration of the compiler and code generator to tease out contributions of different aspects of our approach. Suppose, for example, the brick library were used without the vector code generation. The brick library uses template expansion for address calculation and incorporates indirection to represent neighboring bricks; consequently, without code generation, backend compilers will introduce redundant index calculations and cannot vectorize on CPUs. In fact, stencils in the naive brick library version might even be slower than the baseline. For the more compute-intensive stencils such as the 125pt, our code generator provides an 18× speedup when compared to the naive brick library version on KNL.




Figure 9: Performance on NVIDIA P100. Brick code generation produces the best performance for many of the stencils. For synthetic stencils, using either 32 threads or full block results has little effect, but smaller block sizes reduce cache pressure for real stencils. Bricks speedup many stencils, for example poisson (1.6 \times) and f3d125pt (7.0 \times).




Figure 10: Profiling metrics on P100. Each bar represent transaction count normalized to the total number of global stores of baseline. brick version reduces global loads by up to 12 \times and L2 loads by up to 2.6 \times implying much better locality.

Even on P100, where these issues are less pronounced, the fully optimized code is 7.5 \times faster.

The thresholds used in the experiment are fixed using lenient values. These thresholds are especially relaxed for simple stencils. Our reuse estimate of $t_1 = 1.5$ applies to all associative operators that have reuse potential in our tests. Aside from the value 1 that denotes no potential reuse, the lowest reuse potential is $12/7 \approx 1.86$ in helm-v2. This is due to the fact that helm-v2 contains terms using multiple grid references that are not perfectly reused after applying the shifts.

Our estimate of $t_2 = 20$, $k = 2$ is a proxy for how much register pressure the algorithm incurs. These criteria are also only effective for complex stencils and are rarely met for simple stencils. All synthetic stencils have a value for Equation 5 of 3. For complex stencils like helm-v4 and hypterm, it is theoretically possible to reach this limit. However, this does not happen as our code generator strategy requires reuse between operators in one stage but limits the number of operators that can be put in one stage. This requirement may be proven to be too greedy and further tuning of these parameters is yet to be explored.

6 RELATED WORK

Optimization efforts for stencil computations can be broadly classified as memory access optimization techniques, and optimization methods to improve computation, although in practice, there is often significant interaction between them.

Most of the optimization effort has focused on stencils applied on large grids that are usually bound by capacity or compulsory cache misses, leading to a variety of studies on spatial and temporal tiling [6, 7, 11, 13, 19–23, 26–28, 33–36, 38, 41, 45]. In addition, domain-specific compilers have recently been developed for parallel code generation from a stylized stencil specification [4, 30, 42, 43] or from a code excerpt [16].

The aforementioned tiling techniques have focused on loop or iteration space tiling. In addition to loop tiling, researchers have also tiled or blocked data (space). Data along with loop tiling efforts have been addressed by [2, 17, 31, 40]. TiDA [31] uses coarse-grained data blocking, where the entire grid is tiled into sub-grids, each with its own ghost zone. Fine-grained data blocking is explored in Bricks [44] and Briquettes [17], YASK [40] and RTM on the Cell processor [2]. All the fine-grained blocking techniques targeted

large, compute-intensive stencils, and the small data blocks (bricks) do not have per-block ghost zones.

The fine-grained data blocks used in our research are similar to briquettes in [17], but there is significant difference in our approaches. Briquettes were designed to perform 3D stencils split into 1D stencils, thus requiring multiple sweeps to compute the output. Furthermore, a data transpose was required between each 1D stencil sweep to ensure good SIMDization. Their code generation required data staging tailored for 1D stencils. In contrast to Briquettes, we optimize 3D stencils without manual dimensional splitting and perform complex stencil reordering in addition to fine-grained data blocking to improve computation by reducing reads and improving SIMDization.

YASK is a C++ template-based approach to generating code for large stencils with fine-grained data blocks. YASK autotuned their data block size, and used smaller data blocks than our method (e.g. $2 \times 2 \times 4$ instead of 8^3). They can generate code for clusters of vectors using unrolling and common expression elimination to improve reuse, which is less feasible for complex stencils. They did not directly target stencil reordering as presented in our paper.

Stencil reordering, one of the main characteristics of our approach, has been explored in different ways. Manual optimization of stencil computations has led to techniques such as *semi-stencils* to reduce loads [8] and using common subexpression elimination after unrolling to reduce floating-point computations, improve register reuse, reduce register pressure, or improve instruction level parallelism [5, 7, 25]. Some works target specific properties of the associative operation in stencils. Deitz et al. [10] describes an automated approach to common subexpression elimination for sum-of-product computation and is not applicable to uniquely weighted or variable coefficient stencils where no common subexpressions exist. Basu et al. [3] uses partial sums to reorder constant-coefficient isotropic stencils. Stock et al. [29] uses statement splitting to enable loop shifting to expose reuse of the same input and autotuning to determine the shift amount. They also do not target code generation for the GPU. In comparison, the research presented in this paper illustrates a new powerful stencil reordering method which works on general stencils without manual optimizations. Our method targets tiled stencil computation to improve cache and register reuse, and is designed to work with fine-grained data blocking on modern architectures with wide SIMD units.

High-order PDEs can also be implemented as dense matrix operations as in [32]; this paper and other previous compiler/code-generation research treats higher-order stencils as computations on structured grids. Using dense linear algebra primitives, although technically feasible for finite difference and finite volume methods, would be inefficient for our stencils as most of the entries in the resultant matrices would be zero.

7 CONCLUSION

High-order stencil computations are simultaneously best-suited to the trends in computer architecture (limited bandwidth coupled with high arithmetic intensity) and most-often underperforming. To that end, in this paper, we introduce a novel compiler optimization to exploit reuse and vectorization in block stencil computations. When coupled with a fine-grained blocked data layout (bricks) this

produces code that reduces vector loads and alignment operations, exposes opportunities to eliminate redundant computation, and reduces the data footprint of stencils in the memory hierarchy. We show our approach improves the performance of real stencils compared to a tiled baseline by up to $3.4\times$ on a Intel Knights Landing (Xeon Phi) processor, up to $1.3\times$ on Intel Xeon Skylake-X, and up to $1.6\times$ on NVIDIA P100.

ACKNOWLEDGMENTS

This research was supported by the Exascale Computing Project (17-SC-20-SC), a joint project of the U.S. Department of Energy's Office of Science and National Nuclear Security Administration. This research used resources in Lawrence Berkeley National Laboratory and the National Energy Research Scientific Computing Center, which are supported by the U.S. Department of Energy Office of Sciences Advanced Scientific Computing Research program under contract number DE-AC02-05CH11231.

REFERENCES

- [1] 2016. High-Performance Geometric Multigrid. <http://hpgmg.org>
- [2] Mauricio Araya-Polo, Félix Rubio, Raúl de la Cruz, Mauricio Hanzich, José María Cela, and Daniele Paolo Scarpazza. 2009. 3D Seismic Imaging Through Reverse-time Migration on Homogeneous and Heterogeneous Multi-core Processors. *Sci. Program.* 17, 1-2 (Jan. 2009), 185–198. <https://doi.org/10.1155/2009/382638>
- [3] Protonu Basu, Mary Hall, Samuel Williams, Brian Van Straalen, Leonid Oliker, and Phillip Colella. 2015. Compiler-directed transformation for higher-order stencils. In *Parallel and Distributed Processing Symposium (IPDPS)*, 2015 *IEEE International*. IEEE, 313–323.
- [4] M. Christen, O. Schenk, and H. Burkhart. 2011. PATUS: A Code Generation and Autotuning Framework for Parallel Iterative Stencil Computations on Modern Microarchitectures. In *Parallel Distributed Processing Symposium (IPDPS)*. <https://doi.org/10.1109/IPDPS.2011.70>
- [5] Kaushik Datta. 2009. *Auto-tuning Stencil Codes for Cache-Based Multicore Platforms*. Ph.D. Dissertation. EECS Department, University of California, Berkeley.
- [6] Kaushik Datta, Shoaib Kamil, Samuel Williams, Leonid Oliker, John Shalf, and Katherine Yelick. 2009. Optimization and Performance Modeling of Stencil Computations on Modern Microprocessors. *SIAM Rev.* 51, 1 (2009), 129–159.
- [7] Kaushik Datta, Mark Murphy, Vasily Volkov, Samuel Williams, Jonathan Carter, Leonid Oliker, David Patterson, John Shalf, and Katherine Yelick. 2008. Stencil Computation Optimization and Auto-Tuning on State-of-the-art Multicore Architectures. In *Supercomputing (SC)*.
- [8] Raúl De La Cruz, Mauricio Araya-Polo, and José María Cela. 2010. Introducing the Semi-stencil Algorithm. In *International Conference on Parallel Processing and Applied Mathematics: Part I (PPAM)*. 11.
- [9] Tom Deakin, James Price, Matt Martineau, and Simon McIntosh-Smith. 2016. GPU-STREAM v2. 0: benchmarking the achievable memory bandwidth of many-core processors across diverse parallel programming models. In *International Conference on High Performance Computing*. Springer, 489–507.
- [10] Steven J Deitz, Bradford L Chamberlain, and Lawrence Snyder. 2001. Eliminating redundancies in sum-of-product array computations. In *Proceedings of the 15th international conference on Supercomputing*. ACM, 65–77.
- [11] Craig C. Douglas, Jonathan Hu, Markus Kowarschik, Ulrich Rüde, and Christian Weiss. 2000. Cache Optimization for Structured and Unstructured Grid Multigrid. *Elect. Trans. Numer. Anal.* 10 (2000), 21–40.
- [12] Matthew Emmett, Weiqun Zhang, and John B Bell. 2014. High-order algorithms for compressible reacting flow with complex chemistry. *Combustion Theory and Modelling* 18, 3 (2014), 361–387.
- [13] M. Frigo and V. Strumpén. 2005. Evaluation of cache-based superscalar and cacheless vector architectures for scientific computations. In *Proc. ACM International Conference on Supercomputing (ICS)*.
- [14] P. Ghysels, P. Kosiewicz, and W. Vanroose. 2012. Improving the arithmetic intensity of multigrid with the help of polynomial smoothers. *Numerical Linear Algebra with Applications* 19, 2 (2012), 253–267.
- [15] Tom Henretty, Kevin Stock, Louis-Noël Pouchet, Franz Franchetti, J Ramanujam, and P Sadayappan. 2011. Data layout transformation for stencil computations on short-vector simd architectures. In *Compiler Construction*. Springer, 225–245.
- [16] Justin Holewinski, Louis-Noël Pouchet, and P. Sadayappan. 2012. High-performance code generation for stencil computations on GPU architectures. In *International Conference on Supercomputing (ICS)*.

- [17] Jagan Jayaraj. 2013. *A strategy for high performance in computational fluid dynamics*. Ph.D. Dissertation. University of Minnesota.
- [18] Jagan Jayaraj, Pei-Hung Lin, Paul R Woodward, and Pen-Chung Yew. 2014. CFD builder: A library builder for computational fluid dynamics. In *Parallel & Distributed Processing Symposium Workshops (IPDPSW), 2014 IEEE International*. IEEE, 1029–1038.
- [19] Markus Kowarschik and Christian WeiSS. 2001. DiMEPACK - A Cache-Optimized Multigrid Library. In *International Conference on Parallel and Distributed Processing Techniques and Applications (PDPTA), volume 1*.
- [20] Sriram Krishnamoorthy, Muthu Baskaran, Uday Bondhugula, J. Ramanujam, Atanas Rountev, and P Sadayappan. 2007. Effective automatic parallelization of stencil computations. In *Proc. ACM SIGPLAN conference on Programming language design and implementation (PLDI)*.
- [21] J. McCalpin and D. Wonnacott. 1999. *Time skewing: A value-based approach to optimizing for memory locality*. Technical Report DCS-TR-379. Department of Computer Science, Rutgers University.
- [22] Paulius Micikevicius. 2009. 3D Finite Difference Computation on GPUs Using CUDA. In *Proceedings of 2nd Workshop on General Purpose Processing on Graphics Processing Units (GPGPU-2)*. 6.
- [23] Anthony Nguyen, Nadathur Satish, Jatin Chhugani, Changkyu Kim, and Pradeep Dubey. 2010. 3.5-D Blocking Optimization for Stencil Computations on Modern CPUs and GPUs. In *Proc. ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC)*.
- [24] Prashant Singh Rawat, Fabrice Rastello, Aravind Sukumaran-Rajam, Louis-Noël Pouchet, Atanas Rountev, and P. Sadayappan. 2018. Register Optimizations for Stencils on GPUs. In *Proceedings of the 23rd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP '18)*. ACM, New York, NY, USA, 168–182. <https://doi.org/10.1145/3178487.3178500>
- [25] Prashant Singh Rawat, Aravind Sukumaran-Rajam, Atanas Rountev, Fabrice Rastello, Louis-Noël Pouchet, and P. Sadayappan. 2018. Associative Instruction Reordering to Alleviate Register Pressure. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage, and Analysis (SC '18)*. IEEE Press, Piscataway, NJ, USA, Article 46, 13 pages. <http://dl.acm.org/citation.cfm?id=3291656.3291718>
- [26] G. Rivera and C. Tseng. 2000. Tiling Optimizations for 3D Scientific Computations. In *Supercomputing (SC)*.
- [27] S. Sellappa and S. Chatterjee. 2004. Cache-Efficient Multigrid Algorithms. *International Journal of High Performance Computing Applications* 18, 1 (2004), 115–133.
- [28] Y. Song and Z. Li. 1999. New tiling techniques to improve cache temporal locality. In *Proc. ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*.
- [29] Kevin Stock, Martin Kong, Tobias Grosser, Louis-Noël Pouchet, Fabrice Rastello, Jagannathan Ramanujam, and Ponnuswamy Sadayappan. 2014. A framework for enhancing data reuse via associative reordering. In *ACM SIGPLAN Notices*, Vol. 49. ACM, 65–76.
- [30] Yuan Tang, Rezaul Alam Chowdhury, Bradley C. Kuszmaul, Chi-Keung Luk, and Charles E. Leiserson. 2011. The pochoir stencil compiler. In *ACM symposium on Parallelism in algorithms and architectures*.
- [31] Didem Unat, Tan Nguyen, Weiqun Zhang, Muhammed Nufail Farooqi, Burak Bastem, George Michelogiannakis, Ann Almgren, and John Shalf. 2016. *TiDA: High-Level Programming Abstractions for Data Locality Management*. Springer International Publishing, Cham, 116–135.
- [32] Jerry E. Watkins, Joshua Romero, and Antony Jameson. 2016. Multi-GPU, Implicit Time Stepping for High-order Methods on Unstructured Grids. In *46th AIAA Fluid Dynamics Conference*. American Institute of Aeronautics and Astronautics. <https://doi.org/10.2514/6.2016-3965>
- [33] Gerhard Wellein, Georg Hager, Thomas Zeiser, Markus Wittmann, and Holger Fehske. 2009. Efficient Temporal Blocking for Stencil Computations by Multicore-Aware Wavefront Parallelization. In *International Computer Software and Applications Conference*. <https://doi.org/10.1109/COMPSAC.2009.82>
- [34] S. Williams, J. Carter, L. Oliker, J. Shalf, and K. Yelick. 2008. Lattice Boltzmann Simulation Optimization on Leading Multicore Platforms. In *International Conference on Parallel and Distributed Computing Systems (IPDPS)*.
- [35] Samuel Williams, Leonid Oliker, Jonathan Carter, and John Shalf. 2011. Extracting ultra-scale Lattice Boltzmann performance via hierarchical and distributed auto-tuning. In *Supercomputing (SC)*.
- [36] S. Williams, J. Shalf, L. Oliker, S. Kamil, P. Husbands, and K. Yelick. 2006. The potential of the Cell processor for scientific computing. In *Proc. Conference on Computing Frontiers*.
- [37] S. Williams, A. Watterman, and D. Patterson. 2009. Roofline: An Insightful Visual Performance Model for Floating-Point Programs and Multicore Architectures. *Commun. ACM* (April 2009).
- [38] D. Wonnacott. 2000. Using Time Skewing to Eliminate Idle Time due to Memory Bandwidth and Network Limitations. In *Proc. International Conference on Parallel and Distributed Computing Systems*.
- [39] C. Yount. 2015. Vector Folding: Improving Stencil Performance via Multi-dimensional SIMD-vector Representation. In *2015 IEEE 17th International Conference on High Performance Computing and Communications, 2015 IEEE 7th International Symposium on Cyberspace Safety and Security, and 2015 IEEE 12th International Conference on Embedded Software and Systems*. 865–870.
- [40] Charles Yount, Josh Tobin, Alexander Breuer, and Alejandro Duran. 2016. YASK-yet Another Stencil Kernel: A Framework for HPC Stencil Code-generation and Tuning. In *Proceedings of the Sixth International Workshop on Domain-Specific Languages and High-Level Frameworks for HPC (WOLFHPC '16)*. 10.
- [41] T. Zeiser, G. Wellein, A. Nitsure, K. Iglberger, U. Rude, and G. Hager. 2008. Introducing a parallel cache oblivious blocking approach for the lattice Boltzmann method. *Progress in Computational Fluid Dynamics* 8 (2008).
- [42] N. Zhang, M. Driscoll, C. Markley, S. Williams, P. Basu, and A. Fox. 2017. Snowflake: A Lightweight Portable Stencil DSL. In *2017 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*. 795–804.
- [43] Yongpeng Zhang and Frank Mueller. 2012. Auto-generation and auto-tuning of 3D stencil codes on GPU clusters. In *International Symposium on Code Generation and Optimization (CGO)*.
- [44] T. Zhao, S. Williams, M. Hall, and H. Johansen. 2018. Delivering Performance-Portable Stencil Computations on CPUs and GPUs Using Bricks. In *2018 IEEE/ACM International Workshop on Performance, Portability and Productivity in HPC (P3HPC)*. 59–70. <https://doi.org/10.1109/P3HPC.2018.00009>
- [45] Xing Zhou, Jean-Pierre Giacalone, María Jesús Garzarán, Robert H. Kuhn, Yang Ni, and David Padua. 2012. Hierarchical overlapped tiling. In *Proc. International Symposium on Code Generation and Optimization (CGO)*.
- [46] Gerhard Zumbusch. 2013. Vectorized Higher Order Finite Difference Kernels. In *Proceedings of the 11th International Conference on Applied Parallel and Scientific Computing (PARA'12)*. Springer-Verlag, Berlin, Heidelberg, 343–357. https://doi.org/10.1007/978-3-642-36803-5_25

Appendix: Artifact Description/Artifact Evaluation

SUMMARY OF THE EXPERIMENTS REPORTED

We developed a code generator using our approach described in the paper and applied to the 24 stencil kernels shown in the experiment section. We run the code generator and compile the output and baseline code individually for each of the platforms we use. These code variants are compiled using Intel Compiler 2019 and run on the KNL node on NERSC's Cori. They are also compiled using Intel Compiler 2019 and run on Notchpeak Skylake-X node from University of Utah's CHPC. For GPU, they are compiled using CUDA 9.2 and run on Kingspeak P100 node from University of Utah's CHPC.

ARTIFACT AVAILABILITY

Software Artifact Availability: All author-created software artifacts are maintained in a public repository under an OSI-approved license.

Hardware Artifact Availability: There are no author-created hardware artifacts.

Data Artifact Availability: There are no author-created data artifacts.

Proprietary Artifacts: None of the associated artifacts, author-created or otherwise, are proprietary.

List of URLs and/or DOIs where artifacts are available:

<https://bitbucket.org/ztuowen/vecscatter-artifact/sr>
↪ c/master/

BASELINE EXPERIMENTAL SETUP, AND MODIFICATIONS MADE FOR THE PAPER

Relevant hardware details: KNL 7250, Xeon Gold 6130, Nvidia P100

Compilers and versions: Intel C++ Compiler 2019.3 (GCC 8.2 compatibility) (KNL), Intel C++ Compiler 2019.2 (GCC 8.1 compatibility) (Skylake-X), CUDA 9.2 (GCC 6.4.0)

Libraries and versions: Intel OpenMP corresponding to each compiler version

Paper Modifications: No hardware or software modification. Our code generation generates source code that is input to the available compiler.

Output from scripts that gathers execution environment information.

P100 node:
LMOD_FAMILY_COMPILER_VERSION=6.4.0
SLURM_NODELIST=kp362
SLURM_CHECKPOINT_IMAGE_DIR=/var/slurm/checkpoint

```
MANPATH=/uufs/chpc.utah.edu/sys/installdir/gcc/6.4.0_
↪ /share/man:/uufs/chpc.utah.edu/sys/installdir/lm_
↪ od/7.7.29/share/man:/usr/kerberos/man:/usr/local_
↪ /share/man:/usr/share/man/overrides:/usr/share/m_
↪ an:/uufs/kingspeak.peaks/sys/pkg/slurm/std/share_
↪ /man
SLURM_JOB_NAME=bash
XDG_SESSION_ID=87380
_ModuleTable003_=UEMtMTgvQ29tcGlsZXIva3AvZ2NjLzYuNC4_
↪ wIiwL3V1ZnMvY2hwYy51dGFoLmVkdS9zeXMvbW9kdWx1Zm1_
↪ sZXMvQ0hQYy0xOC9Db21waWx1ci9nY2MvNi40LjAiLCIvdXV_
↪ mcy9jaHBjLnV0YWguZWR1L3N5cy9tb2R1bGVmaWx1cy9DSFB_
↪ DLTE4L0NvbXBpGVyL2N1ZGEvOS4yIiwL3V1ZnMvY2hwYy5_
↪ 1dGFoLmVkdS9zeXMvbW9kdWx1Zm1sZXMvQ0hQYy0xOC9MaW5_
↪ 1eCIsIi91dWZzL2NocGMudXRhaC5lZHUvc3lzl21vZHVzZWZ_
↪ pbGVzL0NIUEMtMTgvQ29yZSIi91dWZzL2NocGMudXRhaC5_
↪ lZHUvc3lzl2luc3RhbGxkaXIvbg1vZC83LjcuMjkvbW9kdWx_
↪ lZm1sZXMvQ29yZSIsfSxbInN5c3R1bUJhc2VNUeFUSCJdPSI_
↪ vdXVmcY9jaHBjLnV0YWguZWR1L3N5cy9tb2R1bGVmaWx1
HOSTNAME=kp362
SLURM_TOPOLOGY_ADDR=kp362
SLURMD_NODENAME=kp362
LMOD_FAMILY_CUDA=cuda
SLURM_PRIO_PROCESS=0
CUDA_BINDIR=/uufs/chpc.utah.edu/sys/installdir/cuda/_
↪ 9.2.148/bin
SLURM_SRUN_COMM_PORT=45416
CUDA_INCDIR=/uufs/chpc.utah.edu/sys/installdir/cuda/_
↪ 9.2.148/include
__LMOD_REF_COUNT_MODULEPATH=/uufs/chpc.utah.edu/sys/_
↪ modulefiles/CHPC-18/Compiler/kp/gcc/6.4.0:1;/uuf_
↪ s/chpc.utah.edu/sys/modulefiles/CHPC-18/Compiler_
↪ /gcc/6.4.0:1;/uufs/chpc.utah.edu/sys/modulefiles_
↪ /CHPC-18/Compiler/cuda/9.2:1;/uufs/chpc.utah.edu_
↪ /sys/modulefiles/CHPC-18/Linux:1;/uufs/chpc.utah_
↪ .edu/sys/modulefiles/CHPC-18/Core:1;/uufs/chpc.u_
↪ tah.edu/sys/installdir/lmod/7.7.29/modulefiles/C_
↪ ore:1
TERM=rxvt-unicode-256color
SHELL=/bin/bash
CUDA_LIBDIR=/uufs/chpc.utah.edu/sys/installdir/cuda/_
↪ 9.2.148/lib64
HISTSIZE=1000
SLURM_PTY_WIN_ROW=55
SLURM_JOB_QOS=QOS
LMOD_SYSTEM_DEFAULT_MODULES=chpc
MODULEPATH_ROOT=/uufs/chpc.utah.edu/sys/modulefiles/_
↪ CHPC-18
SSH_CLIENT=155.98.69.105 45432 22
SLURM_TOPOLOGY_ADDR_PATTERN=node
TMPDIR=/tmp
```



```

CUDA_ROOTDIR=/uufs/chpc.utah.edu/sys/installdir/cuda
↪ /9.2.148
LMOD_FAMILY_CUDA_VERSION=9.2
LMOD_PKG=/uufs/chpc.utah.edu/sys/installdir/lmod/7.7
↪ .29
SLURM_CPU_BIND_VERBOSE=quiet
QTDIR=/usr/lib64/qt-3.3
QTINC=/usr/lib64/qt-3.3/include
LMOD_VERSION=7.7.29
SSH_TTY=/dev/pts/82
__LMOD_REF_COUNT_LOADEDMODULES=chpc/1.0:1;cuda/9.2:1
↪ ;gcc/6.4.0:1
SLURM_CPU_BIND_LIST=
QT_GRAPHICSSYSTEM_CHECKED=1
OSVER=7.6.1810
USER=USER
SLURM_NNODES=1
LD_LIBRARY_PATH=/uufs/chpc.utah.edu/sys/installdir/g
↪ cc/6.4.0/lib64:/uufs/chpc.utah.edu/sys/installdi
↪ r/gcc/6.4.0/lib:/uufs/chpc.utah.edu/sys/installd
↪ ir/cuda/9.2.148/lib64:/uufs/kingspeak.peaks/sys/
↪ pkg/slurm/std/lib

```

```

LS_COLORS=rs=0:di=38;5;27:ln=38;5;51:mh=44;38;5;15:p
↪ i=40;38;5;11:so=38;5;13:do=38;5;5:bd=48;5;232;38
↪ ;5;11:cd=48;5;232;38;5;3:or=48;5;232;38;5;9:mi=0
↪ 5;48;5;232;38;5;15:su=48;5;196;38;5;15:sg=48;5;1
↪ 1;38;5;16:ca=48;5;196;38;5;226:tw=48;5;10;38;5;1
↪ 6:ow=48;5;10;38;5;21:st=48;5;21;38;5;15:ex=38;5;
↪ 34:*.tar=38;5;9:*.tgz=38;5;9:*.arc=38;5;9:*.arj=
↪ 38;5;9:*.taz=38;5;9:*.lha=38;5;9:*.lz4=38;5;9:*.
↪ lzh=38;5;9:*.lzma=38;5;9:*.tlz=38;5;9:*.txz=38;5
↪ ;9:*.tzo=38;5;9:*.t7z=38;5;9:*.zip=38;5;9:*.z=38
↪ ;5;9:*.Z=38;5;9:*.dz=38;5;9:*.gz=38;5;9:*.lrz=38
↪ ;5;9:*.lz=38;5;9:*.lzo=38;5;9:*.xz=38;5;9:*.bz2=
↪ 38;5;9:*.bz=38;5;9:*.tbz=38;5;9:*.tbz2=38;5;9:*.
↪ tz=38;5;9:*.deb=38;5;9:*.rpm=38;5;9:*.jar=38;5;9
↪ :*.war=38;5;9:*.ear=38;5;9:*.sar=38;5;9:*.rar=38
↪ ;5;9:*.alz=38;5;9:*.ace=38;5;9:*.zoo=38;5;9:*.cp
↪ io=38;5;9:*.7z=38;5;9:*.rz=38;5;9:*.cab=38;5;9:*.
↪ .jpg=38;5;13:*.jpeg=38;5;13:*.gif=38;5;13:*.bmp=
↪ 38;5;13:*.pbm=38;5;13:*.pgm=38;5;13:*.ppm=38;5;1
↪ 3:*.tga=38;5;13:*.xbm=38;5;13:*.xpm=38;5;13:*.ti
↪ f=38;5;13:*.tiff=38;5;13:*.png=38;5;13:*.svg=38;
↪ 5;13:*.svgz=38;5;13:*.mng=38;5;13:*.pcx=38;5;13:
↪ *.mov=38;5;13:*.mpg=38;5;13:*.mpeg=38;5;13:*.m2v
↪ =38;5;13:*.mkv=38;5;13:*.webm=38;5;13:*.ogm=38;5
↪ ;13:*.mp4=38;5;13:*.m4v=38;5;13:*.mp4v=38;5;13:*.
↪ .vob=38;5;13:*.qt=38;5;13:*.nuv=38;5;13:*.wmv=38
↪ ;5;13:*.asf=38;5;13:*.rm=38;5;13:*.rmvb=38;5;13:
↪ *.flc=38;5;13:*.avi=38;5;13:*.fli=38;5;13:*.flv=
↪ 38;5;13:*.gl=38;5;13:*.dl=38;5;13:*.xcf=38;5;13:
↪ *.xwd=38;5;13:*.yuv=38;5;13:*.cgm=38;5;13:*.emf=
↪ 38;5;13:*.axv=38;5;13:*.anx=38;5;13:*.ogv=38;5;1
↪ 3:*.ogx=38;5;13:*.aac=38;5;45:*.au=38;5;45:*.fla
↪ c=38;5;45:*.mid=38;5;45:*.midi=38;5;45:*.mka=38;
↪ 5;45:*.mp3=38;5;45:*.mpc=38;5;45:*.ogg=38;5;45:*.
↪ .ra=38;5;45:*.wav=38;5;45:*.axa=38;5;45:*.oga=38
↪ ;5;45:*.spx=38;5;45:*.xspf=38;5;45:
LMOD_sys=Linux
UUFSCELL=kingspeak.peaks
SLURM_STEP_NUM_NODES=1
SSH_AUTH_SOCK=/tmp/ssh-bviKxD3bkd/agent.23093
SRUN_DEBUG=3
SLURM_JOBID=6956232
_ModuleTable004_=cy9DSFBDLTE4L0xpbnV40i91dWZl2NocGM
↪ udXRhaC5lZHUvc3lzl21vZHVszWZpbGVzL0NIUEMtMTgvQ29
↪ yZTovdXVmcy9jaHBjLnV0YWguZWRL1L3N5cy9pbnN0YXxsZG1
↪ yL2xtb2QvNy43LjI5L21vZHVszWZpbGVzL0NvcnUiLH0=
TMOUT=0
__LMOD_REF_COUNT__LMFILES_=/uufs/chpc.utah.edu/sys/m
↪ odulefiles/CHPC-18/Core/chpc/1.0.lua:1;/uufs/chp
↪ c.utah.edu/sys/modulefiles/CHPC-18/Core/cuda/9.2
↪ .lua:1;/uufs/chpc.utah.edu/sys/modulefiles/CHPC-
↪ 18/Core/gcc/6.4.0.lua:1
SLURM_NTASKS=1
SLURM_LAUNCH_NODE_IPADDR=10.242.67.31
SLURM_STEP_ID=0

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```
__LMOD_REF_COUNT_NLSPATH=/uufs/chpc.utah.edu/sys/ins_
↳ talldir/intel/compilers_and_libraries_2018.1.163_
↳ /linux/compiler/lib/intel64/locale/%l_%t/%N:2;/u_
↳ ufs/chpc.utah.edu/sys/installdir/intel/compilers_
↳ _and_libraries_2018.1.163/linux/mkl/lib/intel64_
↳ lin/locale/%l_%t/%N:2;/uufs/chpc.utah.edu/sys/in_
↳ stalldir/intel/debugger_2018/gdb/intel64/share/l_
↳ ocale/%l_%t/%N:2
_ModuleTable001_=X01vZHVsvZVRHymx1Xz17WyJNVHZlcnNpb24_
↳ iXT0zLFsiY19yZWJ1aWxkVGltZSJdPWZhbnN1LFsiY19zaG9_
↳ ydFRpbWUiX1mYwzZSxkZXB0aFQ9e30sZmFtaWx5PXBtbikN_
↳ vbXBpbGVyIl09ImdjYyIsWyJjdWRHIl09ImN1ZGEiLH0sbVQ_
↳ 9e2NocGM9e1siZm4iXT0iL3V1ZnMvY2hwYy51dGFoLmVkdS9_
↳ zeXMvbW9kdWx1ZmlsZXMvQ0hQYy0xOC9Db3JlL2NocGMvMS4_
↳ wLmx1YSIsWyJmdWxsTmFtZSJdPSJjaHBjLzEuMCIsWyJsb2F_
↳ kT3JkZXIiXT0xLHByb3BUPXt9LFsic3RHY2tEZXB0aCJdPTA_
↳ sWyJzdGF0dXMiXT0iYWNoaXZlIixbInVzZXJOYw11l09ImN_
↳ ocGMiLH0sY3VkyYT17WyJmbiJdPSIvdXVmcY9jaHBjLnV0YWg_
↳ uZWR1L3N5cy9tb2R1bGVmaWxlcY9DSFBDLTE4L0NvcmlUv
MAIL=/var/spool/mail/USER
PATH=/uufs/chpc.utah.edu/sys/installdir/gcc/6.4.0/bi_
↳ n:/uufs/chpc.utah.edu/sys/installdir/cuda/9.2.14_
↳ 8/bin:/usr/lib64/qt-3.3/bin:/usr/kerberos/sbin:/_
↳ usr/kerberos/bin:/usr/lib64/ccache:/usr/local/bi_
↳ n:/usr/bin:/usr/local/sbin:/usr/sbin:/uufs/kings_
↳ peak.peaks/sys/pkg/slurm/std/bin:/opt/dell/srvad_
↳ min/bin:/uufs/chpc.utah.edu/common/home/USER/bin
SLURM_TASKS_PER_NODE=1
SLURM_STEP_LAUNCHER_PORT=45416
SLURM_WORKING_CLUSTER=kingspeak:10.242.67.13:6817:81_
↳ 92
_=/usr/bin/env
SLURM_JOB_ID=6956232
SLURM_STEP_GPUS=1
PWD=/uufs/chpc.utah.edu/common/home/USER/brickv2/Aut_
↳ hor-Kit
SLURM_STEPID=0
SLURM_JOB_USER=USER
_LMFILES=/uufs/chpc.utah.edu/sys/modulefiles/CHPC-1_
↳ 8/Core/chpc/1.0.lua:/uufs/chpc.utah.edu/sys/modu_
↳ lefiles/CHPC-18/Core/cuda/9.2.lua:/uufs/chpc.uta_
↳ h.edu/sys/modulefiles/CHPC-18/Core/gcc/6.4.0.lua
SLURM_SRUN_COMM_HOST=10.242.67.31
CUDA_VISIBLE_DEVICES=1
SLURM_CPU_BIND_TYPE=none
LANG=en_US.UTF-8
MODULEPATH=/uufs/chpc.utah.edu/sys/modulefiles/CHPC-_
↳ 18/Compiler/kp/gcc/6.4.0:/uufs/chpc.utah.edu/sys_
↳ /modulefiles/CHPC-18/Compiler/gcc/6.4.0:/uufs/ch_
↳ pc.utah.edu/sys/modulefiles/CHPC-18/Compiler/cud_
↳ a/9.2:/uufs/chpc.utah.edu/sys/modulefiles/CHPC-1_
↳ 8/Linux:/uufs/chpc.utah.edu/sys/modulefiles/CHPC_
↳ -18/Core:/uufs/chpc.utah.edu/sys/installdir/lmod/_
↳ 7.7.29/modulefiles/Core
SLURM_UMASK=0022
SLURM_PTY_WIN_COL=212
```

```
KDEDIRS=/usr
_ModuleTable_Sz_=4
LOADEDMODULES=chpc/1.0:cuda/9.2:gcc/6.4.0
SLURM_JOB_UID=1080628
SLURM_NODEID=0
SLURM_SUBMIT_DIR=/uufs/chpc.utah.edu/common/home/USE_
↳ R/brickv2/Author-Kit
LMOD_CMD=/uufs/chpc.utah.edu/sys/installdir/lmod/7.7_
↳ .29/libexec/lmod
SLURM_NPROCS=1
SLURM_TASK_PID=191863
SLURM_DISTRIBUTION=cyclic
SLURM_CPUS_ON_NODE=1
SSH_ASKPASS=/usr/libexec/openssh/gnome-ssh-askpass
HISTCONTROL=ignoredups
KRB5CCNAME=KEYRING:persistent:1080628
SLURM_PROCID=0
SLURM_JOB_NODELIST=kp362
SHLVL=3
HOME=/uufs/chpc.utah.edu/common/home/USER
SLURM_PTY_PORT=46262
__LMOD_REF_COUNT_PATH=/uufs/chpc.utah.edu/sys/instal_
↳ ldir/gcc/6.4.0/bin:1;/uufs/chpc.utah.edu/sys/ins_
↳ talldir/cuda/9.2.148/bin:1;/usr/lib64/qt-3.3/bin_
↳ :1;/usr/kerberos/sbin:1;/usr/kerberos/bin:1;/usr_
↳ /lib64/ccache:1;/usr/local/bin:1;/usr/bin:1;/usr_
↳ /local/sbin:1;/usr/sbin:1;/uufs/kingspeak.peaks/_
↳ sys/pkg/slurm/std/bin:2;/opt/dell/srvadmin/bin:1_
↳ ;/uufs/chpc.utah.edu/common/home/USER/bin:2
SLURM_LOCALID=0
_ModuleTable002_=Y3VkyS85LjIubHVhIixbImZ1bGx0YW11l0_
↳ 9ImN1ZGEvOS4yIixbImxvYWRPcmRlciJdPTIscHJvcFQ9e2F_
↳ yY2g9e1siZ3B1l09MSx9LH0sWyJzdGFja0RlcHRoIl09MCx_
↳ bIn0YXR1cyJdPSJhY3RpdmlfFsidXNlck5hbWUiXT0iY3V_
↳ kYS85LjIiLH0sZ2NjPXBtbImZuIl09Ii91dWZlL2NocGMudXR_
↳ haC5lZHUvc3lzl21vZHVsvZWZpbGVzL0NIUEMtMtgVQ29yZS9_
↳ nY2MvNi40LjAubHVhIixbImZ1bGx0YW11l09ImdjYy82LjQ_
↳ uMCIsWyJsb2FkT3JkZXIiXT0zLHByb3BUPXt9LFsic3RHY2t_
↳ EZXB0aCJdPTAsWyJzdGF0dXMiXT0iYWNoaXZlIixbInVzZXJ_
↳ OYw11l09ImdjYy82LjQuMCIsfSx9LG1wYXR0QT17Ii91dWZ_
↳ zL2NocGMudXRhaC5lZHUvc3lzl21vZHVsvZWZpbGVzL0NI
SLURM_CLUSTER_NAME=kingspeak
SLURM_JOB_CPUS_PER_NODE=1
SLURM_JOB_GID=2030
SLURM_SUBMIT_HOST=kingspeak1
SLURM_GTIDS=0
GCC_LIB=/uufs/chpc.utah.edu/sys/installdir/gcc/6.4.0_
↳ /lib64
__LMOD_REF_COUNT_INCLUDE=/uufs/chpc.utah.edu/sys/ins_
↳ talldir/gcc/6.4.0/include:1
BASH_ENV=/uufs/chpc.utah.edu/sys/installdir/lmod/7.7_
↳ .29/init/bash
SLURM_JOB_PARTITION=soc-gpu-kp
MODULERCFILE=/uufs/chpc.utah.edu/sys/modulefiles/etc_
↳ /rc
LOGNAME=USER
```

```

QTLIB=/usr/lib64/qt-3.3/lib
CVS_RSH=ssh
SLURM_STEP_NUM_TASKS=1
BASHRC_LOADED=1
XDG_DATA_DIRS=/uufs/chpc.utah.edu/common/home/USER/.
↳ local/share/flatpak/exports/share:/var/lib/flatp
↳ ak/exports/share:/usr/local/share:/usr/share
SSH_CONNECTION=155.98.69.105 45432 155.101.26.20 22
SLURM_JOB_ACCOUNT=ACCT
GPU_DEVICE_ORDINAL=1
MODULESHOME=/uufs/chpc.utah.edu/sys/installldir/lmod/
↳ 7.7.29
SLURM_JOB_NUM_NODES=1
__LMOD_REF_COUNT_LD_LIBRARY_PATH=/uufs/chpc.utah.edu
↳ /sys/installldir/gcc/6.4.0/lib64:1;/uufs/chpc.uta
↳ h.edu/sys/installldir/gcc/6.4.0/lib:1;/uufs/chpc.
↳ utah.edu/sys/installldir/cuda/9.2.148/lib64:1;/uu
↳ fs/kingspeak.peaks/sys/pkg/slurm/std/lib:2
LESSOPEN=||/usr/bin/lesspipe.sh %s
LMOD_SETTARG_FULL_SUPPORT=no
__Init_Default_Modules=1
SLURM_STEP_TASKS_PER_NODE=1
LMOD_FAMILY_COMPILER=gcc
SLURM_STEP_NODELIST=kp362
XDG_RUNTIME_DIR=/run/user/1080628
QT_PLUGIN_PATH=/usr/lib64/kde4/plugins:/usr/lib/kde4
↳ /plugins
OSREL=CentOS
__LMOD_REF_COUNT_MANPATH=/uufs/chpc.utah.edu/sys/ins
↳ talldir/gcc/6.4.0/share/man:1;/uufs/chpc.utah.ed
↳ u/sys/installldir/lmod/7.7.29/share/man:1;/usr/ke
↳ rberos/man:1;/usr/local/share/man:1;/usr/share/m
↳ an/overrides:1;/usr/share/man:1;/uufs/kingspeak.
↳ peaks/sys/pkg/slurm/std/share/man:3
LMOD_DIR=/uufs/chpc.utah.edu/sys/installldir/lmod/7.7
↳ .29/libexec
SLURM_CPU_BIND=quiet,none
INCLUDE=/uufs/chpc.utah.edu/sys/installldir/gcc/6.4.0
↳ /include
BASH_FUNC_module()=() { eval "$($LMOD_CMD bash "$@")
↳ && eval "$($LMOD_SETTARG_CMD:-) -s sh)
}
BASH_FUNC_ml()=() { eval "$($LMOD_DIR/ml_cmd "$@")
}
+ lsb_release -a
LSB Version: :core-4.1-amd64:core-4.1-noarch:
↳ cxx-4.1-amd64:cxx-4.1-noarch:desktop-4.1-amd64:d
↳ esktop-4.1-noarch:languages-4.1-amd64:languages-
↳ 4.1-noarch:printing-4.1-amd64:printing-4.1-noarch
Distributor ID: CentOS
Description: CentOS Linux release 7.6.1810
↳ (Core)
Release: 7.6.1810
Codename: Core
+ uname -a

```

```

Linux kp362 3.10.0-957.1.3.el7.x86_64 #1 SMP Thu Nov
↳ 29 14:49:43 UTC 2018 x86_64 x86_64 x86_64
↳ GNU/Linux
+ lscpu
Architecture: x86_64
CPU op-mode(s): 32-bit, 64-bit
Byte Order: Little Endian
CPU(s): 56
On-line CPU(s) list: 0-55
Thread(s) per core: 2
Core(s) per socket: 14
Socket(s): 2
NUMA node(s): 2
Vendor ID: GenuineIntel
CPU family: 6
Model: 79
Model name: Intel(R) Xeon(R) CPU E5-2680 v4
↳ @ 2.40GHz
Stepping: 1
CPU MHz: 3277.001
CPU max MHz: 3300.0000
CPU min MHz: 1200.0000
BogoMIPS: 4799.66
Virtualization: VT-x
L1d cache: 32K
L1i cache: 32K
L2 cache: 256K
L3 cache: 35840K
NUMA node0 CPU(s): 0,2,4,6,8,10,12,14,16,18,20,22
↳ ,24,26,28,30,32,34,36,38,40,42,44,46,48,50,52,54
NUMA node1 CPU(s): 1,3,5,7,9,11,13,15,17,19,21,23
↳ ,25,27,29,31,33,35,37,39,41,43,45,47,49,51,53,55
Flags: fpu vme de pse tsc msr pae mce
↳ cx8 apic sep mtrr pge mca cmov pat pse36 clflush
↳ dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx
↳ pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs
↳ bts rep_good nopl xtopology nonstop_tsc aperfmperf
↳ eagerfpu pni pclmulqdq dtes64 monitor ds_cpl vmx
↳ smx est tm2 ssse3 sdbg fma cx16 xtpr pdcm pcid dca
↳ sse4_1 sse4_2 x2apic movbe popcnt aes xsave avx
↳ f16c rdrand lahf_lm abm 3dnowprefetch epb cat_l3
↳ cdp_l3 intel_pt tpr_shadow vnmi flexpriority ept
↳ vpid fsgsbase tsc_adjust bmi1 hle avx2 smep bmi2
↳ erms invpcid rtm cqm rdt_a rdseed adx smap
↳ xsaveopt cqm_llc cqm_occup_llc cqm_mbm_total
↳ cqm_mbm_local dtherm ida arat pln pts
+ cat /proc/meminfo
MemTotal: 264033424 kB
MemFree: 241238380 kB
MemAvailable: 239329644 kB
Buffers: 0 kB
Cached: 19087192 kB
SwapCached: 0 kB
Active: 978212 kB
Inactive: 19121056 kB
Active(anon): 900108 kB

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```
Inactive(anon): 19025540 kB
Active(file): 78104 kB
Inactive(file): 95516 kB
Unevictable: 0 kB
Mlocked: 0 kB
SwapTotal: 67108860 kB
SwapFree: 67108860 kB
Dirty: 1508 kB
Writeback: 0 kB
AnonPages: 1012392 kB
Mapped: 231452 kB
Shmem: 18913568 kB
Slab: 530476 kB
SReclaimable: 129912 kB
SUnreclaim: 400564 kB
KernelStack: 14976 kB
PageTables: 16092 kB
NFS_Unstable: 1600 kB
Bounce: 0 kB
WritebackTmp: 0 kB
CommitLimit: 199125572 kB
Committed_AS: 20403576 kB
VmallocTotal: 34359738367 kB
VmallocUsed: 907760 kB
VmallocChunk: 34224568316 kB
HardwareCorrupted: 0 kB
AnonHugePages: 4096 kB
CmaTotal: 0 kB
CmaFree: 0 kB
HugePages_Total: 0
HugePages_Free: 0
HugePages_Rsvd: 0
HugePages_Surp: 0
Hugepagesize: 2048 kB
DirectMap4k: 9759568 kB
DirectMap2M: 254386176 kB
DirectMap1G: 6291456 kB
+ inxi -F -c0
./collect_environment.sh: line 14: inxi: command not
found
+ lsblk -a
NAME MAJ:MIN RM SIZE RO TYPE
└─ sda 8:0 0 1.8T 0 disk
   └─ sda1 8:1 0 1.8T 0 part
      └─ vg_peaks-lv_swap 253:0 0 64G 0 lvm [SWAP]
         └─ vg_peaks-lv_scratch 253:1 0 1.8T 0 lvm
            └─ /scratch/local
+ lsscsi -s
[0:0:0:0] disk ATA TOSHIBA MG04ACA2 FJ2D
└─ /dev/sda 2.00TB
+ module list
++ /uufs/chpc.utah.edu/sys/installdir/lmod/7.7.29/li
└─ bexec/lmod bash
└─ list
```

Currently Loaded Modules:
1) chpc/1.0 2) cuda/9.2 (g) 3) gcc/6.4.0

Where:
g: built for GPU

```
+ eval 'MODULEPATH="/uufs/chpc.utah.edu/sys/modulefi
└─ les/CHPC-18/Compiler/kp/gcc/6.4.0:/uufs/chpc.uta
└─ h.edu/sys/modulefiles/CHPC-18/Compiler/gcc/6.4.0
└─ :/uufs/chpc.utah.edu/sys/modulefiles/CHPC-18/Com
└─ piler/cuda/9.2:/uufs/chpc.utah.edu/sys/modulefil
└─ es/CHPC-18/Linux:/uufs/chpc.utah.edu/sys/modulef
└─ iles/CHPC-18/Core:/uufs/chpc.utah.edu/sys/instal
└─ ldir/lmod/7.7.29/modulefiles/Core";'
export 'MODULEPATH;' '_ModuleTable001_='X01vZHVszVRh
└─ Ymx1Xz17WyJNVHZlcnNpb24iXT0zLWZlZmVzZWJ1aWxkVGlt
└─ ZSJdPWZhbHNlFsiY19zaG9yZFRpbWUiXT1mYWxzZXsXZXB0
└─ aFQ9e30sZmFtaWx5PXBtIkNvbXBpbGVyI109ImdjYyIsWyJj
└─ dWRhI109ImN1ZGEiLH0sbVQ9e2NocGM9e1siZm4iXT0iL3V1
└─ ZnMvY2hwYy51dGFoLmVkdS9zeXMvbW9kdWx1Zm1sZXMvQ0hQ
└─ Qy0xOC9Db3JlL2NocGMvMS4wLmx1YSIsWyJmdWxsTmFtZSJd
└─ PSJjaHBjLzEuMCIsWyJsb2FkT3JkZXIiXT0xLHByb3BUPXt9
└─ LFsic3RhY2tEZXB0aCJdPTAsWyJzdGF0dXMiXT0iYWNoaXZl
└─ IixbInVzZXJOYW1lI109ImNocGMiLH0sY3VhY2t17WYjmbiJd
└─ PSiVdXVmcY9jaHBjLnV0YWguZWZWR1L3N5cy9tb2R1bGVmaWx1
└─ cy9DSFBDLTE4L0NvcmlH0=";'
export '_ModuleTable001_'; '_ModuleTable002_='Y3VhYS
└─ 85LjIubHVhIixbImZ1bGx0YW1lI109ImN1ZGEvOS4yIixbIm
└─ xvYWRPcmRlciJdPTIscHJvcFQ9e2FyY2g9e1siZ3B1I109MS
└─ x9LH0sWyJzdGFja0RlchRoI109MCxbInN0YXR1cyJdPSJhY3
└─ RpdmlUiLWZlZmVzZWJ1aWxkS9zeXMvbW9kdWx1Zm1sZXMvQ0
└─ tbImZU109Ii91dWZzL2NocGMudXRhaC51ZHUvc3l2L21vZV
└─ VsZWZpbGVzL0NlUEMtMTgvQ29yZS9nY2MvNi40LjAubHVhIi
└─ xbImZ1bGx0YW1lI109ImdjYy82LjJqUmCIIsWyJsb2FkT3JkZX
└─ IiXT0zLHByb3BUPXt9LFsic3RhY2tEZXB0aCJdPTAsWyJzdG
└─ F0dXMiXT0iYWNoaXZlIixbInVzZXJOYW1lI109ImdjYy82Lj
└─ QUMCIIsfSx9LGI1wYXR0QT17Ii91dWZzL2NocGMudXRhaC51Z
└─ Uvc3l2L21vZHVszZWZpbGVzL0NlU=";'
export '_ModuleTable002_'; '_ModuleTable003_='UEMTMT
└─ gvQ29tcGlsZXIva3AvZ2NjLzYuNC4wIiwiL3V1ZnMvY2hwYy
└─ 51dGFoLmVkdS9zeXMvbW9kdWx1Zm1sZXMvQ0hQYy0xOC9Db2
└─ 1waWx1ci9nY2MvNi40LjAilCIvdXVmcY9jaHBjLnV0YWguZW
└─ R1L3N5cy9tb2R1bGVmaWx1cy9DSFBDLTE4L0NvbXBpbGVyL2
└─ N1ZGEvOS4yIiwiL3V1ZnMvY2hwYy51dGFoLmVkdS9zeXMvbW
└─ 9kdWx1Zm1sZXMvQ0hQYy0xOC9MaW51eCIIsIi91dWZzL2NocG
└─ MudXRhaC51ZHUvc3l2L21vZHVszZWZpbGVzL0NlUEMtMTgvQ2
└─ 9yZSIIsIi91dWZzL2NocGMudXRhaC51ZHUvc3l2L2luc3Rhbg
└─ xkaXIvbG1vZC83LjcuMjkvbW9kdWx1Zm1sZXMvQ29yZSIIsfS
└─ xbInN5c3R1bUJhc2VNUeFUSCJdPSiVdXVmcY9jaHBjLnV0YW
└─ guZWZWR1L3N5cy9tb2R1bGVmaWx1";'
export '_ModuleTable003_'; '_ModuleTable004_='cy9DSF
└─ BDLTE4L0xpbnV40i91dWZzL2NocGMudXRhaC51ZHUvc3l2L2
└─ 1vZHVszZWZpbGVzL0NlUEMtMTgvQ29yZTovdXVmcY9jaHBjLn
└─ V0YWguZWZWR1L3N5cy9pbmN0YXsZGlyL2xtb2QvNy43LjI5L2
└─ 1vZHVszZWZpbGVzL0NvcmlH0=";'
```

```

export '_ModuleTable004_'; '_ModuleTable_Sz_='4";'
export '_ModuleTable_Sz_';
++ MODULEPATH=/uufs/chpc.utah.edu/sys/modulefiles/CH_
↳ PC-18/Compiler/kp/gcc/6.4.0:/uufs/chpc.utah.edu/_
↳ sys/modulefiles/CHPC-18/Compiler/gcc/6.4.0:/uufs_
↳ /chpc.utah.edu/sys/modulefiles/CHPC-18/Compiler/_
↳ cuda/9.2:/uufs/chpc.utah.edu/sys/modulefiles/CHP_
↳ C-18/Linux:/uufs/chpc.utah.edu/sys/modulefiles/C_
↳ HPC-18/Core:/uufs/chpc.utah.edu/sys/installdir/l_
↳ mod/7.7.29/modulefiles/Core
++
export MODULEPATH
++ _ModuleTable001_=X01vZHVszVRhYmxlXz17WYJNVHZlcnNp_
↳ b24iXT0zLFsiY19yZWJ1aWxkVGltZSJdPWZhbHNlFsiY19z_
↳ aG9ydFRpbWUiXT1mYWxzZSxkZXB0aFQ9e30sZmFtaWx5PXBt_
↳ IkNvbXBpbGVyI109ImdjYyIsWyJjdWRhI109ImN1ZGEiLH0s_
↳ bVQ9e2NocGM9e1siZm4iXT0iL3V1ZnMvY2hwYy51dGFoLmVk_
↳ dS9zeXMvbW9kdWx1Zm1sZXNvQ0hQQy0xOC9Db3JlL2NocGMv_
↳ MS4wLmx1YSIsWyJmdWxsTmFtZSJdPSJjaHBjLzEuMCIsWyJs_
↳ b2FkT3JkZXIiXT0xLHByb3BUPXt9LFsic3RhY2tEZXB0aCJd_
↳ PTAsWyJzdGF0dXMiXT0iYWN0aXZlIixbInVzZXJOYW11I109_
↳ ImNocGMiLH0sY3VkyYT17WyJmbiJdPSlvdXVmcY9jaHBjLnV0_
↳ YWguZWw1L3N5cy9tb2R1bGVmaWxlcY9DSFBDLTE4L0NvcnV_
++
export _ModuleTable001_
++ _ModuleTable002_=Y3VkYS85LjIubHVhIixbImZ1bGx0YW11_
↳ I109ImN1ZGEvOS4yIixbImxvYWRPcmRlciJdPTIscHJvcFQ9_
↳ e2FyY2g9e1siZ3B1I109MSx9LH0sWyJzdGFja0RlcHRoI109_
↳ MCxbInN0YXR1cyJdPSJhY3RpdmlUiLlFsidXNlck5hbWUiXT0i_
↳ Y3VkYS85LjIiLH0sZ2NjPXBtImZuI109Ii91dWZlL2NocGMu_
↳ dXRhaC5lZHUvc3lzl21vZHVszZWZpbGVzL0NIUEMtMTgv29y_
↳ ZS9nY2MvNi40LjAubHVhIixbImZ1bGx0YW11I109ImdjYy82_
↳ LjQuMCIsWyJsb2FkT3JkZXIiXT0xLHByb3BUPXt9LFsic3Rh_
↳ Y2tEZXB0aCJdPTAsWyJzdGF0dXMiXT0iYWN0aXZlIixbInVz_
↳ ZXJOYW11I109ImdjYy82LjQuMCIsfSx9LG1wYXR0QT17Ii91_
↳ dWZlL2NocGMudXRhaC5lZHUvc3lzl21vZHVszZWZpbGVzL0NI_
++
export _ModuleTable002_
++ _ModuleTable003_=UEMTMTgvQ29tcGlsZXIva3AvZ2NjLzYu_
↳ NC4wIiw1L3V1ZnMvY2hwYy51dGFoLmVkdS9zeXMvbW9kdWx1_
↳ Zm1sZXNvQ0hQQy0xOC9Db21waWxlcj9nY2MvNi40LjAilCIv_
↳ dXVmcY9jaHBjLnV0YWguZWw1L3N5cy9tb2R1bGVmaWxlcY9D_
↳ SFBDLTE4L0NvbXBpbGVyL2N1ZGEvOS4yIiw1L3V1ZnMvY2hw_
↳ Yy51dGFoLmVkdS9zeXMvbW9kdWx1Zm1sZXNvQ0hQQy0xOC9M_
↳ aW51eCIsIi91dWZlL2NocGMudXRhaC5lZHUvc3lzl21vZHVz_
↳ ZWZpbGVzL0NIUEMTMTgvQ29yZSIsIi91dWZlL2NocGMudXRh_
↳ aC5lZHUvc3lzl2luc3RhbkxkaXlvdGV1vZC83LjcuMjYvY2h_
↳ dWx1Zm1sZXNvQ29yZSIsfSxbInN5c3R1bUJhc2VNUeFUSCJd_
↳ PSIvdXVmcY9jaHBjLnV0YWguZWw1L3N5cy9tb2R1bGVmaWx1_
++
export _ModuleTable003_
++ _ModuleTable004_=cy9DSFBDLTE4L0xpbnV4oi91dWZlL2No_
↳ cGMudXRhaC5lZHUvc3lzl21vZHVszZWZpbGVzL0NIUEMTMTgv_
↳ Q29yZTovdXVmcY9jaHBjLnV0YWguZWw1L3N5cy9pbN0YWs_
↳ ZGlyL2xtb2QvNy43LjI5L21vZHVszZWZpbGVzL0NvcnVlLH0=
++

```

```

export _ModuleTable004_
++ _ModuleTable_Sz_=4
++
export _ModuleTable_Sz_
++ : -s sh
+ eval
+ nvidia-smi
Wed Apr 10 14:40:46 2019
+-----+
↳ -----+
| NVIDIA-SMI 410.79      Driver Version: 410.79
↳  CUDA Version: 10.0      |
|-----+-----+
↳ ---+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A
↳ | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage
↳ | GPU-Util  Compute M. |
|=====+=====+
↳ ===+=====+
|  0  Tesla P100-PCIE... On      | 00000000:04:00.0 Off
↳ |
↳ | N/A    51C    P0     109W / 250W |      527MiB / 16280MiB
↳ |      94%      Default |
+-----+-----+
↳ ---+-----+
|  1  Tesla P100-PCIE... On      | 00000000:82:00.0 Off
↳ |
↳ | N/A    35C    P0     28W / 250W |      0MiB / 16280MiB
↳ |      0%      Default |
+-----+-----+
↳ ---+-----+
+-----+
↳ -----+
| Processes:
↳
↳      GPU Memory |
| GPU      PID   Type   Process name
↳
↳      Usage      |
|=====+=====+
↳ =====+
|  0      125071    C
↳ ...mber/GIT/amber-notchpeak/bin/pmemd.cuda
↳ 517MiB |
+-----+
↳ -----+
+ lshw -short -quiet -sanitize
+ cat
WARNING: you should run this program as super-user.
H/W path          Device Class          Description
=====+=====+
↳ ===
                                system      Computer
/0                bus          Motherboard

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

/0/0	memory	256GiB	/0/100/1c.7	bridge	C610/X99
↪ System memory			↪ series chipset PCI Express Root Port #8		
/0/4	processor	Intel(R)	/0/100/1c.7/0	bridge	SH7758 PCIe
↪ Xeon(R) CPU E5-2680 v4 @ 2.40GHz			↪ Switch [PS]		
/0/6	processor	Intel(R)	/0/100/1c.7/0/0	bridge	SH7758
↪ Xeon(R) CPU E5-2680 v4 @ 2.40GHz			↪ PCIe Switch [PS]		
/0/100	bridge	Xeon E7	/0/100/1c.7/0/0/0	bridge	SH7758
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DMI2			↪ PCIe-PCI Bridge [PPB]		
/0/100/1	bridge	Xeon E7	/0/100/1c.7/0/0/0/0	display	G200eR2
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			/0/100/1d	bus	C610/X99
↪ Port 1			↪ series chipset USB Enhanced Host Controller #1		
/0/100/2	bridge	Xeon E7	/0/100/1f	bridge	C610/X99
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			↪ series chipset LPC Controller		
↪ Port 2			/0/100/1f.2	storage	C610/X99
/0/100/2/0	display	GP100GL	↪ series chipset 6-Port SATA Controller [AHCI mode]		
↪ [Tesla P100 PCIe 16GB]			/0/7	generic	Xeon E7
/0/100/3	bridge	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 0		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			/0/8	generic	Xeon E7
↪ Port 3			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 0		
/0/100/3/0	eth2 network	NetXtreme	/0/9	generic	Xeon E7
↪ BCM5720 Gigabit Ethernet PCIe			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 0		
/0/100/3/0.1	eth3 network	NetXtreme	/0/a	generic	Xeon E7
↪ BCM5720 Gigabit Ethernet PCIe			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 1		
/0/100/3.1	bridge	Xeon E7	/0/b	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 1		
↪ Port 3			/0/c	generic	Xeon E7
/0/100/3.1/0	eth0 network	NetXtreme	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 1		
↪ BCM5720 Gigabit Ethernet PCIe			/0/d	generic	Xeon E7
/0/100/3.1/0.1	eth1 network	NetXtreme	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link 0/1		
↪ BCM5720 Gigabit Ethernet PCIe			/0/e	generic	Xeon E7
/0/100/3.2	bridge	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link 0/1		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			/0/f	generic	Xeon E7
↪ Port 3			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link 0/1		
/0/100/5	generic	Xeon E7	/0/10	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link Debug		
↪ Map/VTd_Misc/System Management			/0/11	generic	Xeon E7
/0/100/5.1	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D IIO Hot Plug			/0/12	generic	Xeon E7
/0/100/5.2	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D IIO RAS/Control			/0/13	generic	Xeon E7
↪ Status/Global Errors			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/100/5.4	generic	Xeon E7	/0/14	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D I/O APIC			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/100/11	generic	C610/X99	/0/15	generic	Xeon E7
↪ series chipset SPSR			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/100/11.4	storage	C610/X99	/0/16	generic	Xeon E7
↪ series chipset sSATA Controller [AHCI mode]			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/100/16	communication	C610/X99	/0/17	generic	Xeon E7
↪ series chipset MEI Controller #1			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/100/16.1	communication	C610/X99	/0/18	generic	Xeon E7
↪ series chipset MEI Controller #2			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/100/1a	bus	C610/X99	/0/19	generic	Xeon E7
↪ series chipset USB Enhanced Host Controller #2			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/100/1c	bridge	C610/X99			
↪ series chipset PCI Express Root Port #1					

/0/1a	generic	Xeon E7	/0/33	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
/0/1b	generic	Xeon E7	↪ 0 - Channel Target Address Decoder		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			/0/34	generic	Xeon E7
/0/1c	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ 0 - Channel Target Address Decoder		
/0/1d	generic	Xeon E7	/0/35	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
/0/1e	generic	Xeon E7	↪ Broadcast		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			/0/36	generic	Xeon E7
/0/1f	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Global		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ Broadcast		
/0/20	generic	Xeon E7	/0/37	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
/0/21	generic	Xeon E7	↪ 0 - Channel 0 Thermal Control		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			/0/38	generic	Xeon E7
/0/22	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ 0 - Channel 1 Thermal Control		
/0/23	generic	Xeon E7	/0/39	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
/0/24	generic	Xeon E7	↪ 0 - Channel 0 Error		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			/0/3a	generic	Xeon E7
/0/25	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ 0 - Channel 1 Error		
/0/26	generic	Xeon E7	/0/3b	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R2PCIe Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
/0/27	generic	Xeon E7	↪ Interface		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R2PCIe Agent			/0/3c	generic	Xeon E7
/0/28	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Ubox			↪ Interface		
/0/29	generic	Xeon E7	/0/3d	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Ubox			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
/0/2a	generic	Xeon E7	↪ Interface		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Ubox			/0/3e	generic	Xeon E7
/0/2b	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 0			↪ Interface		
/0/2c	generic	Xeon E7	/0/3f	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 0			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Target		
/0/2d	generic	Xeon E7	↪ Address/Thermal/RAS		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 0 Debug			/0/40	generic	Xeon E7
/0/2e	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Target		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 1			↪ Address/Thermal/RAS		
/0/2f	generic	Xeon E7	/0/41	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 1			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Channel Target		
/0/30	generic	Xeon E7	↪ Address Decoder		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 1 Debug			/0/42	generic	Xeon E7
/0/31	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Channel Target		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↪ Address Decoder		
↪ 0 - Target Address/Thermal/RAS			/0/43	generic	Xeon E7
/0/32	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↪ Broadcast		
↪ 0 - Target Address/Thermal/RAS			/0/44	generic	Xeon E7
			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Global		
			↪ Broadcast		
			/0/45	generic	Xeon E7
			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
			↪ 1 - Channel 0 Thermal Control		

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

/0/46	generic	Xeon E7	/0/5.2	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D IIO RAS/Control		
↳ 1 - Channel 1 Thermal Control			↳ Status/Global Errors		
/0/47	generic	Xeon E7	/0/5.4	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D I/O APIC		
↳ 1 - Channel 0 Error			/0/54	generic	Xeon E7
/0/48	generic	Xeon E7	↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 0		
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			/0/55	generic	Xeon E7
↳ 1 - Channel 1 Error			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 0		
/0/49	generic	Xeon E7	/0/56	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 0		
↳ Interface			/0/57	generic	Xeon E7
/0/4a	generic	Xeon E7	↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 1		
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3			/0/58	generic	Xeon E7
↳ Interface			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 1		
/0/4b	generic	Xeon E7	/0/59	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D QPI Link 1		
↳ Interface			/0/5a	generic	Xeon E7
/0/4c	generic	Xeon E7	↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link 0/1		
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3			/0/5b	generic	Xeon E7
↳ Interface			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link 0/1		
/0/4d	generic	Xeon E7	/0/5c	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link 0/1		
/0/4e	generic	Xeon E7	/0/5d	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R3 QPI Link Debug		
/0/4f	generic	Xeon E7	/0/5e	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/50	generic	Xeon E7	/0/5f	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/51	generic	Xeon E7	/0/60	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/52	generic	Xeon E7	/0/61	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/53	generic	Xeon E7	/0/62	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/1	bridge	Xeon E7	/0/63	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↳ Port 1			/0/64	generic	Xeon E7
/0/2	bridge	Xeon E7	↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			/0/65	generic	Xeon E7
↳ Port 2			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/2/0	display	GP100GL	/0/66	generic	Xeon E7
↳ [Tesla P100 PCIe 16GB]			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/3	bridge	Xeon E7	/0/67	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↳ Port 3			/0/68	generic	Xeon E7
/0/3/0	ib0	network	↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↳ Family [ConnectX-3]			/0/69	generic	Xeon E7
/0/3.2	bridge	Xeon E7	↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D PCI Express Root			/0/6a	generic	Xeon E7
↳ Port 3			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
/0/5	generic	Xeon E7	/0/6b	generic	Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D			↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent		
↳ Map/VTd_Misc/System Management					
/0/5.1	generic	Xeon E7			
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D IIO Hot Plug					

/0/6c	generic	Xeon E7	/0/84	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
/0/6d	generic	Xeon E7	↪ 0 - Channel 0 Thermal Control		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			/0/85	generic	Xeon E7
/0/6e	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ 0 - Channel 1 Thermal Control		
/0/6f	generic	Xeon E7	/0/86	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
/0/70	generic	Xeon E7	↪ 0 - Channel 0 Error		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			/0/87	generic	Xeon E7
/0/71	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ 0 - Channel 1 Error		
/0/72	generic	Xeon E7	/0/88	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Caching Agent			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
/0/73	generic	Xeon E7	↪ Interface		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R2PCIE Agent			/0/89	generic	Xeon E7
/0/74	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D R2PCIE Agent			↪ Interface		
/0/75	generic	Xeon E7	/0/8a	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Ubox			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
/0/76	generic	Xeon E7	↪ Interface		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Ubox			/0/8b	generic	Xeon E7
/0/77	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Ubox			↪ Interface		
/0/78	generic	Xeon E7	/0/8c	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 0			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Target		
/0/79	generic	Xeon E7	↪ Address/Thermal/RAS		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 0			/0/8d	generic	Xeon E7
/0/7a	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Target		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 0 Debug			↪ Address/Thermal/RAS		
/0/7b	generic	Xeon E7	/0/8e	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 1			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Channel Target		
/0/7c	generic	Xeon E7	↪ Address Decoder		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 1			/0/8f	generic	Xeon E7
/0/7d	generic	Xeon E7	↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Channel Target		
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Home Agent 1 Debug			↪ Address Decoder		
/0/7e	generic	Xeon E7	/0/90	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3		
↪ 0 - Target Address/Thermal/RAS			↪ Broadcast		
/0/7f	generic	Xeon E7	/0/91	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Global		
↪ 0 - Target Address/Thermal/RAS			↪ Broadcast		
/0/80	generic	Xeon E7	/0/92	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ 0 - Channel Target Address Decoder			↪ 1 - Channel 0 Thermal Control		
/0/81	generic	Xeon E7	/0/93	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ 0 - Channel Target Address Decoder			↪ 1 - Channel 1 Thermal Control		
/0/82	generic	Xeon E7	/0/94	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 0/1			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ Broadcast			↪ 1 - Channel 0 Error		
/0/83	generic	Xeon E7	/0/95	generic	Xeon E7
↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Global			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Memory Controller		
↪ Broadcast			↪ 1 - Channel 1 Error		
			/0/96	generic	Xeon E7
			↪ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3		
			↪ Interface		

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```

/0/97          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3
↳ Interface
/0/98          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3
↳ Interface
/0/99          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D DDRIO Channel 2/3
↳ Interface
/0/9a          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit
/0/9b          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit
/0/9c          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit
/0/9d          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit
/0/9e          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit
/0/9f          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit
/0/a0          generic      Xeon E7
↳ v4/Xeon E5 v4/Xeon E3 v4/Xeon D Power Control Unit
/0/a1          system       PnP device
↳ PNP0b00
/0/a2          system       PnP device
↳ PNP0c02
/0/a3          communication PnP device
↳ PNP0501
/0/a4          communication PnP device
↳ PNP0501
/0/a5          generic      PnP device
↳ IPI0001

```

WARNING: output may be incomplete or inaccurate, you
↳ should run this program as super-user.

```

=====
↳ =====

```

Skylake-X node:

```

SLURM_NODELIST=notch035
SLURM_CHECKPOINT_IMAGE_DIR=/var/slurm/checkpoint
LMOD_FAMILY_COMPILER_VERSION=2019.2.187
MKLR00T=/uufs/chpc.utah.edu/sys/installdir/intel/com
↳ pilers_and_libraries_2019.2.187/linux/mkl
MANPATH=/uufs/chpc.utah.edu/sys/installdir/intel/doc
↳ umentation_2019/en/man/common:/uufs/chpc.utah.ed
↳ u/sys/installdir/intel/compilers_and_libraries_2
↳ 019.2.187/linux/mpi/man:/uufs/chpc.utah.edu/sys/
↳ installdir/intel/documentation_2019/en/debugger/
↳ gdb-ia/man:/uufs/chpc.utah.edu/sys/installdir/lm
↳ od/7.7.29/share/man:/usr/kerberos/man:/usr/local
↳ /share/man:/usr/share/man/overrides:/usr/share/m
↳ an:/uufs/notchpeak.peaks/sys/installdir/slurm/st
↳ d/share/man
SLURM_JOB_NAME=hello

```

```

XDG_SESSION_ID=127512
HOSTNAME=notch035
SLURM_TOPOLOGY_ADDR=notch035
SLURMD_NODENAME=notch035
_ModuleTable003_=cmUiLH0sWyJzeXN0ZW1CYXNlTVBBVEgiXT0
↳ iL3V1ZnMvY2hwYy51dGFoLmVkdS9zeXMvbW9kdWx1ZmlsZXM
↳ vQ0hQQy0xOC9MaW51eDovdXVmc9jaHBjLnV0YWguZWZWR1L3N
↳ 5cy9tb2R1bGVmaWxlc9DSFBDLTE4L0NvcmU6L3V1ZnMvY2h
↳ wYy51dGFoLmVkdS9zeXMvaW5zdGFsbGRpci9sbW9kLzcuNy4
↳ yOS9tb2R1bGVmaWxlc9Db3JlIix9
SLURM_PRIORITY_PROCESS=0
SLURM_NODE_ALIASES=(null)
INTEL_LICENSE_FILE=/uufs/chpc.utah.edu/sys/installdi
↳ r/intel/licenses
IPPROOT=/uufs/chpc.utah.edu/sys/installdir/intel/com
↳ pilers_and_libraries_2019.2.187/linux/ipp
__LMOD_REF_COUNT_MODULEPATH=/uufs/chpc.utah.edu/sys/
↳ modulefiles/CHPC-18/Compiler/intel/2019.2.187:1;
↳ /uufs/chpc.utah.edu/sys/modulefiles/CHPC-18/Linu
↳ x:1;/uufs/chpc.utah.edu/sys/modulefiles/CHPC-18/
↳ Core:1;/uufs/chpc.utah.edu/sys/installdir/lmod/7
↳ .7.29/modulefiles/Core:1
TERM=xterm-256color
SHELL=/bin/bash
HISTSIZE=1000
SLURM_JOB_QOS=notchpeak-freecycle
KMP_HOT_TEAMS_MODE=1
LMOD_SYSTEM_DEFAULT_MODULES=chpc
MODULEPATH_ROOT=/uufs/chpc.utah.edu/sys/modulefiles/
↳ CHPC-18
SSH_CLIENT=155.97.232.235 53200 22
SLURM_TOPOLOGY_ADDR_PATTERN=node
TMPDIR=/tmp
LIBRARY_PATH=/uufs/chpc.utah.edu/sys/installdir/inte
↳ l//compilers_and_libraries_2019.2.187/linux/mpi/
↳ intel64/libfabric/lib:/uufs/chpc.utah.edu/sys/in
↳ stalldir/intel/compilers_and_libraries_2019.2.18
↳ 7/linux/ipp/lib/intel64:/uufs/chpc.utah.edu/sys/
↳ installdir/intel/compilers_and_libraries_2019.2.
↳ 187/linux/compiler/lib/intel64_lin:/uufs/chpc.ut
↳ ah.edu/sys/installdir/intel/compilers_and_librar
↳ ies_2019.2.187/linux/mkl/lib/intel64_lin:/uufs/c
↳ hpc.utah.edu/sys/installdir/intel/compilers_and_
↳ libraries_2019.2.187/linux/tbb/lib/intel64/gcc4.
↳ 7:/uufs/chpc.utah.edu/sys/installdir/intel/comp
↳ lers_and_libraries_2019.2.187/linux/tbb/lib/inte
↳ l64/gcc4.7:/uufs/chpc.utah.edu/sys/installdir/in
↳ tel/compilers_and_libraries_2019.2.187/linux/daa
↳ l/lib/intel64_lin:/uufs/chpc.utah.edu/sys/instal
↳ ldir/intel/compilers_and_libraries_2019.2.187/li
↳ nux/daal/./tbb/lib/intel64_lin/gcc4.4
LMOD_PKG=/uufs/chpc.utah.edu/sys/installdir/lmod/7.7
↳ .29
QTDIR=/usr/lib64/qt-3.3
QTINC=/usr/lib64/qt-3.3/include
LMOD_VERSION=7.7.29

```

```

SSH_TTY=/dev/pts/8
__LMOD_REF_COUNT_LOADEDMODULES=chpc/1.0:1:intel/2019
↳ .2.187:1
QT_GRAPHICSSYSTEM_CHECKED=1
OSVER=7.6.1810
USER=USER
SLURM_NNODES=1
LD_LIBRARY_PATH=/uufs/chpc.utah.edu/sys/installdir/g
↳ cc/8.1.0/lib64:/uufs/chpc.utah.edu/sys/installdi
↳ r/gcc/8.1.0/lib:/uufs/chpc.utah.edu/sys/installd
↳ ir/intel/compilers_and_libraries_2019.2.187/linu
↳ x/compiler/lib/intel64_lin:/uufs/chpc.utah.edu/s
↳ ys/installdir/intel//compilers_and_libraries_201
↳ 9.2.187/linux/mpi/intel64/libfabric/lib:/uufs/ch
↳ pc.utah.edu/sys/installdir/intel//compilers_and_
↳ libraries_2019.2.187/linux/mpi/intel64/lib/relea
↳ se:/uufs/chpc.utah.edu/sys/installdir/intel//com
↳ pilers_and_libraries_2019.2.187/linux/mpi/intel6
↳ 4/lib:/uufs/chpc.utah.edu/sys/installdir/intel/c
↳ ompilers_and_libraries_2019.2.187/linux/ipp/lib/
↳ intel64:/uufs/chpc.utah.edu/sys/installdir/intel
↳ /compilers_and_libraries_2019.2.187/linux/compil
↳ er/lib/intel64_lin:/uufs/chpc.utah.edu/sys/insta
↳ lldir/intel/compilers_and_libraries_2019.2.187/l
↳ inux/mkl/lib/intel64_lin:/uufs/chpc.utah.edu/sys
↳ /installdir/intel/compilers_and_libraries_2019.2
↳ .187/linux/tbb/lib/intel64/gcc4.7:/uufs/chpc.uta
↳ h.edu/sys/installdir/intel/compilers_and_librari
↳ es_2019.2.187/linux/tbb/lib/intel64/gcc4.7:/uufs
↳ /chpc.utah.edu/sys/installdir/intel/debugger_201
↳ 9/libipt/intel64/lib:/uufs/chpc.utah.edu/sys/ins
↳ talldir/intel/compilers_and_libraries_2019.2.187
↳ /linux/daal/lib/intel64_lin:/uufs/chpc.utah.edu/
↳ sys/installdir/intel/compilers_and_libraries_201
↳ 9.2.187/linux/daal/./tbb/lib/intel64_lin/gcc4.4

```

```

LS_COLORS=rs=0:di=38;5;27:ln=38;5;51:mh=44;38;5;15:p
↳ i=40;38;5;11:so=38;5;13:do=38;5;5:bd=48;5;232;38
↳ ;5;11:cd=48;5;232;38;5;3:or=48;5;232;38;5;9:mi=0
↳ 5;48;5;232;38;5;15:su=48;5;196;38;5;15:sg=48;5;1
↳ 1;38;5;16:ca=48;5;196;38;5;226:tw=48;5;10;38;5;1
↳ 6:ow=48;5;10;38;5;21:st=48;5;21;38;5;15:ex=38;5;
↳ 34:*.tar=38;5;9:*.tgz=38;5;9:*.arc=38;5;9:*.arj=
↳ 38;5;9:*.taz=38;5;9:*.lha=38;5;9:*.lz4=38;5;9:*.
↳ lzh=38;5;9:*.lzma=38;5;9:*.tlz=38;5;9:*.txz=38;5
↳ ;9:*.tzo=38;5;9:*.t7z=38;5;9:*.zip=38;5;9:*.z=38
↳ ;5;9:*.Z=38;5;9:*.dz=38;5;9:*.gz=38;5;9:*.lrz=38
↳ ;5;9:*.lz=38;5;9:*.lzo=38;5;9:*.xz=38;5;9:*.bz2=
↳ 38;5;9:*.bz=38;5;9:*.tbz=38;5;9:*.tbz2=38;5;9:*.
↳ tz=38;5;9:*.deb=38;5;9:*.rpm=38;5;9:*.jar=38;5;9
↳ :*.war=38;5;9:*.ear=38;5;9:*.sar=38;5;9:*.rar=38
↳ ;5;9:*.alz=38;5;9:*.ace=38;5;9:*.zoo=38;5;9:*.cp
↳ io=38;5;9:*.7z=38;5;9:*.rz=38;5;9:*.cab=38;5;9:*.
↳ .jpg=38;5;13:*.jpeg=38;5;13:*.gif=38;5;13:*.bmp=
↳ 38;5;13:*.pbm=38;5;13:*.pgm=38;5;13:*.ppm=38;5;1
↳ 3:*.tga=38;5;13:*.xbm=38;5;13:*.xpm=38;5;13:*.ti
↳ f=38;5;13:*.tiff=38;5;13:*.png=38;5;13:*.svg=38;
↳ 5;13:*.svgz=38;5;13:*.mng=38;5;13:*.pcx=38;5;13:
↳ *.mov=38;5;13:*.mpg=38;5;13:*.mpeg=38;5;13:*.m2v
↳ =38;5;13:*.mkv=38;5;13:*.webm=38;5;13:*.ogm=38;5
↳ ;13:*.mp4=38;5;13:*.m4v=38;5;13:*.mp4v=38;5;13:*.
↳ .vob=38;5;13:*.qt=38;5;13:*.nuv=38;5;13:*.wmv=38
↳ ;5;13:*.asf=38;5;13:*.rm=38;5;13:*.rmvb=38;5;13:
↳ *.flc=38;5;13:*.avi=38;5;13:*.fli=38;5;13:*.flv=
↳ 38;5;13:*.gl=38;5;13:*.dl=38;5;13:*.xcf=38;5;13:
↳ *.xwd=38;5;13:*.yuv=38;5;13:*.cgm=38;5;13:*.emf=
↳ 38;5;13:*.axv=38;5;13:*.anx=38;5;13:*.ogv=38;5;1
↳ 3:*.ogx=38;5;13:*.aac=38;5;45:*.au=38;5;45:*.fla
↳ c=38;5;45:*.mid=38;5;45:*.midi=38;5;45:*.mka=38;
↳ 5;45:*.mp3=38;5;45:*.mpc=38;5;45:*.ogg=38;5;45:*.
↳ .ra=38;5;45:*.wav=38;5;45:*.axa=38;5;45:*.oga=38
↳ ;5;45:*.spx=38;5;45:*.xspf=38;5;45:
LMOD_sys=Linux
UUFSCELL=notchpeak.peaks
PSTLR00T=/uufs/chpc.utah.edu/sys/installdir/intel/co
↳ mpilers_and_libraries_2019.2.187/linux/pstl
FI_PROVIDER_PATH=/uufs/chpc.utah.edu/sys/installdir/
↳ intel//compilers_and_libraries_2019.2.187/linux/
↳ mpi/intel64/libfabric/lib/prov
SSH_AUTH_SOCKET=/tmp/ssh-ElHXvX90p0/agent.252968
SLURM_JOBID=289502

```

```

CPATH=/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/ipp/include:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/mkl/include:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/pstl/include:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/tbb/include:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/tbb/include:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/daal/include
TMOUT=0
__LMOD_REF_COUNT__LMFILES=/uufs/chpc.utah.edu/sys/modulefiles/CHPC-18/Core/chpc/1.0.lua:1:/uufs/chpc.utah.edu/sys/modulefiles/CHPC-18/Core/intel/2019.2.187.lua:1
KMP_HOT_TEAMS_MAX_LEVEL=2
SLURM_NTASKS=1
VIRTUAL_ENV=/uufs/chpc.utah.edu/common/home/USER/brickv2/env
NLSPATH=/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/compiler/lib/intel64/locale/%l_1t/%N:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/mkl/lib/intel64_lin/locale/%l_1t/%N:/uufs/chpc.utah.edu/sys/installdir/intel/debugger_2019/gdb/intel64/share/locale/%l_1t/%N
ModuleTable001=X0i1vZHVsvZVRHymXlXz17WyJNVHZlcnNpb24iX0ZLFSiY19yZWJ1aWxkVGltZSJsZDpWZhbHNlFSiY19zaG9yYdFRpbWUjX0t1mYwXzZSxkZXB0aG9Qe30sZmFtaWx5PXBtIknvYbXBPbGVYI09ImIudGVsIix9LGIUPXtjaHBjPXBtImZuI09yIi91dWZlZ2NocGMudXRhaC5lZHUvc3ZlZ21vZHVsvZWZpbGVyZl09IUUeMTMTgvQ29yZS9jaHBjLzEuMC5sdWEiLFSiZnVsbE5hbWUjX0t0iY2hwYy8xLjAiLFSibG9hZE9yZGVyI09MSXwcm9wVD17fSxbInN0YWNnRGVwdGgiX0t0wLFSic3RhdHVzI09ImFjdG1Z2ZSIswYj1c2VyTmFtZSJsZDpSjjaHBjIix9LGIudGVsPXBtImZuI09yIi91dWZlZ2NocGMudXRhaC5lZHUvc3ZlZ21vZHVsvZWZpbGVyZl09IUUeMTMTgvQ29yZS9pbmRlbC8yMDE5LjIu
MAIL=/var/spool/mail/USER
PATH=/uufs/chpc.utah.edu/sys/installdir/gcc/8.1.0/bin:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/bin/intel64:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/mpi/intel64/libfabric/bin:/uufs/chpc.utah.edu/sys/installdir/intel/compilers_and_libraries_2019.2.187/linux/mpi/intel64/bin:/uufs/chpc.utah.edu/sys/installdir/intel/debugger_2019/gdb/intel64/bin:/uufs/chpc.utah.edu/common/home/USER/brickv2/env/bin:/usr/lib64/qt-3.3/bin:/usr/kerberos/sbin:/usr/kerberos/bin:/usr/lib64/ccache:/usr/local/bin:/usr/bin:/usr/local/sbin:/usr/sbin:/uufs/notchpeak.peak/sys/installdir/slurm/std/bin:/opt/dell/srvadmin/bin:/uufs/chpc.utah.edu/common/home/USER/bin
SLURM_TASKS_PER_NODE=1

```

```
SLURM_WORKING_CLUSTER=notchpeak:10.242.67.21:6817:81
_ 92
_=/usr/bin/env
SLURM_JOB_ID=289502
TBBROOT=/uufs/chpc.utah.edu/sys/installdir/intel/com
_ pilers_and_libraries_2019.2.187/linux/tbb
PWD=/uufs/chpc.utah.edu/common/home/USER/brickv2/Aut
_ hor-Kit
SLURM_JOB_USER=USER
_LMFILES_=/uufs/chpc.utah.edu/sys/modulefiles/CHPC-1
_ 8/Core/chpc/1.0.lua:/uufs/chpc.utah.edu/sys/modu
_ lefiles/CHPC-18/Core/intel/2019.2.187.lua
LANG=en_US.UTF-8
MODULEPATH=/uufs/chpc.utah.edu/sys/modulefiles/CHPC-
_ 18/Compiler/intel/2019.2.187:/uufs/chpc.utah.edu
_ /sys/modulefiles/CHPC-18/Linux:/uufs/chpc.utah.e
_ du/sys/modulefiles/CHPC-18/Core:/uufs/chpc.utah.
_ edu/sys/installdir/lmod/7.7.29/modulefiles/Core
KDEDIRS=/usr
_ModuleTable_Sz=3
LOADEDMODULES=chpc/1.0:intel/2019.2.187
SLURM_JOB_UID=1080628
SLURM_NODEID=0
SLURM_SUBMIT_DIR=/uufs/chpc.utah.edu/common/home/USE
_ R/brickv2/Author-Kit
LMOD_CMD=/uufs/chpc.utah.edu/sys/installdir/lmod/7.7
_ .29/libexec/lmod
SLURM_NPROCS=1
SLURM_TASK_PID=385164
SLURM_CPUS_ON_NODE=32
DAALROOT=/uufs/chpc.utah.edu/sys/installdir/intel/co
_ mpilers_and_libraries_2019.2.187/linux/daal
SSH_ASKPASS=/usr/libexec/openssh/gnome-ssh-askpass
HISTCONTROL=ignoredups
KRB5CCNAME=KEYRING:persistent:1080628
ENVIRONMENT=BATCH
SLURM_PROCID=0
SLURM_JOB_NODELIST=notch035
INTEL_PYTHONHOME=/uufs/chpc.utah.edu/sys/installdir/
_ intel/debugger_2019/python/intel64/
SHLVL=3
HOME=/uufs/chpc.utah.edu/common/home/USER
SLURM_LOCALID=0
```



```

__LMOD_REF_COUNT_PATH=/uufs/chpc.utah.edu/sys/install
↳ ldir/intel/compilers_and_libraries_2019.2.187/li
↳ nux/bin/intel64:1;/uufs/chpc.utah.edu/sys/install
↳ ldir/intel/compilers_and_libraries_2019.2.187/li
↳ nux/mpi/intel64/libfabric/bin:1;/uufs/chpc.utah.
↳ edu/sys/installdir/intel/compilers_and_libraries
↳ _2019.2.187/linux/mpi/intel64/bin:1;/uufs/chpc.u
↳ tah.edu/sys/installdir/intel/debugger_2019/gdb/i
↳ intel64/bin:1;/uufs/chpc.utah.edu/common/home/USE
↳ R/brickv2/env/bin:1;/usr/lib64/qt-3.3/bin:1;/usr
↳ /kerberos/sbin:1;/usr/kerberos/bin:1;/usr/lib64/
↳ ccache:1;/usr/local/bin:1;/usr/bin:1;/usr/local/
↳ sbin:1;/usr/sbin:1;/uufs/notchpeak.peaks/sys/ins
↳ talldir/slurm/std/bin:1;/opt/dell/srvadmin/bin:1
↳ ;/uufs/chpc.utah.edu/common/home/USER/bin:1
_ModuleTable002_=MTg3Lmx1YSIsWyJmdWxsTmFtZSJdPSJpbmR
↳ lbC8yMDE5LjIuMTg3IixbImxvYWRPcmRlciJdPTIsCHJvcFQ
↳ 9e30sWyJzdGFja0RlCHRoI09MCxbInN0YXR1cyJdPSJhY3R
↳ pdmUiLFsidXNlck5hbWUiXT0iaW50ZWwvMjAxOS4yLjE4NyI
↳ sfSx9LG1wYXR0QT17Ii9ldWZzL2NocGMudXRhaC5lZHUvc3l
↳ zL21vZHVzZWZpbGVzL0NIUEMtMTgvQ29tcGlsZXIvaW50ZWw
↳ vMjAxOS4yLjE4NyIsIi9ldWZzL2NocGMudXRhaC5lZHUvc3l
↳ zL21vZHVzZWZpbGVzL0NIUEMtMTgvTGluXG9ldXVmc3l
↳ jaHBjLnV0YWguZWZlL3N5cy9tb2R1bGVmaWxlc3l5dSFDL
↳ 4L0NvcUilCIdXVmc3l5jaHBjLnV0YWguZWZlL3N5cy9pbmN
↳ 0YWxsZGlyL2xtb2QvNy43LjI5L21vZHVzZWZpbGVzL0Nv
SLURM_CLUSTER_NAME=notchpeak
SLURM_JOB_CPUS_PER_NODE=32
SLURM_JOB_GID=2030
SLURM_SUBMIT_HOST=notchpeak1
SLURM_GTIDS=0
BASH_ENV=/uufs/chpc.utah.edu/sys/installdir/lmod/7.7
↳ .29/init/bash
SLURM_JOB_PARTITION=notchpeak-freecycle
MODULERCFILE=/uufs/chpc.utah.edu/sys/modulefiles/etc
↳ /rc
LOGNAME=USER
QTLIB=/usr/lib64/qt-3.3/lib
CVS_RSH=ssh
BASHRC_LOADED=1
XDG_DATA_DIRS=/uufs/chpc.utah.edu/common/home/USER/.
↳ local/share/flatpak/exports/share:/var/lib/flatp
↳ ak/exports/share:/usr/local/share:/usr/share
SSH_CONNECTION=155.97.232.235 53200 155.101.26.78 22
SLURM_JOB_ACCOUNT=ACCT
CLASSPATH=/uufs/chpc.utah.edu/sys/installdir/intel/
↳ compilers_and_libraries_2019.2.187/linux/mpi/int
↳ el64/lib/mpi.jar:/uufs/chpc.utah.edu/sys/install
↳ dir/intel/compilers_and_libraries_2019.2.187/lin
↳ ux/daal/lib/daal.jar
KMP_HW_SUBSET=1s
MODULESHOME=/uufs/chpc.utah.edu/sys/installdir/lmod/
↳ 7.7.29
SLURM_JOB_NUM_NODES=1
LESSOPEN=||/usr/bin/lesspipe.sh %s
LMOD_SETTARG_FULL_SUPPORT=no

```

```

PKG_CONFIG_PATH=/uufs/chpc.utah.edu/sys/installdir/i
↳ intel/compilers_and_libraries_2019.2.187/linux/mk
↳ l/bin/pkgconfig
OMP_PLACES=cores
__Init_Default_Modules=1
INTELR00T=/uufs/chpc.utah.edu/sys/installdir/intel/c
↳ ompilers_and_libraries_2019.2.187/linux
arch=intel64
INFOPATH=/uufs/chpc.utah.edu/sys/installdir/intel/do
↳ cumentation_2019/en/debugger/gdb-ia/info/
LMOD_FAMILY_COMPILER=intel
XDG_RUNTIME_DIR=/run/user/1080628
QT_PLUGIN_PATH=/usr/lib64/kde4/plugins:/usr/lib/kde4
↳ /plugins
OSREL=CentOS
LMOD_DIR=/uufs/chpc.utah.edu/sys/installdir/lmod/7.7
↳ .29/libexec
__LMOD_REF_COUNT_MANPATH=/uufs/chpc.utah.edu/sys/ins
↳ talldir/intel/documentation_2019/en/man/common:1
↳ ;/uufs/chpc.utah.edu/sys/installdir/intel/compil
↳ ers_and_libraries_2019.2.187/linux/mpi/man:1;/uu
↳ fs/chpc.utah.edu/sys/installdir/intel/documentat
↳ ion_2019/en/debugger/gdb-ia/man:1;/uufs/chpc.uta
↳ h.edu/sys/installdir/lmod/7.7.29/share/man:1;/us
↳ r/kerberos/man:1;/usr/local/share/man:1;/usr/sha
↳ re/man/overrides:1;/usr/share/man:1;/uufs/notchp
↳ eak.peaks/sys/installdir/slurm/std/share/man:2
SLURM_MEM_PER_NODE=102400
I_MPI_ROOT=/uufs/chpc.utah.edu/sys/installdir/intel/
↳ /compilers_and_libraries_2019.2.187/linux/mpi
BASH_FUNC_module()=( ) { eval $(($LMOD_CMD bash "$@")
↳ && eval $((${LMOD_SETTARG_CMD:-}) -s sh)
}
BASH_FUNC_ml()=( ) { eval $(($LMOD_DIR/ml_cmd "$@")
}
+ lsb_release -a
LSB Version: :core-4.1-amd64:core-4.1-noarch:
↳ cxx-4.1-amd64:cxx-4.1-noarch:desktop-4.1-amd64:d
↳ esktop-4.1-noarch:languages-4.1-amd64:languages-
↳ 4.1-noarch:printing-4.1-amd64:printing-4.1-noarch
Distributor ID: CentOS
Description: CentOS Linux release 7.6.1810
↳ (Core)
Release: 7.6.1810
Codename: Core
+ uname -a
Linux notch035 3.10.0-957.1.3.el7.x86_64 #1 SMP Thu
↳ Nov 29 14:49:43 UTC 2018 x86_64 x86_64 x86_64
↳ GNU/Linux
+ lscpu
Architecture: x86_64
CPU op-mode(s): 32-bit, 64-bit
Byte Order: Little Endian
CPU(s): 64
On-line CPU(s) list: 0-63
Thread(s) per core: 2

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```
Core(s) per socket:    16
Socket(s):             2
NUMA node(s):         2
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 85
Model name:            Intel(R) Xeon(R) Gold 6130 CPU
↳ @ 2.10GHz
Stepping:              4
CPU MHz:               1000.012
CPU max MHz:           3700.0000
CPU min MHz:           1000.0000
BogoMIPS:              4200.00
Virtualization:        VT-x
L1d cache:             32K
L1i cache:             32K
L2 cache:              1024K
L3 cache:              22528K
NUMA node0 CPU(s):
↳ 0,2,4,6,8,10,12,14,16,18,20,22,24,26,28,30,32,34
↳ 36,38,40,42,44,46,48,50,52,54,56,58,60,62
NUMA node1 CPU(s):
↳ 1,3,5,7,9,11,13,15,17,19,21,23,25,27,29,31,33,35
↳ 37,39,41,43,45,47,49,51,53,55,57,59,61,63
Flags:                 fpu vme de pse tsc msr pae mce
↳ cx8 apic sep mtrr pge mca cmov pat pse36 clflush
↳ dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx
↳ pdpe1gb rdtscp lm constant_tsc art arch_perfmon
↳ pebs bts rep_good nopl xtopology nonstop_tsc
↳ aperfmperf eagerfpu pni pclmulqdq dtes64 monitor
↳ ds_cpl vmx smx est tm2 ssse3 sdbg fma cx16 xtpr
↳ pdcm pcid dca sse4_1 sse4_2 x2apic movbe popcnt
↳ tsc_deadline_timer aes xsave avx f16c rdrand
↳ lahf_lm abm 3dnowprefetch epb cat_l3 cdp_l3
↳ intel_ppin intel_pt mba ibrs ibpb stibp
↳ tpr_shadow vnmi flexpriority ept vpid fsgsbase
↳ tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid
↳ rtm cqm mpx rdt_a avx512f avx512dq rdseed adx smap
↳ clflushopt clwb avx512cd avx512bw avx512vl
↳ xsaveopt xsavec xgetbv1 cqm_llc cqm_occup_llc
↳ cqm_mbm_total cqm_mbm_local dtherm ida arat pln
↳ pts hwp hwp_act_window hwp_epp hwp_pkg_req pku
↳ ospke spec_ctrl intel_stibp
+ cat /proc/meminfo
MemTotal:              196548732 kB
MemFree:               192091916 kB
MemAvailable:          190250680 kB
Buffers:               0 kB
Cached:                168004 kB
SwapCached:            0 kB
Active:                886296 kB
Inactive:              144728 kB
Active(anon):          864000 kB
Inactive(anon):        85436 kB
Active(file):          22296 kB
Inactive(file):        59292 kB
```

```
Unevictable:           0 kB
Mlocked:               0 kB
SwapTotal:             50331644 kB
SwapFree:              50331644 kB
Dirty:                 36 kB
Writeback:             0 kB
AnonPages:             863088 kB
Mapped:                79952 kB
Shmem:                 86416 kB
Slab:                  669784 kB
SReclaimable:          90504 kB
SUnreclaim:            579280 kB
KernelStack:           20528 kB
PageTables:            12992 kB
NFS_Unstable:          4 kB
Bounce:                0 kB
WritebackTmp:          0 kB
CommitLimit:           148606008 kB
Committed_AS:          1299492 kB
VmallocTotal:          34359738367 kB
VmallocUsed:            875196 kB
VmallocChunk:          34258257916 kB
HardwareCorrupted:     0 kB
AnonHugePages:         4096 kB
CmaTotal:              0 kB
CmaFree:               0 kB
HugePages_Total:       0
HugePages_Free:        0
HugePages_Rsvd:        0
HugePages_Surp:        0
Hugepagesize:          2048 kB
DirectMap4k:           569120 kB
DirectMap2M:           8415232 kB
DirectMap1G:           192937984 kB
+ inxi -F -c0
./collect_environment.sh: line 14: inxi: command not
↳ found
+ lslblk -a
NAME                                MAJ:MIN RM  SIZE RO TYPE
↳ MOUNTPOINT
sda                                8:0    0  1.8T  0 disk
└─sda1                             8:1    0  1.8T  0 part
    └─vg_peaks-lv_swap              253:0    0   48G  0 lvm [SWAP]
        └─vg_peaks-lv_scratch       253:1    0  1.8T  0 lvm
↳ /scratch/local
+ lsscsi -s
[0:0:0:0]    disk      SEAGATE  ST2000NM0155      DT31
↳ /dev/sda  2.00TB
+ module list
++ /uufs/chpc.utah.edu/sys/installdir/lmod/7.7.29/li
↳ bexec/lmod bash
↳ list

Currently Loaded Modules:
  1) chpc/1.0  2) intel/2019.2.187
```


Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

/0/4	processor	Intel(R)	/0/2	bridge	Sky Lake-E PCI
↪ Xeon(R) Gold 6130 CPU @ 2.10GHz			↪ Express Root Port C		
/0/100	bridge	Sky Lake-E DMI3	/0/2/0	storage	MegaRAID SAS-3
↪ Registers			↪ 3008 [Fury]		
/0/100/5	generic	Sky Lake-E	/0/6	generic	Intel Corporation
↪ MM/Vt-d Configuration Registers			/0/7	generic	Sky Lake-E RAS
/0/100/5.2	generic	Intel	↪ Configuration Registers		
↪ Corporation			/0/9	generic	Intel Corporation
/0/100/5.4	generic	Intel	/0/a	generic	Sky Lake-E CHA
↪ Corporation			↪ Registers		
/0/100/8	generic	Sky Lake-E Ubox	/0/b	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/100/8.1	generic	Sky Lake-E Ubox	/0/c	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/100/8.2	generic	Sky Lake-E Ubox	/0/d	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/100/11	generic	C620 Series	/0/e	generic	Sky Lake-E CHA
↪ Chipset Family MROM 0			↪ Registers		
/0/100/11.5	storage	C620 Series	/0/f	generic	Sky Lake-E CHA
↪ Chipset Family SSATA Controller [AHCI mode]			↪ Registers		
/0/100/14	bus	C620 Series	/0/10	generic	Sky Lake-E CHA
↪ Chipset Family USB 3.0 xHCI Controller			↪ Registers		
/0/100/14.2	generic	C620 Series	/0/11	generic	Sky Lake-E CHA
↪ Chipset Family Thermal Subsystem			↪ Registers		
/0/100/16	communication	C620 Series	/0/12	generic	Sky Lake-E CHA
↪ Chipset Family MEI Controller #1			↪ Registers		
/0/100/16.1	communication	C620 Series	/0/13	generic	Sky Lake-E CHA
↪ Chipset Family MEI Controller #2			↪ Registers		
/0/100/16.4	communication	C620 Series	/0/14	generic	Sky Lake-E CHA
↪ Chipset Family MEI Controller #3			↪ Registers		
/0/100/17	storage	C620 Series	/0/15	generic	Sky Lake-E CHA
↪ Chipset Family SATA Controller [AHCI mode]			↪ Registers		
/0/100/1c	bridge	C620 Series	/0/16	generic	Sky Lake-E CHA
↪ Chipset Family PCI Express Root Port #1			↪ Registers		
/0/100/1c.4	bridge	C620 Series	/0/17	generic	Sky Lake-E CHA
↪ Chipset Family PCI Express Root Port #5			↪ Registers		
/0/100/1c.4/0	bridge	PLDA	/0/18	generic	Sky Lake-E CHA
/0/100/1c.4/0/0	display	Integrated	↪ Registers		
↪ Matrox G200eW3 Graphics Controller			/0/19	generic	Sky Lake-E CHA
/0/100/1c.5	bridge	C620 Series	↪ Registers		
↪ Chipset Family PCI Express Root Port #6			/0/1a	generic	Sky Lake-E CHA
/0/100/1c.5/0	eth0 network	NetXtreme	↪ Registers		
↪ BCM5720 Gigabit Ethernet PCIe			/0/1b	generic	Sky Lake-E CHA
/0/100/1c.5/0.1	eth1 network	NetXtreme	↪ Registers		
↪ BCM5720 Gigabit Ethernet PCIe			/0/1c	generic	Sky Lake-E CHA
/0/100/1f	bridge	C621 Series	↪ Registers		
↪ Chipset LPC/eSPI Controller			/0/1d	generic	Sky Lake-E CHA
/0/100/1f.2	memory	Memory	↪ Registers		
↪ controller			/0/1e	generic	Sky Lake-E CHA
/0/100/1f.4	bus	C620 Series	↪ Registers		
↪ Chipset Family SMBus			/0/1f	generic	Sky Lake-E CHA
/0/100/1f.5	bus	C620 Series	↪ Registers		
↪ Chipset Family SPI Controller			/0/20	generic	Sky Lake-E CHA
			↪ Registers		

/0/21	generic	Sky Lake-E CHA	/0/3c	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/22	generic	Sky Lake-E CHA	/0/3d	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/23	generic	Sky Lake-E CHA	/0/3e	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/24	generic	Sky Lake-E CHA	/0/3f	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/25	generic	Sky Lake-E CHA	/0/40	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/26	generic	Sky Lake-E CHA	/0/41	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/27	generic	Sky Lake-E CHA	/0/42	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/28	generic	Sky Lake-E CHA	/0/43	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/29	generic	Sky Lake-E CHA	/0/44	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/2a	generic	Sky Lake-E CHA	/0/45	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/2b	generic	Sky Lake-E CHA	/0/46	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/2c	generic	Sky Lake-E CHA	/0/47	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/2d	generic	Sky Lake-E CHA	/0/48	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/2e	generic	Sky Lake-E CHA	/0/49	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/2f	generic	Sky Lake-E CHA	/0/4a	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/30	generic	Sky Lake-E CHA	/0/4b	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/31	generic	Sky Lake-E CHA	/0/4c	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/32	generic	Sky Lake-E CHA	/0/4d	generic	Intel Corporation
↪ Registers			/0/4e	generic	Sky Lake-E RAS
/0/33	generic	Sky Lake-E CHA	↪ Configuration Registers		
↪ Registers			/0/4f	generic	Intel Corporation
/0/34	generic	Sky Lake-E CHA	/0/50	generic	Intel Corporation
↪ Registers			/0/51	generic	Intel Corporation
/0/35	generic	Sky Lake-E CHA	/0/52	generic	Intel Corporation
↪ Registers			/0/53	generic	Intel Corporation
/0/36	generic	Sky Lake-E CHA	/0/54	generic	Intel Corporation
↪ Registers			/0/55	generic	Intel Corporation
/0/37	generic	Sky Lake-E CHA	/0/56	generic	Intel Corporation
↪ Registers			/0/57	generic	Intel Corporation
/0/38	generic	Sky Lake-E CHA	/0/58	generic	Intel Corporation
↪ Registers			/0/59	generic	Intel Corporation
/0/39	generic	Sky Lake-E CHA	/0/5a	generic	Intel Corporation
↪ Registers			/0/5b	generic	Intel Corporation
/0/3a	generic	Sky Lake-E CHA	/0/5c	generic	Intel Corporation
↪ Registers			/0/5d	generic	Intel Corporation
/0/3b	generic	Sky Lake-E CHA	/0/5e	generic	Intel Corporation
↪ Registers			/0/5f	generic	Intel Corporation
			/0/60	generic	Intel Corporation

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

/0/61	generic	Intel Corporation	/0/8.2	generic	Sky Lake-E CHA
/0/62	generic	Intel Corporation	↪ Registers		
/0/63	generic	Intel Corporation	/0/86	generic	Sky Lake-E CHA
/0/64	generic	Intel Corporation	↪ Registers		
/0/65	generic	Intel Corporation	/0/87	generic	Sky Lake-E CHA
/0/66	generic	Intel Corporation	↪ Registers		
/0/67	generic	Intel Corporation	/0/88	generic	Sky Lake-E CHA
/0/68	generic	Intel Corporation	↪ Registers		
/0/69	generic	Intel Corporation	/0/89	generic	Sky Lake-E CHA
/0/6a	generic	Intel Corporation	↪ Registers		
/0/6b	generic	Sky Lake-E RAS	/0/8a	generic	Sky Lake-E CHA
↪ Configuration Registers			↪ Registers		
/0/6c	generic	Intel Corporation	/0/8b	generic	Sky Lake-E CHA
/0/6d	generic	Intel Corporation	↪ Registers		
/0/6e	generic	Intel Corporation	/0/8c	generic	Sky Lake-E CHA
/0/6f	generic	Intel Corporation	↪ Registers		
/0/70	generic	Intel Corporation	/0/8d	generic	Sky Lake-E CHA
/0/71	generic	Intel Corporation	↪ Registers		
/0/72	generic	Intel Corporation	/0/8e	generic	Sky Lake-E CHA
/0/73	generic	Sky Lake-E	↪ Registers		
↪ M3KTI Registers			/0/8f	generic	Sky Lake-E CHA
/0/74	generic	Sky Lake-E	↪ Registers		
↪ M3KTI Registers			/0/90	generic	Sky Lake-E CHA
/0/75	generic	Sky Lake-E	↪ Registers		
↪ M3KTI Registers			/0/91	generic	Sky Lake-E CHA
/0/76	generic	Sky Lake-E	↪ Registers		
↪ M3KTI Registers			/0/92	generic	Sky Lake-E CHA
/0/77	generic	Sky Lake-E	↪ Registers		
↪ M3KTI Registers			/0/93	generic	Sky Lake-E CHA
/0/78	generic	Sky Lake-E	↪ Registers		
↪ M2PCI Registers			/0/94	generic	Sky Lake-E CHA
/0/79	generic	Sky Lake-E	↪ Registers		
↪ M2PCI Registers			/0/95	generic	Sky Lake-E CHA
/0/7a	generic	Sky Lake-E	↪ Registers		
↪ M2PCI Registers			/0/96	generic	Sky Lake-E CHA
/0/7b	generic	Sky Lake-E	↪ Registers		
↪ M2PCI Registers			/0/97	generic	Sky Lake-E CHA
/0/7c	generic	Sky Lake-E	↪ Registers		
↪ MM/Vt-d Configuration Registers			/0/98	generic	Sky Lake-E CHA
/0/7d	generic	Intel Corporation	↪ Registers		
/0/7e	generic	Intel Corporation	/0/99	generic	Sky Lake-E CHA
/0/7f	generic	Sky Lake-E Ubox	↪ Registers		
↪ Registers			/0/9a	generic	Sky Lake-E CHA
/0/80	generic	Sky Lake-E Ubox	↪ Registers		
↪ Registers			/0/9b	generic	Sky Lake-E CHA
/0/81	generic	Sky Lake-E Ubox	↪ Registers		
↪ Registers			/0/9c	generic	Sky Lake-E CHA
/0/82	generic	Intel Corporation	↪ Registers		
/0/83	generic	Sky Lake-E RAS	/0/9d	generic	Sky Lake-E CHA
↪ Configuration Registers			↪ Registers		
/0/84	generic	Intel Corporation	/0/9e	generic	Sky Lake-E CHA
/0/85	generic	Sky Lake-E CHA	↪ Registers		
↪ Registers			/0/9f	generic	Sky Lake-E CHA
/0/8.1	generic	Sky Lake-E CHA	↪ Registers		
↪ Registers					

/0/a0	generic	Sky Lake-E CHA	/0/bb	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/a1	generic	Sky Lake-E CHA	/0/bc	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/a2	generic	Sky Lake-E CHA	/0/bd	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/a3	generic	Sky Lake-E CHA	/0/be	generic	Sky Lake-E CHA
↪ Registers			↪ Registers		
/0/a4	generic	Sky Lake-E CHA	/0/bf	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/a5	generic	Sky Lake-E CHA	/0/c0	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/a6	generic	Sky Lake-E CHA	/0/c1	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/a7	generic	Sky Lake-E CHA	/0/c2	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/a8	generic	Sky Lake-E CHA	/0/c3	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/a9	generic	Sky Lake-E CHA	/0/c4	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/aa	generic	Sky Lake-E CHA	/0/c5	generic	Sky Lake-E PCU
↪ Registers			↪ Registers		
/0/ab	generic	Sky Lake-E CHA	/0/0	bridge	Sky Lake-E PCI
↪ Registers			↪ Express Root Port A		
/0/ac	generic	Sky Lake-E CHA	/0/0/0	ib0 network	MT27700 Family
↪ Registers			↪ [ConnectX-4]		
/0/ad	generic	Sky Lake-E CHA	/0/c6	generic	Intel Corporation
↪ Registers			/0/c7	generic	Sky Lake-E RAS
/0/ae	generic	Sky Lake-E CHA	↪ Configuration Registers		
↪ Registers			/0/c8	generic	Intel Corporation
/0/af	generic	Sky Lake-E CHA	/0/8	generic	Intel Corporation
↪ Registers			/0/c9	generic	Intel Corporation
/0/b0	generic	Sky Lake-E CHA	/0/ca	generic	Intel Corporation
↪ Registers			/0/cb	generic	Intel Corporation
/0/b1	generic	Sky Lake-E CHA	/0/cc	generic	Intel Corporation
↪ Registers			/0/cd	generic	Intel Corporation
/0/b2	generic	Sky Lake-E CHA	/0/ce	generic	Intel Corporation
↪ Registers			/0/cf	generic	Intel Corporation
/0/b3	generic	Sky Lake-E CHA	/0/d0	generic	Intel Corporation
↪ Registers			/0/d1	generic	Intel Corporation
/0/b4	generic	Sky Lake-E CHA	/0/d2	generic	Intel Corporation
↪ Registers			/0/d3	generic	Intel Corporation
/0/b5	generic	Sky Lake-E CHA	/0/d4	generic	Intel Corporation
↪ Registers			/0/d5	generic	Intel Corporation
/0/b6	generic	Sky Lake-E CHA	/0/d6	generic	Intel Corporation
↪ Registers			/0/d7	generic	Intel Corporation
/0/b7	generic	Sky Lake-E CHA	/0/d8	generic	Intel Corporation
↪ Registers			/0/d9	generic	Intel Corporation
/0/b8	generic	Sky Lake-E CHA	/0/da	generic	Intel Corporation
↪ Registers			/0/db	generic	Intel Corporation
/0/b9	generic	Sky Lake-E CHA	/0/dc	generic	Intel Corporation
↪ Registers			/0/dd	generic	Intel Corporation
/0/ba	generic	Sky Lake-E CHA	/0/de	generic	Intel Corporation
↪ Registers			/0/df	generic	Intel Corporation
			/0/e0	generic	Intel Corporation

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

/0/e1	generic	Intel Corporation	ZAP_LIBPATH=/opt/ovis/lib64/ovis-lib
/0/5	generic	Intel Corporation	KSH_AUTOLOAD=1
/0/5.2	generic	Sky Lake-E RAS	MODULE_VERSION_STACK=3.2.10.6
↪ Configuration Registers			LESSKEY=/etc/lesskey.bin
/0/5.4	generic	Intel	SLURM_NODELIST=nid05865
↪ Corporation			SLURM_CHECKPOINT_IMAGE_DIR=/var/slurm/checkpoint
/0/e2	generic	Intel Corporation	GNU_VERSION=8.2.0
/0/e3	generic	Intel Corporation	DVS_MAXNODES=1__
/0/e4	generic	Intel Corporation	PE_CXX_PKGCONFIG_LIBS=mpichcxx
/0/e5	generic	Intel Corporation	PE_TPSSL_DEFAULT_GENCOMPS_INTEL_x86_skylake=160
/0/e6	generic	Intel Corporation	PE_PETSC_DEFAULT_GENCOMPS_CRAY_skylake=86
/0/e7	generic	Intel Corporation	PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_sandybridge=8.6
/0/e8	generic	Sky Lake-E	PE_PAPI_DEFAULT_ACCEL_FAMILY_LIBS_nvidia=-lcupti,-l_
↪ M3KTI Registers			↪ cudart,-lcuda
/0/e9	generic	Sky Lake-E	NNTPSERVER=news
↪ M3KTI Registers			MANPATH=/usr/common/software/man:/usr/common/mss/man_
/0/ea	generic	Sky Lake-E	↪ :/usr/common/nsg/man:/opt/gcc/8.2.0/snos/share/m_
↪ M3KTI Registers			↪ an:/usr/common/software/man:/usr/common/mss/man:_
/0/eb	generic	Sky Lake-E	↪ /usr/common/nsg/man:/opt/cray/pe/mpt/7.7.3/gni/m_
↪ M3KTI Registers			↪ an/mpich:/opt/cray/pe/atp/2.1.3/man:/opt/cray/al_
/0/ec	generic	Sky Lake-E	↪ ps/6.6.43-6.0.7.1_5.45__ga796da32.ari/man:/opt/c_
↪ M3KTI Registers			↪ ray/job/2.2.3-6.0.7.1_5.43__g6c4e934.ari/man:/op_
/0/ed	generic	Sky Lake-E	↪ t/cray/pe/pmi/5.0.14/man:/opt/cray/pe/libsci/18._
↪ M2PCI Registers			↪ 07.1/man:/opt/cray/pe/man/csmlversion:/opt/cray/_
/0/ee	generic	Sky Lake-E	↪ pe/craype/2.5.15/man:/opt/intel/compilers_and_li_
↪ M2PCI Registers			↪ braries_2019.3.199/linux/man/common:/usr/syscom/_
/0/ef	generic	Sky Lake-E	↪ nsg/man:/opt/cray/pe/modules/3.2.10.6/share/man:_
↪ M2PCI Registers			↪ /usr/local/man:/usr/share/man:/opt/cray/share/ma_
/0/f0	generic	Sky Lake-E	↪ n:/opt/cray/pe/man:/opt/cray/share/man:/opt/cray_
↪ M2PCI Registers			↪ /share/man
/0/f1	system	PnP device	SLURM_JOB_NAME=run8x8x8.sh
↪ PNP0b00			XDG_SESSION_ID=4331
/0/f2	system	PnP device	PE_TRILINOS_DEFAULT_GENCOMPS_CRAY_x86_64=87
↪ PNP0c02			PE_TPSSL_64_DEFAULT_GENCOMPS_INTEL_interlagos=160
/0/f3	communication	PnP device	PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_mic_knl=16.0
↪ PNP0501			PE_FFTW_DEFAULT_TARGET_mic_knl=mic_knl
/0/f4	communication	PnP device	CRAY_UDREG_INCLUDE_OPTS=-I/opt/cray/udreg/2.3.2-6.0._
↪ PNP0501			↪ 7.1_5.13__g5196236.ari/include
/0/f5	system	PnP device	HOSTNAME=nid05865
↪ PNP0c02			SLURM_TOPOLOGY_ADDR=s34.s21.nid05865
/0/f6	system	PnP device	SLURMD_NODENAME=nid05865
↪ PNP0c02			PE_TRILINOS_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cra_
WARNING: output may be incomplete or inaccurate, you			↪ y/pe/trilinos/12.12.1.1/@PRGENV@/@PE_TRILINOS_DE_
↪ should run this program as super-user.			↪ FAULT_GENCOMPS@/@PE_TRILINOS_DEFAULT_TARGET@/lib_
=====			↪ /pkgconfig
↪ =====			PE_SMA_DEFAULT_COMPFLAG_GNU=-fcray-pointer
			PE_PARALLEL_NETCDF_DEFAULT_VOLATILE_PKGCONFIG_PATH=/_
			↪ opt/cray/pe/parallel-netcdf/1.8.1.3/@PRGENV@/@PE_
			↪ _PARALLEL_NETCDF_DEFAULT_GENCOMPS@/lib/pkgconfig
			PE_NETCDF_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/_
			↪ pe/netcdf/4.6.1.3/@PRGENV@/@PE_NETCDF_DEFAULT_GE_
			↪ NCOMPS@/lib/pkgconfig

KNL node:
PE_TPSSL_64_DEFAULT_GENCOMPS_INTEL_mic_knl=160
PE_SMA_DEFAULT_PKGCONFIG_VARIABLES=PE_SMA_COMPFLAG@_
↪ prgenv@
PE_LIBSCI_VOLATILE_PRGENV=CRAY GNU INTEL
MKLR00T=/opt/intel/compilers_and_libraries_2019.3.19_
↪ 9/linux/mkl

```

LIBRARYMODULES=acml:alps:cray-dwarf:cray-fftw:cray-g
↪ a:cray-hdf5:cray-hdf5-parallel:cray-libsci:cray-
↪ libsci_acc:cray-mpich:cray-mpich-abi:cray-mpich2
↪ :cray-netcdf:cray-netcdf-hdf5parallel:cray-paral
↪ lel-netcdf:cray-petsc:cray-petsc-complex:cray-sh
↪ mem:cray-tpsl:cray-trilinos:cuda toolkit:fftw:ga:
↪ hdf5:hdf5-parallel:iobuf:libfast:netcdf:netcdf-h
↪ df5parallel:ntk:onesided:papi:petsc:petsc-comple
↪ x:pmi:tpsl:trilinos:xt-libsci:xt-mpich2:xt-mpt:x
↪ t-papi
CRAY_SITE_LIST_DIR=/etc/opt/cray/pe/modules
XKEYSYMDB=/usr/X11R6/lib/X11/XKeysymDB
SLURM_PRIO_PROCESS=0
RCLOCAL_BASEOPTS=true
INTEL_LICENSE_FILE=28518@crayintel.licenses.nersc.go
↪ v:28518@intel.licenses.nersc.gov
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
PE_SMA_DEFAULT_COMPFLAG=
PE_MPICH_ALTERNATE_LIBS_dpm=dpm
PE_HDF5_DEFAULT_GENCOMPILERS_GNU=8.2 7.1 6.1 5.3 4.9
PE_ENV=INTEL
SLURM_NODE_ALIASES=(null)
PKGCONFIG_ENABLED=1
PE_TPSL_DEFAULT_GENCOMPS_CRAY_x86_skylake=86
HOST=cori11
TERM=xterm-256color
SHELL=/bin/bash
PE_TPSL_DEFAULT_GENCOMPILERS_GNU_x86_skylake=8.2 7.1
↪ 6.1
PE_PETSC_DEFAULT_GENCOMPS_CRAY_sandybridge=86
INTEL_MINOR_VERSION=19
PROFILEREAD=true
HISTSIZE=1000
SLURM_JOB_QOS=regular_1
KMP_HOT_TEAMS_MODE=1
TMPDIR=/tmp
PE_TRILINOS_DEFAULT_VOLATILE_PRGENV=CRAY GNU INTEL
PE_TPSL_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_LIBSCI
PE_TPSL_DEFAULT_GENCOMPS_GNU_sandybridge=82 71 53 49
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_x86_skylake=160
PE_PETSC_DEFAULT_GENCOMPS_INTEL_haswell=160
PE_PETSC_DEFAULT_GENCOMPS_GNU_haswell=71 53 49
PE_PARALLEL_NETCDF_DEFAULT_VOLATILE_PRGENV=GNU
PE_NETCDF_DEFAULT_VOLATILE_PRGENV=GNU
CRAY_XPMEM_POST_LINK_OPTS=-L/opt/cray/xpmem/2.2.15-6
↪ .0.7.1.5.11__g7549d06.ari/lib64
CRAY_UGNI_POST_LINK_OPTS=-L/opt/cray/ugni/6.0.14.0-6
↪ .0.7.1.3.13__gea11d3d.ari/lib64
CRAYPE_DIR=/opt/cray/pe/craype/2.5.15
SSH_CLIENT=155.97.232.235 60438 22
SLURM_TOPOLOGY_ADDR_PATTERN=switch.switch.node
SITE_MODULE_NAMES=darshan
ALT_LINKER=/usr/common/software/altd/2.0/bin/ld
ALT_SELECT_OFF_USERS=
CRAY_MPICH2_DIR=/opt/cray/pe/mpt/7.7.3/gni/mpich-int
↪ el/16.0

```

```

PE_PETSC_DEFAULT_GENCOMPS_CRAY_interlagos=86
PE_NETCDF_HDF5PARALLEL_DEFAULT_VOLATILE_PKGCONFIG_PA
↪ TH=/opt/cray/pe/netcdf-hdf5parallel/4.6.1.3/@PRG
↪ ENV@/@PE_NETCDF_HDF5PARALLEL_DEFAULT_GENCOMPS@/l
↪ ib/pkgconfig
PE_HDF5_PARALLEL_DEFAULT_VOLATILE_PKGCONFIG_PATH=/op
↪ t/cray/pe/hdf5-parallel/1.10.2.0/@PRGENV@/@PE_HD
↪ F5_PARALLEL_DEFAULT_GENCOMPS@/lib/pkgconfig
PE_HDF5_DEFAULT_VOLATILE_PRGENV=GNU
PE_FFTW_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe
↪ /fftw/3.3.8.1/@PE_FFTW_DEFAULT_TARGET@/lib/pkgco
↪ nfig
LIBRARY_PATH=/opt/intel/compilers_and_libraries_2019
↪ .3.199/linux/compiler/lib/intel64:/opt/intel/com
↪ pilers_and_libraries_2019.3.199/linux/mkl/lib/in
↪ tel64
CONDA_SHLVL=1
PYTHON_DIR=/usr/common/software/python/3.6-anaconda-
↪ 4.4
ALT_SELECT_ON=0
PE_TPSL_DEFAULT_GENCOMPS_CRAY_mic_knl=86
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_interlagos=8.6
PE_LIBSCI_DEFAULT_GENCOMPS_GNU_x86_64=71 61 51 49
PE_GA_DEFAULT_VOLATILE_PRGENV=GNU
INTEL_PATH=/opt/intel/compilers_and_libraries_2019.3
↪ .199
CONDA_PROMPT_MODIFIER=(stenv)
PE_MPICH_GENCOMPS_GNU=71 51 49
PE_TPSL_DEFAULT_GENCOMPS_INTEL_x86_64=160
PE_PKGCONFIG_PRODUCTS=PE_MPICH:PE_LIBSCI
PE_MPICH_DEFAULT_GENCOMPILERS_GNU=7.1 5.1 4.9
FPATH=/opt/cray/pe/modules/3.2.10.6/init/sh_funcs/n
↪ o_redirect:/opt/cray/pe/modules/3.2.10.6/init/sh
↪ _funcs/no_redirect:/opt/cray/pe/modules/3.2.10.6
↪ /init/sh_funcs/no_redirect
MORE=-s1
ALT_VERBOSE=0
PE_TPSL_64_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray
↪ /pe/tpsl/18.06.1/@PRGENV@64/@PE_TPSL_64_DEFAULT_
↪ GENCOMPS@/@PE_TPSL_64_DEFAULT_TARGET@/lib/pkgcon
↪ fig
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_haswell=86
PE_PETSC_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_LIBSC
↪ I:PE_HDF5_PARALLEL:PE_TPSL
PE_PAPI_DEFAULT_ACCEL_LIBS_nvidia35=-lcupti,-lcudar
↪ t,-lcuda
SLURM_SPANK_NERSC_ZONESORT_INTERVAL=0
PE_TRILINOS_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.7
PE_CRAY_DEFAULT_FIXED_PKGCONFIG_PATH=/opt/cray/pe/pa
↪ rallel-netcdf/1.8.1.3/CRAY/8.6/lib/pkgconfig:/op
↪ t/cray/pe/netcdf-hdf5parallel/4.6.1.3/CRAY/8.6/l
↪ ib/pkgconfig:/opt/cray/pe/netcdf/4.6.1.3/CRAY/8.
↪ 6/lib/pkgconfig:/opt/cray/pe/hdf5-parallel/1.10.
↪ 2.0/CRAY/8.6/lib/pkgconfig:/opt/cray/pe/hdf5/1.1
↪ 0.2.0/CRAY/8.6/lib/pkgconfig:/opt/cray/pe/ga/5.3
↪ .0.8/CRAY/8.6/lib/pkgconfig

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```
PE_FORTRAN_PKGCONFIG_LIBS=mpichf90
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_sandybridge=8.6
PE_PETSC_DEFAULT_GENCOMPS_CRAY_x86_64=86
PE_LIBSCI_DEFAULT_OMP_REQUIRES_openmp=_mp
SSH_TTY=/dev/pts/1
ALLINEA_QUEUE_DLL=/opt/cray/pe/mpt/7.7.3/gni/mpich-i
↳ ntel/16.0/lib/libtvmptich.so.3.0.1
PE_SMA_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/
↳ mpt/7.7.3/gni/sma@PE_SMA_DEFAULT_DIR_DEFAULT64@/
↳ lib64/pkgconfig
PYTHONUSERBASE=/global/homes/z/USER/.local/cori/3.6-
↳ anaconda-4.4
CRAY_MPICH_BASEDIR=/opt/cray/pe/mpt/7.7.3/gni
PE_TRILINOS_DEFAULT_GENCOMPS_INTEL_x86_64=160
ALPS_APP_ID=18446744065140005382
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_haswell=160
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_x86_skylake=86
PE_NETCDF_HDF5PARALLEL_DEFAULT_GENCOMPILERS_GNU=8.2
↳ 7.1 6.1 5.3 4.9
PE_HDF5_PARALLEL_DEFAULT_GENCOMPILERS_GNU=8.2 7.1 6.1
↳ 5.3 4.9
JRE_HOME=/usr/lib64/jvm/java/jre
USER=USER
SLURM_NNODES=1
PE_TRILINOS_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
PE_TRILINOS_DEFAULT_GENCOMPILERS_GNU_x86_64=8.2 7.3
↳ 5.1 4.9
PE_TPSL_DEFAULT_GENCOMPS_CRAY_x86_64=86
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_mic_knl=16.0
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_interlagos=16.0
PE_LIBSCI_DEFAULT_VOLATILE_PRGENV=CRAY GNU INTEL
PE_FFTW_DEFAULT_TARGET_interlagos=interlagos
LS_COLORS=no=00:fi=00:di=01;34:ln=00;36:pi=40;33:so=
↳ 01;35:do=01;35:bd=40;33;01:cd=40;33;01:or=41;33;
↳ 01:ex=00;32:*.cmd=00;32:*.exe=01;32:*.com=01;32:
↳ *.bat=01;32:*.btm=01;32:*.dll=01;32:*.tar=00;31:
↳ *.tbz=00;31:*.tgz=00;31:*.rpm=00;31:*.deb=00;31:
↳ *.arj=00;31:*.taz=00;31:*.lzh=00;31:*.lzma=00;31
↳ :*.zip=00;31:*.zoo=00;31:*.z=00;31:*.Z=00;31:*.g
↳ z=00;31:*.bz2=00;31:*.tb2=00;31:*.tz2=00;31:*.tb
↳ z2=00;31:*.xz=00;31:*.avi=01;35:*.bmp=01;35:*.fl
↳ i=01;35:*.gif=01;35:*.jpg=01;35:*.jpeg=01;35:*.m
↳ ng=01;35:*.mov=01;35:*.mpg=01;35:*.pcx=01;35:*.p
↳ bm=01;35:*.pgm=01;35:*.png=01;35:*.ppm=01;35:*.t
↳ ga=01;35:*.tif=01;35:*.xbm=01;35:*.xpm=01;35:*.d
↳ l=01;35:*.gl=01;35:*.wmv=01;35:*.aiff=00;32:*.au
↳ =00;32:*.mid=00;32:*.mp3=00;32:*.ogg=00;32:*.voc
↳ =00;32:*.wav=00;32:
LD_LIBRARY_PATH=/opt/gcc/8.2.0/snos/lib64:/opt/cray/
↳ job/2.2.3-6.0.7.1_5.43_g6c4e934.ari/lib64:/opt/
↳ intel/compilers_and_libraries_2019.3.199/linux/c
↳ ompiler/lib/intel64:/opt/intel/compilers_and_lib
↳ raries_2019.3.199/linux/mkl/lib/intel64:/usr/sys
↳ com/nsng/lib
CSCRATCH=/global/cscratch1/sd/USER
CSHRCREAD=true
```

```
PE_MPICH_FIXED_PRGENV=INTEL
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_haswell=16.0
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_sandybridge=8.2
↳ 7.1 5.3 4.9
PE_PKGCONFIG_LIBS=darshan-runtime:mpich:AtpSigHandle
↳ r:cray-rca:libsci:mpi:libsci
PE_PETSC_DEFAULT_VOLATILE_PRGENV=CRAY CRAY64 GNU
↳ GNU64 INTEL INTEL64
PE_LIBSCI_PKGCONFIG_VARIABLES=PE_LIBSCI_OMP_REQUIRES
↳ _@openmp@:PE_SCI_EXT_LIBPATH:PE_SCI_EXT_LIBNAME
CRAY_RCA_POST_LINK_OPTS=-L/opt/cray/rca/2.2.18-6.0.7
↳ .1_5.47_g2aa4f39.ari/lib64
↳ -lrca
PE_TPSL_DEFAULT_GENCOMPS_GNU_haswell=82 71 53 49
PE_PETSC_DEFAULT_GENCOMPS_INTEL_sandybridge=160
PE_PETSC_DEFAULT_GENCOMPS_INTEL_interlagos=160
PE_PETSC_DEFAULT_GENCOMPS_GNU_sandybridge=71 53 49
PE_PETSC_DEFAULT_GENCOMPS_GNU_interlagos=71 53 49
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_skylake=16.0
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_mic_knl=8.6
XNLSPATH=/usr/share/X11/nls
CONDA_EXE=/usr/common/software/python/3.6-anaconda-4
↳ .4/bin/conda
ALTD_ON=1
MPICH_DIR=/opt/cray/pe/mpt/7.7.3/gni/mpich-intel/16.0
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_sandybridge=160
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_interlagos=8.2
↳ 7.1 5.3 4.9
PE_PETSC_DEFAULT_GENCOMPS_INTEL_mic_knl=160
PE_PETSC_DEFAULT_GENCOMPS_GNU_mic_knl=53
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_haswell=8.6
PE_PAPI_DEFAULT_PKGCONFIG_VARIABLES=PE_PAPI_ACCEL_LI
↳ BS_@accelerator@
PE_LIBSCI_DEFAULT_GENCOMPS_CRAY_x86_64=86
MPICH_ABORT_ON_ERROR=1
INTEL_VERSION=19.0.3.199
LDAPTLS_REQCERT=never
MPICH_MPIIO_DVS_MAXNODES=32
PE_MPICH_FORTRAN_PKGCONFIG_LIBS=mpichf90
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_haswell=8.6
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_sandybridge=16.0
PE_NETCDF_HDF5PARALLEL_DEFAULT_REQUIRED_PRODUCTS=PE_
↳ HDF5_PARALLEL
PE_HDF5_PARALLEL_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH
PE_FFTW_DEFAULT_TARGET_sandybridge=sandybridge
PE_FFTW_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH
CPATH=/opt/intel/compilers_and_libraries_2019.3.199/
↳ linux/mkl/include
ATP_POST_LINK_OPTS=-Wl,-L/opt/cray/pe/atp/2.1.3/libA
↳ pp/
HOSTTYPE=x86_64
SSH_AUTH_SOCKET=/tmp/ssh-6PdT5S1Wat/agent.21653
SLURM_JOBID=20388358
RCLOCAL_PRGENV=true
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_mic_knl=5.3
```

```

KMP_HOT_TEAMS_MAX_LEVEL=2
GCC_VERSION=8.2.0
PE_TPSSL_DEFAULT_GENCOMPS_GNU_interlagos=82 71 53 49
PE_TPSSL_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
PE_PRODUCT_LIST=CRAYPE_MIC-KNL:CRAY_RCA:CRAY_ALPS:DV
↪ S:CRAY_XPMEM:CRAY_DMAPP:CRAY_PMI:CRAY_UGNI:CRAY_
↪ UDREG:CRAY_LIBSCI:CRAYPE:INTEL
PE_LIBSCI_GENCOMPS_INTEL_x86_64=160
PE_LIBSCI_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
FROM_HEADER=
CRAY_MPICH_ROOTDIR=/opt/cray/pe/mpit/7.7.3
PE_TPSSL_DEFAULT_GENCOMPS_GNU_x86_skylake=82 71 61
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_x86_64=7.1 5.3 4.9
PAGER=less
ALPS_LLI_STATUS_OFFSET=1
PE_MPICH_MODULE_NAME=cray-mpich
PE_MPICH_GENCOMPILERS_CRAY=8.6
PE_TPSSL_64_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
PE_PETSC_DEFAULT_GENCOMPS_INTEL_skylake=160
PE_PETSC_DEFAULT_GENCOMPS_GNU_skylake=61
PE_LIBSCI_GENCOMPILERS_GNU_x86_64=7.1 6.1 5.1 4.9
CSHEDIT=emacs
ALPS_APP_PE=0
PE_TPSSL_DEFAULT_GENCOMPS_CRAY_sandybridge=86
PE_TPSSL_DEFAULT_GENCOMPS_CRAY_haswell=86
PE_TPSSL_64_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_LIB
↪ SCI
PE_MPICH_TARGET_VAR_nvidia20=-lcudart
PE_MPICH_DEFAULT_VOLATILE_PRGENV=CRAY_GNU
PE_LIBSCI_GENCOMPS_CRAY_x86_64=86
PE_LIBSCI_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
INTEL_MAJOR_VERSION=19
XDG_CONFIG_DIRS=/etc/xdg
PE_TPSSL_64_DEFAULT_GENCOMPS_GNU_x86_64=82 71 53 49
PE_TPSSL_64_DEFAULT_GENCOMPS_GNU_mic_knl=71 53
PE_PARALLEL_NETCDF_DEFAULT_GENCOMPS_GNU=51 49
PE_NETCDF_DEFAULT_GENCOMPS_GNU=
PE_LIBSCI_PKGCONFIG_LIBS=libsci_mpi:libsci
NLSPATH=/opt/intel/compilers_and_libraries_2019.3.19
↪ 9/linux/compiler/lib/intel64/locale/%l_%t/%N:/op
↪ t/intel/compilers_and_libraries_2019.3.199/linux
↪ /mkl/lib/intel64/locale/%l_%t/%N
DVS_VERSION=0.9.0
CRAY_LIBSCI_DIR=/opt/cray/pe/libsci/18.07.1
CRAY_LIBSCI_BASE_DIR=/opt/cray/pe/libsci/18.07.1
CRAY_DMAPP_INCLUDE_OPTS=-I/opt/cray/dmapp/7.1.1-6.0.
↪ 7.1_5.45__g5a674e0.ari/include
↪ -I/opt/cray/gni-headers/5.0.12.0-6.0.7.1_3.11__g
↪ 3b1768f.ari/include

```

```

USERMODULES=PrgEnv-cray:PrgEnv-gnu:PrgEnv-intel:PrgE
↪ nv-pathscale:PrgEnv-pgi:acml:alps:apprentice:app
↪ rentice2:atp:blcr:cce:chapel:cray-cddb:cray-fftw
↪ :cray-ga:cray-hdf5:cray-hdf5-parallel:cray-lgdb:
↪ cray-libsci:cray-libsci_acc:cray-mpich:cray-mpic
↪ h-compat:cray-mpich2:cray-netcdf:cray-netcdf-hdf
↪ 5parallel:cray-parallel-netcdf:cray-petsc:cray-p
↪ etsc-complex:cray-shmem:cray-snp launcher:cray-tp
↪ sl:cray-trilinos:craypat:craype:craypkg-gen:cuda
↪ toolkit:ddt:fftw:ga:gcc:hdf5:hdf5-parallel:intel
↪ :iobuf:java:lgdb:libfast:libsci_acc:mpich1:netcd
↪ f:netcdf-hdf5parallel:netcdf-nofsync:netcdf-nofs
↪ ync-hdf5parallel:ntk:onesided:papi:parallel-netc
↪ df:pathscale:perf tools:perf tools-lite:petsc:pets
↪ c-complex:pgi:pmi:stat:totalview:tpsl:trilinos:x
↪ t-asyncpe:xt-craypat:xt-lgdb:xt-libsci:xt-mpich2
↪ :xt-mpt:xt-papi:xt-shmem:xt-totalview
LIBGL_DEBUG=quiet
MINICOM=-c on
LIBGL_ALWAYS_INDIRECT=1
PE_MPICH_GENCOMPILERS_GNU=7.1 5.1 4.9
PE_TPSSL_DEFAULT_GENCOMPS_CRAY_interlagos=86
PE_TPSSL_DEFAULT_GENCOMPILERS_GNU_x86_64=8.2 7.1 5.3
↪ 4.9
PE_PKGCONFIG_DEFAULT_PRODUCTS=PE_TRILINOS:PE_TPSSL_64
↪ :PE_TPSSL:PE_PETSC:PE_PARALLEL_NETCDF:PE_NETCDF_H
↪ DF5PARALLEL:PE_NETCDF:PE_MPICH:PE_LIBSCI:PE_HDF5
↪ _PARALLEL:PE_HDF5:PE_GA:PE_FFTW2:PE_FFTW
PE_HDF5_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe
↪ /hdf5/1.10.2.0/@PRGENV@/@PE_HDF5_DEFAULT_GENCOMP
↪ S@/lib/pkgconfig
MODULE_VERSION=3.2.10.6
MAIL=/var/mail/USER
PATH=/global/homes/z/USER/.conda/envs/stenv/bin:/usr
↪ /common/software/python/3.6-anaconda-4.4/condaabi
↪ n:/usr/common/software/bin:/usr/common/mss/bin:/
↪ usr/common/nsg/bin:/usr/common/software/python/3
↪ .6-anaconda-4.4/bin:/usr/common/software/python/
↪ 3.6-anaconda-4.4/lib/python3.6/site-packages/mpi
↪ 4py/bin:/opt/gcc/8.2.0/bin:/usr/common/software/
↪ darshan/3.1.4/bin:/usr/common/software/altd/2.0/
↪ bin:/usr/common/software/bin:/usr/common/mss/bin
↪ :/usr/common/nsg/bin:/opt/cray/pe/mpit/7.7.3/gni/
↪ bin:/opt/cray/rca/2.2.18-6.0.7.1_5.47__g2aa4f39.
↪ ari/bin:/opt/cray/alps/6.6.43-6.0.7.1_5.45__ga79
↪ 6da32.ari/sbin:/opt/cray/job/2.2.3-6.0.7.1_5.43_
↪ _g6c4e934.ari/bin:/opt/cray/pe/craype/2.5.15/bin
↪ :/opt/intel/compilers_and_libraries_2019.3.199/l
↪ inux/bin/intel64:/opt/ovis/bin:/opt/ovis/sbin:/u
↪ sr/syscom/nsg/sbin:/usr/syscom/nsg/bin:/opt/cray
↪ /pe/modules/3.2.10.6/bin:/usr/local/bin:/usr/bin
↪ :/bin:/usr/bin/X11:/usr/games:/usr/lib/mit/bin:/
↪ usr/lib/mit/sbin:/opt/cray/pe/bin
SLURM_TASKS_PER_NODE=1
PE_TPSSL_DEFAULT_GENCOMPILERS_GNU_haswell=8.2 7.1 5.3
↪ 4.9

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_x86_skylake=8.2
↪ 7.1 6.1
PE_PETSC_DEFAULT_GENCOMPS_CRAY_mic_knl=86
PE_PARALLEL_NETCDF_DEFAULT_GENCOMPILERS_GNU=5.1 4.9
PE_NETCDF_DEFAULT_GENCOMPILERS_GNU=8.2 7.1 6.1 5.3 4.9
PE_FFTW_DEFAULT_TARGET_abudhabi=abudhabi
ATP_IGNORE_SIGTERM=1
XTPE_NETWORK_TARGET=aries
CPU=x86_64
SLURM_WORKING_CLUSTER=cori:ctl1:6817:8448
_=/usr/bin/env
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_x86_skylake=8.6
PE_SMA_DEFAULT_DIR_CRAY_DEFAULT64=64
PE_NETCDF_HDF5PARALLEL_DEFAULT_GENCOMPS_GNU=
PE_NETCDF_HDF5PARALLEL_DEFAULT_FIXED_PRGENV=CRAY
↪ INTEL
PE_HDF5_PARALLEL_DEFAULT_GENCOMPS_GNU=
PE_HDF5_PARALLEL_DEFAULT_FIXED_PRGENV=CRAY INTEL
LDMSD_PLUGIN_LIBPATH=/opt/ovis/lib64/ovis-ldms
JAVA_BINDIR=/usr/lib64/jvm/java/bin
SLURM_CPUS_PER_TASK=256
SLURM_JOB_ID=20388358
PMI_NO_FORK=1
PE_TPSL_DEFAULT_GENCOMPS_INTEL_interlagos=160
PE_TPSL_DEFAULT_GENCOMPILERS_CRAY_mic_knl=8.6
PE_TPSL_64_DEFAULT_VOLATILE_PRGENV=CRAY CRAY64 GNU
↪ GNU64 INTEL INTEL64
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_sandybridge=86
CRAY_UDREG_POST_LINK_OPTS=-L/opt/cray/udreg/2.3.2-6.
↪ 0.7.1_5.13__g5196236.ari/lib64
CONDA_PREFIX=/global/homes/z/USER/.conda/envs/stenv
PE_MPICH_VOLATILE_PRGENV=CRAY GNU
PE_TPSL_DEFAULT_GENCOMPS_GNU_mic_knl=71 53
CRAY_ALPS_POST_LINK_OPTS=-L/opt/cray/alps/6.6.43-6.0
↪ .7.1_5.45__ga796da32.ari/lib64
CRAYPE_VERSION=2.5.15
INPUTRC=/etc/inputrc
PWD=/global/homes/z/USER/brickv2/stest/Author-Kit
SLURM_JOB_USER=USER
PE_TPSL_DEFAULT_GENCOMPILERS_INTEL_haswell=16.0
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_sandybridge=7.1
↪ 5.3 4.9
PE_MPICH_DEFAULT_GENCOMPS_CRAY=86
PE_LIBSCI_DEFAULT_OMP_REQUIRES=
_LMFILES=/opt/cray/pe/modulefiles/modules/3.2.10.6:
↪ /usr/syscom/nsg/modulefiles/nsg/1.2.0:/opt/modul
↪ efiles/intel/19.0.3.199:/opt/cray/pe/craype/2.5.
↪ 15/modulefiles/craype-network-aries:/opt/cray/pe
↪ /modulefiles/craype/2.5.15:/opt/cray/pe/modulefi
↪ les/cray-libsci/18.07.1:/opt/cray/ari/modulefile
↪ s/udreg/2.3.2-6.0.7.1_5.13__g5196236.ari:/opt/cr
↪ ay/ari/modulefiles/ugni/6.0.14.0-6.0.7.1_3.13__g
↪ ea11d3d.ari:/opt/cray/pe/modulefiles/pmi/5.0.14:
↪ /opt/cray/ari/modulefiles/dmapp/7.1.1-6.0.7.1_5.
↪ 45__g5a674e0.ari:/opt/cray/ari/modulefiles/gni-h
↪ eaders/5.0.12.0-6.0.7.1_3.11__g3b1768f.ari:/opt/
↪ cray/ari/modulefiles/xpmem/2.2.15-6.0.7.1_5.11__
↪ g7549d06.ari:/opt/cray/ari/modulefiles/job/2.2.3
↪ -6.0.7.1_5.43__g6c4e934.ari:/opt/cray/ari/modulef
↪ iles/dvs/2.7.2.2.118-6.0.7.1_10.1__g58b37a2:/opt
↪ /cray/ari/modulefiles/alps/6.6.43-6.0.7.1_5.45__
↪ ga796da32.ari:/opt/cray/ari/modulefiles/rca/2.2.
↪ 18-6.0.7.1_5.47__g2aa4f39.ari:/opt/cray/pe/modul
↪ efiles/atp/2.1.3:/opt/cray/pe/modulefiles/PrgEnv
↪ -intel/6.0.4:/opt/cray/pe/craype/2.5.15/modulefil
↪ es/craype-mic-knl:/opt/cray/pe/modulefiles/cray-
↪ mpich/7.7.3:/usr/common/software/modulefiles/alt
↪ d/2.0:/usr/common/software/modulefiles/darshan/3
↪ .1.4:/opt/modulefiles/gcc/8.2.0:/usr/common/soft
↪ ware/modulefiles/python/3.6-anaconda-4.4:/opt/mo
↪ dulefiles/Base-opts/2.4.135-6.0.7.1_5.6__g718f89
↪ 1.ari
TARGETMODULES=craype-abudhabi:craype-abudhabi-cu:cra
↪ ype-accel-host:craype-accel-nvidia20:craype-acce
↪ l-nvidia30:craype-accel-nvidia35:craype-barcelon
↪ a:craype-broadwell:craype-haswell:craype-hugepag
↪ es128K:craype-hugepages128M:craype-hugepages16M:
↪ craype-hugepages256M:craype-hugepages2M:craype-h
↪ ugepages32M:craype-hugepages4M:craype-hugepages5
↪ 12K:craype-hugepages512M:craype-hugepages64M:cra
↪ ype-hugepages8M:craype-intel-knc:craype-interlag
↪ os:craype-interlagos-cu:craype-istanbul:craype-i
↪ vybridge:craype-mc12:craype-mc8:craype-mic-knl:c
↪ raype-network-aries:craype-network-gemini:craype
↪ -network-infiniband:craype-network-none:craype-ne
↪ twork-seastar:craype-sandybridge:craype-shanghai
↪ :craype-target-compute_node:craype-target-local_
↪ host:craype-target-native:craype-xeon:xtpe-barce
↪ lona:xtpe-interlagos:xtpe-interlagos-cu:xtpe-ist
↪ anbul:xtpe-mc12:xtpe-mc8:xtpe-network-gemini:xtp
↪ e-network-seastar:xtpe-shanghai:xtpe-target-nati
↪ ve:xtpe-xeon
JAVA_HOME=/usr/lib64/jvm/java
PE_TPSL_DEFAULT_GENCOMPILERS_GNU_mic_knl=7.1 5.3
PE_TPSL_DEFAULT_GENCOMPILERS_CRAY_interlagos=8.6
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_skylake=8.6
PE_LIBSCI_MODULE_NAME=cray-libsci/18.07.1
PE_INTEL_FIXED_PKGCONFIG_PATH=/opt/cray/pe/mpt/7.7.3
↪ /gni/mpich-intel/16.0/lib/pkgconfig
PE_TPSL_64_DEFAULT_GENCOMPS_GNU_x86_skylake=82 71 61
```

```

PE_MPICH_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/mpt/7.
↳ 7.3/gni/mpich-@PRGENV@PE_MPICH_DIR_DEFAULT64@/@
↳ PE_MPICH_GENCOMPS@/lib/pkgconfig
PE_MPICH_NV_LIBS_nvidia20=-lcudart
PE_LIBSCI_GENCOMPILERS_CRAY_x86_64=8.6
MODULEPATH=/opt/cray/pe/craype/2.5.15/modulefiles:/o
↳ pt/cray/pe/modulefiles:/opt/cray/modulefiles:/op
↳ t/modulefiles:/usr/common/software/modulefiles:/
↳ usr/syscom/nsg/modulefiles:/usr/syscom/nsg/opt/m
↳ odulefiles:/usr/common/das/modulefiles:/usr/comm
↳ on/ftg/modulefiles:/opt/cray/craype/default/modu
↳ lefiles:/opt/cray/ari/modulefiles
MAN_POSIXLY_CORRECT=1
SHMEM_ABORT_ON_ERROR=1
NSG_HOME=/usr/syscom/nsg
LOADEDMODULES=modules/3.2.10.6:nsg/1.2.0:intel/19.0.
↳ 3.199:craype-network-aries:craype/2.5.15:cray-li
↳ bsci/18.07.1:udreg/2.3.2-6.0.7.1_5.13__g5196236.
↳ ari:ugni/6.0.14.0-6.0.7.1_3.13__gea11d3d.ari:pmi
↳ /5.0.14:dmapp/7.1.1-6.0.7.1_5.45__g5a674e0.ari:g
↳ ni-headers/5.0.12.0-6.0.7.1_3.11__g3b1768f.ari:x
↳ pnem/2.2.15-6.0.7.1_5.11__g7549d06.ari:job/2.2.3
↳ -6.0.7.1_5.43__g6c4e934.ari:dvs/2.7.2.2.118-6.0.7
↳ .1_10.1__g58b37a2:alps/6.6.43-6.0.7.1_5.45__ga79
↳ 6da32.ari:rca/2.2.18-6.0.7.1_5.47__g2aa4f39.ari:
↳ atp/2.1.3:PrgEnv-intel/6.0.4:craype-mic-kl:cray
↳ -mpich/7.7.3:altd/2.0:darshan/3.1.4:gcc/8.2.0:pyt
↳ hon/3.6-anaconda-4.4:Base-opts/2.4.135-6.0.7.1_5
↳ .6__g718f891.ari
TZ=US/Pacific
SDK_HOME=/usr/lib64/jvm/java
SLURM_JOB_UID=74457
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_mic_knl=16.0
PE_TPSSL_64_DEFAULT_GENCOMPS_GNU_interlagos=82 71 53
↳ 49
PE_PKG_CONFIG_PATH=/opt/cray/pe/cti/1.0.7/lib/pkgcon
↳ fig:/opt/cray/pe/cti/1.0.6/lib/pkgconfig
PE_FFTW_DEFAULT_TARGET_x86_skylake=x86_skylake
PE_FFTW_DEFAULT_TARGET_share=share
PE_FFTW_DEFAULT_TARGET_ivybridge=ivybridge
CRAY_DMAPP_POST_LINK_OPTS=-L/opt/cray/dmapp/7.1.1-6.
↳ 0.7.1_5.45__g5a674e0.ari/lib64
SLURM_NODEID=0
PE_TPSSL_DEFAULT_GENCOMPILERS_CRAY_x86_skylake=8.6
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_skylake=6.1
PE_LIBSCI_OMP_REQUIRES_openmp=_mp
CRAY_RCA_INCLUDE_OPTS=-I/opt/cray/rca/2.2.18-6.0.7.1
↳ _5.47__g2aa4f39.ari/include
↳ -I/opt/cray/krca/2.2.4-6.0.7.1_5.43__g8505b97.ar
↳ i/include
↳ -I/opt/cray-hss-devel/9.0.0/include
SLURM_SUBMIT_DIR=/global/u1/z/USER/brickv2/stest/Aut
↳ hor-Kit
PE_MPICH_CXX_PKGCONFIG_LIBS=mpichcxx
CRAY_MPICH_DIR=/opt/cray/pe/mpt/7.7.3/gni/mpich-inte
↳ 1/16.0

```

```

PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
PE_TPSSL_64_DEFAULT_GENCOMPS_CRAY_mic_knl=86
SLURM_TASK_PID=27877
PE_MPICH_PKGCONFIG_VARIABLES=PE_MPICH_NV_LIBS_@accel
↳ erator@:PE_MPICH_ALTERNATE_LIBS_@multithreaded@:
↳ PE_MPICH_ALTERNATE_LIBS_@dpm@
PE_LIBSCI_DEFAULT_GENCOMPS_INTEL_x86_64=160
SLURM_CPUS_ON_NODE=256
PE_MPICH_PKGCONFIG_LIBS=mpich
CRAY_MPICH2_VER=7.7.3
PE_TPSSL_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe
↳ /tpsl/18.06.1/@PRGENV@/@PE_TPSSL_DEFAULT_GENCOMPS
↳ @/@PE_TPSSL_DEFAULT_TARGET@/lib/pkgconfig
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_x86_skylake=16.0
PE_TPSSL_64_DEFAULT_GENCOMPILERS_CRAY_mic_knl=8.6
PE_HDF5_DEFAULT_FIXED_PRGENV=CRAY_INTEL
CRAY_PMI_POST_LINK_OPTS=-L/opt/cray/pe/pmi/5.0.14/li
↳ b64
ENVIRONMENT=BATCH
SLURM_PROCID=0
PE_PARALLEL_NETCDF_DEFAULT_FIXED_PRGENV=CRAY_INTEL
PE_NETCDF_DEFAULT_FIXED_PRGENV=CRAY_INTEL
PE_MPICH_ALTERNATE_LIBS_multithreaded=_mt
PE_LIBSCI_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/libsc
↳ i/18.07.1/@PRGENV@/@PE_LIBSCI_GENCOMPS@/@PE_LIBS
↳ CI_TARGET@/lib/pkgconfig
PE_GA_DEFAULT_GENCOMPILERS_GNU=5.3 4.9
GPG_TTY=not a tty
SLURM_JOB_NODELIST=nid05865
PE_TPSSL_64_DEFAULT_GENCOMPS_GNU_haswell=82 71 53 49
PE_PKGCONFIG_PRODUCTS_DEFAULT=PE_PAPI
PE_NETCDF_HDF5PARALLEL_DEFAULT_VOLATILE_PRGENV=GNU
PE_MPICH_TARGET_VAR_nvidia35=-lcudart
PE_HDF5_PARALLEL_DEFAULT_VOLATILE_PRGENV=GNU
CRAY_LIBSCI_VERSION=18.07.1
QT_SYSTEM_DIR=/usr/share/desktop-data
JDK_HOME=/usr/lib64/jvm/java
SHLVL=3
HOME=/global/homes/z/USER
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_interlagos=16.0
LESS_ADVANCED_PREPROCESSOR=no
OSTYPE=linux
SLURM_LOCALID=0
ALTD_PATH=/usr/common/software/altd/2.0
PE_TPSSL_DEFAULT_VOLATILE_PRGENV=CRAY CRAY64 GNU GNU64
↳ INTEL INTEL64
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_interlagos=8.6
PE_MPICH_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/p
↳ e/mpt/7.7.3/gni/mpich-@PRGENV@PE_MPICH_DEFAULT_
↳ DIR_DEFAULT64@/@PE_MPICH_DEFAULT_GENCOMPS@/lib/p
↳ kgconfig
PE_TPSSL_DEFAULT_GENCOMPS_INTEL_sandybridge=160
PE_TPSSL_64_DEFAULT_GENCOMPS_CRAY_interlagos=86
CRAY_PMI_INCLUDE_OPTS=-I/opt/cray/pe/pmi/5.0.14/incl
↳ ude
LS_OPTIONS=-N --color=none -T 0

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```
XCURSOR_THEME=DMZ
SLURM_CLUSTER_NAME=cori
SLURM_JOB_CPUS_PER_NODE=256
SLURM_JOB_GID=74457
GCC_PATH=/opt/gcc/8.2.0
PE_MPICH_DIR_CRAY_DEFAULT64=64
PKG_CONFIG_PATH_DEFAULT=/opt/cray/pe/papi/5.6.0.3/lib_
↳ b64/pkgconfig
PE_TPSL_DEFAULT_GENCOMPILERS_CRAY_haswell=8.6
ATP_MRNET_COMM_PATH=/opt/cray/pe/atp/2.1.3/libexec/a_
↳ tp_mrnet_commnodewrapper
CRAYPE_NETWORK_TARGET=aries
WINDOWMANAGER=
PRGENVMODULES=PrgEnv-cray:PrgEnv-gnu:PrgEnv-intel:Pr_
↳ gEnv-pathscale:PrgEnv-pgi
SLURM_SUBMIT_HOST=cori11
SLURM_GTIDS=0
BASH_ENV=/global/homes/z/USER/.bashrc
PE_TPSL_DEFAULT_GENCOMPILERS_INTEL_sandybridge=16.0
PE_TPSL_DEFAULT_GENCOMPILERS_GNU_interlagos=8.2 7.1
↳ 5.3 4.9
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_mic_knl=7.1 5.3
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_haswell=7.1 5.3 4.9
SLURM_JOB_PARTITION=regular
ALTD_SELECT_USERS=
PE_TRILINOS_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_HD_
↳ F5_PARALLEL:PE_NETCDF:HDF5PARALLEL:PE_LIBSCI:PE_
↳ TPSL
PE_TPSL_DEFAULT_GENCOMPS_GNU_x86_64=82 71 53 49
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_sandybridge=16_
↳ .0
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_haswell=8.2 7.1
↳ 5.3 4.9
PE_NETCDF_DEFAULT_REQUIRED_PRODUCTS=PE_HDF5
PE_MPICH_NV_LIBS=
PE_HDF5_DEFAULT_GENCOMPS_GNU=
CRAY_LIBSCI_PREFIX_DIR=/opt/cray/pe/libsci/18.07.1/I_
↳ NTEL/16.0/x86_64
CRAY_GNI_HEADERS_INCLUDE_OPTS=-I/opt/cray/gni-header_
↳ s/5.0.12.0-6.0.7.1.3.11__g3b1768f.ari/include
PYTHONPATH=/opt/ovis/lib/python2.7/site-packages
LESS=-M -I -R
MACHTYPE=x86_64-suse-linux
LOGNAME=USER
CONDA_PYTHON_EXE=/usr/common/software/python/3.6-ana_
↳ conda-4.4/bin/python
PE_MPICH_GENCOMPS_CRAY=86
PE_TRILINOS_DEFAULT_GENCOMPS_GNU_x86_64=82 73 51 49
PE_MPICH_DEFAULT_GENCOMPILERS_CRAY=8.6
PE_LIBSCI_OMP_REQUIRES=
DMAPP_ABORT_ON_ERROR=1
CVS_RSH=ssh
PE_TPSL_DEFAULT_GENCOMPILERS_CRAY_sandybridge=8.6
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_interlagos=16.0
PE_MPICH_DEFAULT_GENCOMPS_GNU=71 51 49
PE_MPICH_DEFAULT_FIXED_PRGENV=INTEL
```

```
PE_LIBSCI_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH
DVS_INCLUDE_OPTS=-I/opt/cray/dvs/2.7.2.2.118-6.0.7.1_
↳ _10.1__g58b37a2/include
XDG_DATA_DIRS=/usr/share
TOOLMODULES=apprentice:apprentice2:atp:chapel:cray-l_
↳ gdb:cray-snplauncher:craypat:craypkg-gen:ddt:gdb_
↳ :iobuf:papi:perftools:perftools-lite:stat:totalv_
↳ iew:xt:craypat:xt-lgdb:xt-papi:xt-totalview
SSH_CONNECTION=155.97.232.235 60438 128.55.209.23 22
SLURM_JOB_ACCOUNT=ACNT
KMP_HW_SUBSET=64c
PE_TPSL_DEFAULT_GENCOMPILERS_GNU_sandybridge=8.2 7.1
↳ 5.3 4.9
PE_LIBSCI_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/_
↳ pe/libsci/18.07.1/@PRGENV@/@PE_LIBSCI_DEFAULT_GE_
↳ NCOMPS@/@PE_LIBSCI_DEFAULT_TARGET@/lib/pkgconfig
PE_GA_DEFAULT_FIXED_PRGENV=CRAY_INTEL
PE_FFTW2_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH
CRAY_PRGENVINTEL=loaded
MODULESHOME=/opt/cray/pe/modules/3.2.10.6
SLURM_JOB_NUM_NODES=1
PELOCAL_PRGENV=true
PKG_CONFIG_PATH=/usr/common/software/darshan/3.1.4/1_
↳ ib/pkgconfig:/opt/cray/rca/2.2.18-6.0.7.1.5.47__
↳ g2aa4f39.ari/lib64/pkgconfig:/opt/cray/alps/6.6._
↳ 43-6.0.7.1.5.45__ga796da32.ari/lib64/pkgconfig:/
↳ opt/cray/xpmem/2.2.15-6.0.7.1.5.11__g7549d06.ari_
↳ /lib64/pkgconfig:/opt/cray/gni-headers/5.0.12.0-_
↳ 6.0.7.1.3.11__g3b1768f.ari/lib64/pkgconfig:/opt/_
↳ cray/dmapp/7.1.1-6.0.7.1.5.45__g5a674e0.ari/lib6_
↳ 4/pkgconfig:/opt/cray/pe/pmi/5.0.14/lib64/pkgcon_
↳ fig:/opt/cray/ugni/6.0.14.0-6.0.7.1.3.13__gea1ld_
↳ 3d.ari/lib64/pkgconfig:/opt/cray/udreg/2.3.2-6.0_
↳ .7.1.5.13__g5196236.ari/lib64/pkgconfig:/opt/cra_
↳ y/pe/craype/2.5.15/pkg-config:/opt/cray/pe/iobuf_
↳ /2.0.8/lib/pkgconfig:/opt/cray/pe/fftw/2.1.5.9/1_
↳ ib/pkgconfig:/opt/cray/pe/atp/2.1.3/lib/pkgconfig
PE_PETSC_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/p_
↳ e/petsc/3.8.4.0/complex/@PRGENV@/@PE_PETSC_DEFAU_
↳ LT_GENCOMPS@/@PE_PETSC_DEFAULT_TARGET@/lib/pkgco_
↳ nfig
PE_MPICH_NV_LIBS_nvidia35=-lcudart
LESSOPEN=lessopen.sh %s
OMP_PLACES=cores
CONDA_DEFAULT_ENV=stenv
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_x86_64=160
LIBSCI_BASE_DIR=/opt/cray/pe/libsci/18.07.1
INFOPATH=/opt/gcc/8.2.0/snos/share/info
PE_TPSL_DEFAULT_GENCOMPS_INTEL_mic_knl=160
PE_TPSL_64_DEFAULT_GENCOMPS_GNU_sandybridge=82 71 53
↳ 49
PE_MPICH_NV_LIBS_nvidia60=-lcudart
PE_LIBSCI_DEFAULT_PKGCONFIG_VARIABLES=PE_LIBSCI_DEFA_
↳ ULT_OMP_REQUIRES_@openmp@:PE_SCI_EXT_LIBPATH:PE_
↳ SCI_EXT_LIBNAME
LIBSCI_VERSION=18.07.1
```



```

CRAY_CPU_TARGET=mic-kenl
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_x86_64=8.2 7.1
↪ 5.3 4.9
PE_LIBSCI_GENCOMPILERS_INTEL_x86_64=16.0
PE_FFTW_DEFAULT_TARGET_broadwell=broadwell
CRAY_ALPS_INCLUDE_OPTS=-I/opt/cray/alps/6.6.43-6.0.7
↪ .1_5.45__ga796da32.ari/include
CRAY_PRE_COMPILE_OPTS=-hnetwork=aries
NERSC_HOST=cori
XDG_RUNTIME_DIR=/run/user/74457
craype_already_loaded=0
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_x86_64=86
PE_PAPI_DEFAULT_ACCELL_FAMILY_LIBS=
PE_LIBSCI_REQUIRED_PRODUCTS=PE_MPICH
CRAY_XPMEM_INCLUDE_OPTS=-I/opt/cray/xpmem/2.2.15-6.0
↪ .7.1_5.11__g7549d06.ari/include
CRAY_UGNI_INCLUDE_OPTS=-I/opt/cray/ugni/6.0.14.0-6.0
↪ .7.1_3.13__gea11d3d.ari/include
PE_TPSL_DEFAULT_GENCOMPS_INTEL_haswell=160
PE_LIBSCI_GENCOMPS_GNU_x86_64=71 61 51 49
PE_LIBSCI_DEFAULT_GENCOMPILERS_GNU_x86_64=7.1 6.1 5.1
↪ 4.9
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
PE_FFTW_DEFAULT_TARGET_x86_64=x86_64
ATP_HOME=/opt/cray/pe/atp/2.1.3
LESSCLOSE=lessclose.sh %s %s
ALTD_WORKDIR=/global/cscratch1/altd/logs
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_x86_skylake=16
↪ .0
PE_SMA_DEFAULT_DIR_PGI_DEFAULT64=64
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_haswell=16.0
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_interlagos=7.1 5.3
↪ 4.9
PE_PAPI_DEFAULT_ACCEL_LIBS=
PE_INTEL_DEFAULT_FIXED_PKGCONFIG_PATH=/opt/cray/pe/p
↪ arallel-netcdf/1.8.1.3/INTEL/16.0/lib/pkgconfig:
↪ /opt/cray/pe/netcdf-hdf5parallel/4.6.1.3/INTEL/1
↪ 6.0/lib/pkgconfig:/opt/cray/pe/netcdf/4.6.1.3/IN
↪ TEL/16.0/lib/pkgconfig:/opt/cray/pe/mpt/7.7.3/gn
↪ i/mpich-intel/16.0/lib/pkgconfig:/opt/cray/pe/hd
↪ f5-parallel/1.10.2.0/INTEL/16.0/lib/pkgconfig:/o
↪ pt/cray/pe/hdf5/1.10.2.0/INTEL/16.0/lib/pkgconfi
↪ g:/opt/cray/pe/ga/5.3.0.8/INTEL/18.0/lib/pkgconf
↪ ig
PE_GA_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/g
↪ a/5.3.0.8/@PRGENV@/PE_GA_DEFAULT_GENCOMPS@/lib/
↪ pkgconfig
PE_GA_DEFAULT_GENCOMPS_GNU=53 49
PE_FFTW_DEFAULT_TARGET_haswell=haswell

```

```

CRAY_LD_LIBRARY_PATH=/usr/common/software/darshan/3.
↪ 1.4/lib:/opt/cray/pe/mpt/7.7.3/gni/mpich-intel/1
↪ 6.0/lib:/opt/cray/rca/2.2.18-6.0.7.1_5.47__g2aa4
↪ f39.ari/lib64:/opt/cray/alps/6.6.43-6.0.7.1_5.45
↪ __ga796da32.ari/lib64:/opt/cray/xpmem/2.2.15-6.0
↪ .7.1_5.11__g7549d06.ari/lib64:/opt/cray/dmapp/7.
↪ 1.1-6.0.7.1_5.45__g5a674e0.ari/lib64:/opt/cray/p
↪ e/pmi/5.0.14/lib64:/opt/cray/ugni/6.0.14.0-6.0.7
↪ .1_3.13__gea11d3d.ari/lib64:/opt/cray/udreg/2.3.
↪ 2-6.0.7.1_5.13__g5196236.ari/lib64:/opt/cray/pe/
↪ libsci/18.07.1/INTEL/16.0/x86_64/lib
G_BROKEN_FILENAMES=1
SCRATCH=/global/cscratch1/sd/USER
SLURM_MEM_PER_NODE=89088
intel_already_loaded=0
PE_PETSC_DEFAULT_GENCOMPS_INTEL_x86_64=160
PE_PETSC_DEFAULT_GENCOMPS_GNU_x86_64=71 53 49
PE_PETSC_DEFAULT_GENCOMPS_CRAY_haswell=86
PE_MPICH_DEFAULT_DIR_CRAY_DEFAULT64=64
JAVA_ROOT=/usr/lib64/jvm/java
COLORTERM=1
BASH_FUNC_module%=() { eval
↪ `opt/cray/pe/modules/3.2.10.6/bin/modulecmd
↪ bash $*`
}
+ lsb_release -a
LSB Version:          n/a
Distributor ID:       SUSE
Description:          SUSE Linux Enterprise Server 12
↪ SP3
Release:              12.3
Codename:              n/a
+ uname -a
Linux nid05865 4.4.103-6.38.4.0.153-cray_ari_c #1 SMP
↪ Thu Nov 1 16:05:05 UTC 2018 (gef8fef) x86_64
↪ x86_64 x86_64 GNU/Linux
+ lscpu
Architecture:         x86_64
CPU op-mode(s):       32-bit, 64-bit
Byte Order:           Little Endian
CPU(s):               272
On-line CPU(s) list:  0-271
Thread(s) per core:   4
Core(s) per socket:   68
Socket(s):            1
NUMA node(s):         1
Vendor ID:            GenuineIntel
CPU family:           6
Model:                87
Model name:           Intel(R) Xeon Phi(TM) CPU 7250
↪ @ 1.40GHz
Stepping:             1
CPU MHz:              1401.000
CPU max MHz:          1401.0000
CPU min MHz:          1000.0000
BogoMIPS:             2799.85

```

Exploiting Reuse and Vectorization in Blocked Stencil Computations on CPUs and GPUs

```

L1d cache:      32K
L1i cache:      32K
L2 cache:       1024K
NUMA node0 CPU(s): 0-271
Flags:          fpu vme de pse tsc msr pae mce
↳ cx8 apic sep mtrr pge mca cmov pat pse36 clflush
↳ dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx
↳ pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs
↳ bts rep_good nopl xtopology nonstop_tsc
↳ ring3mwait aperfmperf eagerfpu pni pclmulqdq
↳ dtes64 monitor ds_cpl est tm2 ssse3 fma cx16 xtpr
↳ pdcm sse4_1 sse4_2 x2apic movbe popcnt
↳ tsc_deadline_timer aes xsave avx f16c rdrand
↳ lahf_lm abm 3dnowprefetch ida arat epb pln pts
↳ dtherm spec_ctrl kaiser fsgsbase tsc_adjust bmi1
↳ avx2 smep bmi2 erms avx512f rdseed adx avx512pf
↳ avx512er avx512cd xsaveopt
+ cat /proc/meminfo
MemTotal:      98844412 kB
MemFree:       95952572 kB
MemAvailable:  95397668 kB
Buffers:       7344 kB
Cached:        356992 kB
SwapCached:    0 kB
Active:        131524 kB
Inactive:      318224 kB
Active(anon):  116272 kB
Inactive(anon): 219656 kB
Active(file):   15252 kB
Inactive(file): 98568 kB
Unevictable:   8316 kB
Mlocked:       8316 kB
SwapTotal:     0 kB
SwapFree:      0 kB
Dirty:         0 kB
Writeback:     0 kB
AnonPages:     93960 kB
Mapped:        59104 kB
Shmem:         249244 kB
Slab:          1013548 kB
SReclaimable:  31336 kB
SUnreclaim:   982212 kB
KernelStack:   40096 kB
PageTables:    4648 kB
NFS_Unstable:  0 kB
Bounce:        0 kB
WritebackTmp:  0 kB
CommitLimit:   49422204 kB
Committed_AS:  588652 kB
VmallocTotal:  34359738367 kB
VmallocUsed:   0 kB
VmallocChunk:  0 kB
HardwareCorrupted: 0 kB
HugePages_Total: 0
HugePages_Free: 0
HugePages_Rsvd: 0
HugePages_Surp: 0

Hugepagesize: 2048 kB
DirectMap4k: 10088 kB
DirectMap2M: 1988608 kB
DirectMap1G: 100663296 kB
+ inxi -F -c0
./collect_environment.sh: line 14: inxi: command not
↳ found
+ lsblk -a
NAME MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
loop0 7:0 0 192K 0 loop /var/opt/cray/imps-d_
↳ istribution/squash/mounts/global
loop1 7:1 0 4.5M 0 loop
↳ /var/opt/cray/imps-distribution/squash/mounts/cl_
↳ e.coricle6.0.up07.20190410084149
loop2 7:2 0 2.4G 1 loop /.rootfs_lower_ro
loop3 7:3 0 22.2G 1 loop
↳ /var/opt/cray/imps-image-binding/PE_x86_64/squas_
↳ h_mounts/squashfs_4EypVE_mount_point
loop4 7:4 0 0 loop
loop5 7:5 0 0 loop
loop6 7:6 0 0 loop
loop7 7:7 0 0 loop
loop8 7:8 0 0 loop
loop9 7:9 0 0 loop
loop10 7:10 0 0 loop
loop11 7:11 0 0 loop
loop12 7:12 0 0 loop
loop13 7:13 0 0 loop
loop14 7:14 0 0 loop
loop15 7:15 0 0 loop
loop16 7:16 0 0 loop
loop17 7:17 0 0 loop
loop18 7:18 0 0 loop
loop19 7:19 0 0 loop
loop20 7:20 0 0 loop
loop21 7:21 0 0 loop
loop22 7:22 0 0 loop
loop23 7:23 0 0 loop
loop24 7:24 0 0 loop
loop25 7:25 0 0 loop
loop26 7:26 0 0 loop
loop27 7:27 0 0 loop
loop28 7:28 0 0 loop
loop29 7:29 0 0 loop
loop30 7:30 0 0 loop
loop31 7:31 0 0 loop
loop32 7:32 0 0 loop
loop33 7:33 0 0 loop
loop34 7:34 0 0 loop
loop35 7:35 0 0 loop
loop36 7:36 0 0 loop
loop37 7:37 0 0 loop
loop38 7:38 0 0 loop
loop39 7:39 0 0 loop
loop40 7:40 0 0 loop
loop41 7:41 0 0 loop

```

```

loop42  7:42  0      0 loop
loop43  7:43  0      0 loop
loop44  7:44  0      0 loop
loop45  7:45  0      0 loop
loop46  7:46  0      0 loop
loop47  7:47  0      0 loop
loop48  7:48  0      0 loop
loop49  7:49  0      0 loop
loop50  7:50  0      0 loop
loop51  7:51  0      0 loop
loop52  7:52  0      0 loop
loop53  7:53  0      0 loop
loop54  7:54  0      0 loop
loop55  7:55  0      0 loop
loop56  7:56  0      0 loop
loop57  7:57  0      0 loop
loop58  7:58  0      0 loop
loop59  7:59  0      0 loop
loop60  7:60  0      0 loop
loop61  7:61  0      0 loop
loop62  7:62  0      0 loop
loop63  7:63  0      0 loop
loop64  7:64  0      0 loop
loop65  7:65  0      0 loop
loop66  7:66  0      0 loop
loop67  7:67  0      0 loop
loop68  7:68  0      0 loop
loop69  7:69  0      0 loop
loop70  7:70  0      0 loop
loop71  7:71  0      0 loop
loop72  7:72  0      0 loop
loop73  7:73  0      0 loop
loop74  7:74  0      0 loop
loop75  7:75  0      0 loop
loop76  7:76  0      0 loop
loop77  7:77  0      0 loop
loop78  7:78  0      0 loop
loop79  7:79  0      0 loop
loop80  7:80  0      0 loop
loop81  7:81  0      0 loop
loop82  7:82  0      0 loop
loop83  7:83  0      0 loop
loop84  7:84  0      0 loop
loop85  7:85  0      0 loop
loop86  7:86  0      0 loop
loop87  7:87  0      0 loop
loop88  7:88  0      0 loop
loop89  7:89  0      0 loop
loop90  7:90  0      0 loop
loop91  7:91  0      0 loop
loop92  7:92  0      0 loop
loop93  7:93  0      0 loop
loop94  7:94  0      0 loop
loop95  7:95  0      0 loop
loop96  7:96  0      0 loop
loop97  7:97  0      0 loop
loop98  7:98  0      0 loop

```

```

loop99  7:99  0      0 loop
loop100 7:100  0      0 loop
loop101 7:101  0      0 loop
loop102 7:102  0      0 loop
loop103 7:103  0      0 loop
loop104 7:104  0      0 loop
loop105 7:105  0      0 loop
loop106 7:106  0      0 loop
loop107 7:107  0      0 loop
loop108 7:108  0      0 loop
loop109 7:109  0      0 loop
loop110 7:110  0      0 loop
loop111 7:111  0      0 loop
loop112 7:112  0      0 loop
loop113 7:113  0      0 loop
loop114 7:114  0      0 loop
loop115 7:115  0      0 loop
loop116 7:116  0      0 loop
loop117 7:117  0      0 loop
loop118 7:118  0      0 loop
loop119 7:119  0      0 loop
loop120 7:120  0      0 loop
loop121 7:121  0      0 loop
loop122 7:122  0      0 loop
loop123 7:123  0      0 loop
loop124 7:124  0      0 loop
loop125 7:125  0      0 loop
loop126 7:126  0      0 loop
loop127 7:127  0      0 loop
+ lsscsi -s
+ module list
++ /opt/cray/pe/modules/3.2.10.6/bin/modulecmd bash
↵ list
Currently Loaded Modulefiles:
  1) modules/3.2.10.6
  2) intel/19.0.3.199
  3) craype-network-aries
  4) craype/2.5.15
  5) cray-libsci/18.07.1
  6) udreg/2.3.2-6.0.7.1_5.13__g5196236.ari
  7) ugni/6.0.14.0-6.0.7.1_3.13__gea11d3d.ari
  8) pmi/5.0.14
  9) dmapp/7.1.1-6.0.7.1_5.45__g5a674e0.ari
 10) gni-headers/5.0.12.0-6.0.7.1_3.11__g3b1768f.ari
 11) xpmem/2.2.15-6.0.7.1_5.11__g7549d06.ari
 12) job/2.2.3-6.0.7.1_5.43__g6c4e934.ari
 13) dvs/2.7_2.2.118-6.0.7.1_10.1__g58b37a2
 14) alps/6.6.43-6.0.7.1_5.45__ga796da32.ari
 15) rca/2.2.18-6.0.7.1_5.47__g2aa4f39.ari
 16) atp/2.1.3
 17) PrgEnv-intel/6.0.4
 18) craype-mic-knl
 19) cray-mpich/7.7.3
 20) altd/2.0
 21) darshan/3.1.4
 22) gcc/8.2.0

```

```
23) python/3.6-anaconda-4.4
24) Base-opts/2.4.135-6.0.7.1_5.6__g718f891.ari
+ eval
+ nvidia-smi
NVIDIA-SMI has failed because it couldn't communicate
↳ with the NVIDIA driver. Make sure that the latest
↳ NVIDIA driver is installed and running.

+ lshw -short -quiet -sanitize
+ cat
./collect_environment.sh: line 19: lshw: command not
↳ found
+ lspci
./collect_environment.sh: line 19: lspci: command not
↳ found
```

ARTIFACT EVALUATION

Verification and validation studies: We applied our approach on synthetic stencils with varying shape and order to expose how these factors affect the performance of our approach.

We also gathered hardware profiling data and looked at the generated assembly to identify the differences between our code and the baseline. We performed detailed analysis based on these results.

The results from these analysis agree with the theory we presented in the paper.

Accuracy and precision of timings: On KNL and Skylake-X, for all the stencil variants we drop the first iteration due to cold start, and time the stencils kernel running consecutively until total time exceeds 5 seconds to reduce variation between each stencil iterations. This is enough time for taking the average of at least 25 iterations with the code generated using our approach. Time is taken using `omp_wtime()`. The throughput is calculated based on the average time taken for one iteration.

On P100, each stencil is run one iteration untimed due to cold start and then timed for 100 iterations using `cudaEvent_t`. The throughput is calculated based on the average time taken for one iteration.

Used manufactured solutions or spectral properties: Not applicable

Quantified the sensitivity of results to initial conditions and/or parameters of the computational environment: No

Controls, statistics, or other steps taken to make the measurements and analyses robust to variability and unknowns in the system. No